

# Linear Models of Regression

K Sri Rama Murty

IIT Hyderabad

`ksrm@ee.iith.ac.in`

January 28, 2022

# Regression

# Regression

- Predict target variable(s)  $t \in \mathbb{R}$  given D-dimensional input vector  $\mathbf{x}$

# Regression

- Predict target variable(s)  $t \in \mathbb{R}$  given D-dimensional input vector  $\mathbf{x}$
- E.g. Weight estimation, Share market prediction, 3D image from 2D

# Regression

- Predict target variable(s)  $t \in \mathbb{R}$  given D-dimensional input vector  $\mathbf{x}$
- E.g. Weight estimation, Share market prediction, 3D image from 2D
- Target can be estimated as a linear combination of inputs

$$\hat{t} = y(\mathbf{x}, \mathbf{w}) = w_0 + w_1 x_1 + w_2 x_2 + \cdots w_D x_D = \mathbf{w}^T \mathbf{x}$$

$$\mathbf{x} = [1 \ x_1 \ x_2 \ \cdots \ x_D]^T \quad \mathbf{w} = [w_0 \ w_1 \ w_2 \ \cdots \ w_D]^T$$

# Regression

- Predict target variable(s)  $t \in \mathbb{R}$  given D-dimensional input vector  $\mathbf{x}$
- E.g. Weight estimation, Share market prediction, 3D image from 2D
- Target can be estimated as a linear combination of inputs

$$\hat{t} = y(\mathbf{x}, \mathbf{w}) = w_0 + w_1 x_1 + w_2 x_2 + \cdots w_D x_D = \mathbf{w}^T \mathbf{x}$$

$$\mathbf{x} = [1 \ x_1 \ x_2 \ \cdots \ x_D]^T \quad \mathbf{w} = [w_0 \ w_1 \ w_2 \ \cdots \ w_D]^T$$

- Determine the model parameters  $\mathbf{w}$  to minimize error on labeled training data

$$\mathcal{S} = \{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \cdots (\mathbf{x}_N, t_N)\}$$

# Regression

- Predict target variable(s)  $t \in \mathbb{R}$  given D-dimensional input vector  $\mathbf{x}$
- E.g. Weight estimation, Share market prediction, 3D image from 2D
- Target can be estimated as a linear combination of inputs

$$\hat{t} = y(\mathbf{x}, \mathbf{w}) = w_0 + w_1 x_1 + w_2 x_2 + \cdots w_D x_D = \mathbf{w}^T \mathbf{x}$$

$$\mathbf{x} = [1 \ x_1 \ x_2 \ \cdots \ x_D]^T \quad \mathbf{w} = [w_0 \ w_1 \ w_2 \ \cdots \ w_D]^T$$

- Determine the model parameters  $\mathbf{w}$  to minimize error on labeled training data

$$\mathcal{S} = \{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \cdots (\mathbf{x}_N, t_N)\}$$

- Need to define a loss function for optimizing model parameters  $\mathbf{w}$

# Least Squares Criterion to Determine $\mathbf{w}$



# Least Squares Criterion to Determine $\mathbf{w}$

- Estimated target of  $n^{th}$  example

$$\hat{t}_n = y(\mathbf{x}_n, \mathbf{w}) = \mathbf{w}^T \mathbf{x}_n$$

# Least Squares Criterion to Determine $\mathbf{w}$

- Estimated target of  $n^{th}$  example

$$\hat{t}_n = y(\mathbf{x}_n, \mathbf{w}) = \mathbf{w}^T \mathbf{x}_n$$

- Error in estimation

$$e_n = t_n - y_n \quad n = 1, 2, \dots, N$$

# Least Squares Criterion to Determine $\mathbf{w}$

- Estimated target of  $n^{th}$  example

$$\hat{t}_n = y(\mathbf{x}_n, \mathbf{w}) = \mathbf{w}^T \mathbf{x}_n$$

- Error in estimation

$$e_n = t_n - y_n \quad n = 1, 2, \dots, N$$

- Overall error on training set

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N e_n^2$$

# Least Squares Criterion to Determine $\mathbf{w}$

- Estimated target of  $n^{th}$  example

$$\hat{t}_n = y(\mathbf{x}_n, \mathbf{w}) = \mathbf{w}^T \mathbf{x}_n$$

- Error in estimation

$$e_n = t_n - y_n \quad n = 1, 2, \dots, N$$

- Overall error on training set

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N e_n^2$$

- Estimate  $\mathbf{w}$  to minimize  $J(\mathbf{w})$

$$\mathbf{w}_* = \arg \min_{\mathbf{w}} J(\mathbf{w})$$

# Least Squares Criterion to Determine $\mathbf{w}$

- Estimated target of  $n^{th}$  example

$$\hat{t}_n = y(\mathbf{x}_n, \mathbf{w}) = \mathbf{w}^T \mathbf{x}_n$$

- Error in estimation

$$e_n = t_n - y_n \quad n = 1, 2, \dots, N$$

- Overall error on training set

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N e_n^2$$

- Estimate  $\mathbf{w}$  to minimize  $J(\mathbf{w})$

$$\mathbf{w}_* = \arg \min_{\mathbf{w}} J(\mathbf{w})$$

- Formulating in matrix notation

$$\mathbf{y}_{N \times 1} = \mathbf{X}_{N \times D+1} \mathbf{w}_{D+1 \times 1}$$

$$\mathbf{X} = [\mathbf{x}_1 \ \mathbf{x}_2 \ \cdots \ \mathbf{x}_N]^T$$

# Least Squares Criterion to Determine $\mathbf{w}$

- Estimated target of  $n^{th}$  example

$$\hat{t}_n = y(\mathbf{x}_n, \mathbf{w}) = \mathbf{w}^T \mathbf{x}_n$$

- Error in estimation

$$e_n = t_n - y_n \quad n = 1, 2, \dots, N$$

- Overall error on training set

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N e_n^2$$

- Estimate  $\mathbf{w}$  to minimize  $J(\mathbf{w})$

$$\mathbf{w}_* = \arg \min_{\mathbf{w}} J(\mathbf{w})$$

- Formulating in matrix notation

$$\mathbf{y}_{N \times 1} = \mathbf{X}_{N \times D+1} \mathbf{w}_{D+1 \times 1}$$

$$\mathbf{X} = [\mathbf{x}_1 \ \mathbf{x}_2 \ \cdots \ \mathbf{x}_N]^T$$

- LSE can be expressed as

$$J(\mathbf{w}) = \frac{1}{2} \text{Tr}[(\mathbf{t} - \mathbf{y})(\mathbf{t} - \mathbf{y})^T]$$

# Least Squares Criterion to Determine $\mathbf{w}$

- Estimated target of  $n^{th}$  example

$$\hat{t}_n = y(\mathbf{x}_n, \mathbf{w}) = \mathbf{w}^T \mathbf{x}_n$$

- Error in estimation

$$e_n = t_n - y_n \quad n = 1, 2, \dots, N$$

- Overall error on training set

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N e_n^2$$

- Estimate  $\mathbf{w}$  to minimize  $J(\mathbf{w})$

$$\mathbf{w}_* = \arg \min_{\mathbf{w}} J(\mathbf{w})$$

- Formulating in matrix notation

$$\mathbf{y}_{N \times 1} = \mathbf{X}_{N \times D+1} \mathbf{w}_{D+1 \times 1}$$

$$\mathbf{X} = [\mathbf{x}_1 \ \mathbf{x}_2 \ \cdots \ \mathbf{x}_N]^T$$

- LSE can be expressed as


$$J(\mathbf{w}) = \frac{1}{2} \text{Tr}[(\mathbf{t} - \mathbf{y})(\mathbf{t} - \mathbf{y})^T]$$

- Equating derivative w.r.t  $\mathbf{w}$  to 0


$$\begin{aligned} \nabla_{\mathbf{w}} J(\mathbf{w}) &= \nabla_{\mathbf{y}} J(\mathbf{w}) \nabla_{\mathbf{w}} \mathbf{y} \\ &= \mathbf{X}^T (\mathbf{t} - \mathbf{X} \mathbf{w}) = \mathbf{0} \end{aligned}$$

$$\boxed{\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{t}}$$

# Geometric Interpretation of Least Squares




Error Surface of a Linear Neuron with Two Input Weights







# Geometric Interpretation of Least Squares



- Given  $N$  examples, the target vector  $\mathbf{t} \in \mathbb{R}^N$  and columns of  $\mathbf{X} \in \mathbb{R}^N$




# Geometric Interpretation of Least Squares




- Given  $N$  examples, the target vector  $\mathbf{t} \in \mathbb{R}^N$  and columns of  $\mathbf{X} \in \mathbb{R}^N$
- Let  $S$  denote a subspace spanned by columns of  $X$  in  $N$ -dim space


Error Surface of a Linear Neuron with Two Input Weights



# Geometric Interpretation of Least Squares

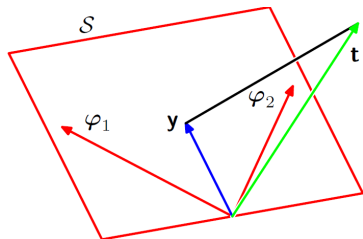


Error Surface of a Linear Neuron with Two Input Weights

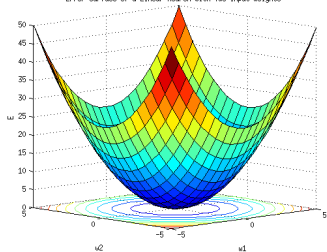


- Given  $N$  examples, the target vector  $\mathbf{t} \in \mathbb{R}^N$  and columns of  $\mathbf{X} \in \mathbb{R}^N$
- Let  $\mathcal{S}$  denote a subspace spanned by columns of  $\mathbf{X}$  in  $N$ -dim space
- $\mathbf{y} = \mathbf{X}\mathbf{w} \in \mathcal{S}$ , being a linear combination of columns of  $\mathbf{X}$

# Geometric Interpretation of Least Squares




Error Surface of a Linear Neuron with Two Input Weights




- Given  $N$  examples, the target vector  $\mathbf{t} \in \mathbb{R}^N$  and columns of  $\mathbf{X} \in \mathbb{R}^N$
- Let  $\mathcal{S}$  denote a subspace spanned by columns of  $\mathbf{X}$  in  $N$ -dim space
- $\mathbf{y} = \mathbf{X}\mathbf{w} \in \mathcal{S}$ , being a linear combination of columns of  $\mathbf{X}$
- For the LS optimality criterion

# Geometric Interpretation of Least Squares




Error Surface of a Linear Neuron with Two Input Weights




- Given  $N$  examples, the target vector  $\mathbf{t} \in \mathbb{R}^N$  and columns of  $\mathbf{X} \in \mathbb{R}^N$
- Let  $\mathcal{S}$  denote a subspace spanned by columns of  $\mathbf{X}$  in  $N$ -dim space
- $\mathbf{y} = \mathbf{X}\mathbf{w} \in \mathcal{S}$ , being a linear combination of columns of  $\mathbf{X}$
- For the LS optimality criterion
  - $\mathbf{y}$  is orthogonal projection of  $\mathbf{t}$  on  $\mathcal{S}$

# Geometric Interpretation of Least Squares




Error Surface of a Linear Neuron with Two Input Weights




- Given  $N$  examples, the target vector  $\mathbf{t} \in \mathbb{R}^N$  and columns of  $\mathbf{X} \in \mathbb{R}^N$
- Let  $\mathcal{S}$  denote a subspace spanned by columns of  $\mathbf{X}$  in  $N$ -dim space
- $\mathbf{y} = \mathbf{X}\mathbf{w} \in \mathcal{S}$ , being a linear combination of columns of  $\mathbf{X}$
- For the LS optimality criterion
  - $\mathbf{y}$  is orthogonal projection of  $\mathbf{t}$  on  $\mathcal{S}$
  - Error surface  $J(\mathbf{w})$  is convex

# Geometric Interpretation of Least Squares




Error Surface of a Linear Neuron with Two Input Weights




- Given  $N$  examples, the target vector  $\mathbf{t} \in \mathbb{R}^N$  and columns of  $\mathbf{X} \in \mathbb{R}^N$
- Let  $\mathcal{S}$  denote a subspace spanned by columns of  $\mathbf{X}$  in  $N$ -dim space
- $\mathbf{y} = \mathbf{X}\mathbf{w} \in \mathcal{S}$ , being a linear combination of columns of  $\mathbf{X}$
- For the LS optimality criterion
  - $\mathbf{y}$  is orthogonal projection of  $\mathbf{t}$  on  $\mathcal{S}$
  - Error surface  $J(\mathbf{w})$  is convex
  - Sim. to Wiener filter:  $\mathbf{w} = \mathbf{R}_{xx}^{-1} \mathbf{r}_{xt}$

# Geometric Interpretation of Least Squares





Error Surface of a Linear Neuron with Two Input Weights



- Given  $N$  examples, the target vector  $\mathbf{t} \in \mathbb{R}^N$  and columns of  $\mathbf{X} \in \mathbb{R}^N$
- Let  $\mathcal{S}$  denote a subspace spanned by columns of  $\mathbf{X}$  in  $N$ -dim space
- $\mathbf{y} = \mathbf{X}\mathbf{w} \in \mathcal{S}$ , being a linear combination of columns of  $\mathbf{X}$
- For the LS optimality criterion
  - $\mathbf{y}$  is orthogonal projection of  $\mathbf{t}$  on  $\mathcal{S}$
  - Error surface  $J(\mathbf{w})$  is convex
  - Sim. to Wiener filter:  $\mathbf{w} = \mathbf{R}_{xx}^{-1} \mathbf{r}_{xt}$
  - Also referred to as pseudo inverse sol.



# Nonlinear Input-Output Relations






- Polynomial curve fitting can be used to model nonlinear i/o relation


$$\begin{aligned}\hat{t} &= y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \cdots + w_Mx^M \\ &= \mathbf{w}^T \phi(\mathbf{x}) \quad (\text{Model is linear in } \mathbf{w})\end{aligned}$$

- $\phi(.) : \mathbb{R}^1 \rightarrow \mathbb{R}^M$  - nonlinear transformation to higher dim. space


# Kernel Examples




# Kernel Examples



$$\phi_j(x) = x^j$$




$$\phi_j(x) = \exp\left(-\frac{x-\mu_j}{2\sigma^2}\right)$$




$$\phi_j(x) = \sigma\left(\frac{x-\mu_j}{s}\right)$$


# Kernel Examples



$$\phi_j(x) = x^j$$




$$\phi_j(x) = \exp\left(-\frac{x-\mu_j}{2\sigma^2}\right)$$




$$\phi_j(x) = \sigma\left(\frac{x-\mu_j}{s}\right)$$


# Kernel Examples



$$\phi_j(x) = x^j$$




$$\phi_j(x) = \exp\left(-\frac{x-\mu_j}{2\sigma^2}\right)$$




$$\phi_j(x) = \sigma\left(\frac{x-\mu_j}{s}\right)$$

- Explicit vs Implicit kernels

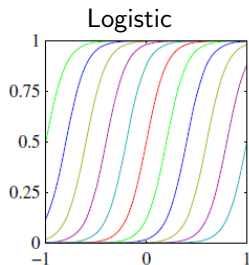
# Kernel Examples



$$\phi_j(x) = x^j$$




$$\phi_j(x) = \exp\left(-\frac{x - \mu_j}{2\sigma^2}\right)$$




$$\phi_j(x) = \sigma\left(\frac{x - \mu_j}{s}\right)$$

- Explicit vs Implicit kernels
  - Explicit representation for  $\phi(\mathbf{x})$  is available or not


# Kernel Examples



$$\phi_j(x) = x^j$$




$$\phi_j(x) = \exp\left(-\frac{x-\mu_j}{2\sigma^2}\right)$$




$$\phi_j(x) = \sigma\left(\frac{x-\mu_j}{s}\right)$$

- Explicit vs Implicit kernels
  - Explicit representation for  $\phi(\mathbf{x})$  is available or not
- Global vs Local kernels


# Kernel Examples



$$\phi_j(x) = x^j$$



$$\phi_j(x) = \exp\left(-\frac{x-\mu_j}{2\sigma^2}\right)$$




$$\phi_j(x) = \sigma\left(\frac{x-\mu_j}{s}\right)$$


- Explicit vs Implicit kernels
  - Explicit representation for  $\phi(\mathbf{x})$  is available or not
- Global vs Local kernels
  - Changes in one region of input space affect all other regions




# Kernel Examples



$$\phi_j(x) = x^j$$



$$\phi_j(x) = \exp\left(-\frac{x-\mu_j}{2\sigma^2}\right)$$



$$\phi_j(x) = \sigma\left(\frac{x-\mu_j}{s}\right)$$

- Explicit vs Implicit kernels
  - Explicit representation for  $\phi(\mathbf{x})$  is available or not
- Global vs Local kernels
  - Changes in one region of input space affect all other regions
  - Local kernels are preferable for functions with varying characteristics

# Least Squares Regression in Kernel Space

# Least Squares Regression in Kernel Space

- If  $t_n$  is nonlinearly related to  $\mathbf{x}_n$ , perform regression in kernel space.

# Least Squares Regression in Kernel Space

- If  $t_n$  is nonlinearly related to  $\mathbf{x}_n$ , perform regression in kernel space.
- Let  $\phi : \mathbb{R}^D \rightarrow \mathbb{R}^M$ ,  $M > D$  is a nonlinear kernel mapping

# Least Squares Regression in Kernel Space

- If  $t_n$  is nonlinearly related to  $\mathbf{x}_n$ , perform regression in kernel space.
- Let  $\phi : \mathbb{R}^D \rightarrow \mathbb{R}^M$ ,  $M > D$  is a nonlinear kernel mapping
- $\mathbf{x}_n = [x_{n1} \ x_{n2}]^T \in \mathbb{R}^2$  can be mapped using 2<sup>nd</sup> order polynomial kernel as  $\phi(\mathbf{x}_n) = [1 \ x_{n1} \ x_{n2} \ x_{n1}^2 \ x_{n2}^2 \ x_{n1}x_{n2}]^T \in \mathbb{R}^6$

# Least Squares Regression in Kernel Space

- If  $t_n$  is nonlinearly related to  $\mathbf{x}_n$ , perform regression in kernel space.
- Let  $\phi : \mathbb{R}^D \rightarrow \mathbb{R}^M$ ,  $M > D$  is a nonlinear kernel mapping
- $\mathbf{x}_n = [x_{n1} \ x_{n2}]^T \in \mathbb{R}^2$  can be mapped using 2<sup>nd</sup> order polynomial kernel as  $\phi(\mathbf{x}_n) = [1 \ x_{n1} \ x_{n2} \ x_{n1}^2 \ x_{n2}^2 \ x_{n1}x_{n2}]^T \in \mathbb{R}^6$
- The target  $t_n$  is regressed from the kernel representation  $\phi(\mathbf{x}_n)$  as

$$\hat{t}_n = \mathbf{w}^T \phi(\mathbf{x}_n) \quad \mathbf{w} \in \mathbb{R}^M$$

# Least Squares Regression in Kernel Space

- If  $t_n$  is nonlinearly related to  $\mathbf{x}_n$ , perform regression in kernel space.
- Let  $\phi : \mathbb{R}^D \rightarrow \mathbb{R}^M$ ,  $M > D$  is a nonlinear kernel mapping
- $\mathbf{x}_n = [x_{n1} \ x_{n2}]^T \in \mathbb{R}^2$  can be mapped using 2<sup>nd</sup> order polynomial kernel as  $\phi(\mathbf{x}_n) = [1 \ x_{n1} \ x_{n2} \ x_{n1}^2 \ x_{n2}^2 \ x_{n1}x_{n2}]^T \in \mathbb{R}^6$
- The target  $t_n$  is regressed from the kernel representation  $\phi(\mathbf{x}_n)$  as

$$\hat{t}_n = \mathbf{w}^T \phi(\mathbf{x}_n) \quad \mathbf{w} \in \mathbb{R}^M$$

- The regression coefficients are given by  $\mathbf{w}_* = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$

# Least Squares Regression in Kernel Space

- If  $t_n$  is nonlinearly related to  $\mathbf{x}_n$ , perform regression in kernel space.
- Let  $\phi : \mathbb{R}^D \rightarrow \mathbb{R}^M$ ,  $M > D$  is a nonlinear kernel mapping
- $\mathbf{x}_n = [x_{n1} \ x_{n2}]^T \in \mathbb{R}^2$  can be mapped using 2<sup>nd</sup> order polynomial kernel as  $\phi(\mathbf{x}_n) = [1 \ x_{n1} \ x_{n2} \ x_{n1}^2 \ x_{n2}^2 \ x_{n1}x_{n2}]^T \in \mathbb{R}^6$
- The target  $t_n$  is regressed from the kernel representation  $\phi(\mathbf{x}_n)$  as

$$\hat{t}_n = \mathbf{w}^T \phi(\mathbf{x}_n) \quad \mathbf{w} \in \mathbb{R}^M$$

- The regression coefficients are given by  $\mathbf{w}_* = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$
- DNNs can be used to learn data-dependent nonlinear transf.  $\phi(\mathbf{x}_n)$





# Least Squares Regression in Kernel Space

- If  $t_n$  is nonlinearly related to  $\mathbf{x}_n$ , perform regression in kernel space.
- Let  $\phi : \mathbb{R}^D \rightarrow \mathbb{R}^M$ ,  $M > D$  is a nonlinear kernel mapping
- $\mathbf{x}_n = [x_{n1} \ x_{n2}]^T \in \mathbb{R}^2$  can be mapped using 2<sup>nd</sup> order polynomial kernel as  $\phi(\mathbf{x}_n) = [1 \ x_{n1} \ x_{n2} \ x_{n1}^2 \ x_{n2}^2 \ x_{n1}x_{n2}]^T \in \mathbb{R}^6$
- The target  $t_n$  is regressed from the kernel representation  $\phi(\mathbf{x}_n)$  as





$$\hat{t}_n = \mathbf{w}^T \phi(\mathbf{x}_n) \quad \mathbf{w} \in \mathbb{R}^M$$

- The regression coefficients are given by  $\mathbf{w}_* = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$
- DNNs can be used to learn data-dependent nonlinear transf.  $\phi(\mathbf{x}_n)$
- The last layer of DNNs typically performs linear regression on  $\phi(\mathbf{x}_n)$


## Effect of Model Order $M$ : $t = \sin(\pi x) + \epsilon$




## Effect of Model Order $M$ : $t = \sin(\pi x) + \epsilon$



# Model Validation




# Model Validation



- Training & test error diverge for higher model orders
- Model 'overfits' to the noise in the training data


# Model Validation



	$M = 0$	$M = 1$	$M = 6$	$M = 9$
$w_0^*$	0.19	0.82	0.31	0.35
$w_1^*$		-1.27	7.99	232.37
$w_2^*$			-25.43	-5321.83
$w_3^*$			17.37	48568.31
$w_4^*$				-231639.30
$w_5^*$				640042.26
$w_6^*$				-1061800.52
$w_7^*$				1042400.18
$w_8^*$				-557682.99
$w_9^*$				125201.43

- Training & test error diverge for higher model orders
- Model 'overfits' to the noise in the training data



# Model Validation



	$M = 0$	$M = 1$	$M = 6$	$M = 9$
$w_0^*$	0.19	0.82	0.31	0.35
$w_1^*$		-1.27	7.99	232.37
$w_2^*$			-25.43	-5321.83
$w_3^*$			17.37	48568.31
$w_4^*$				-231639.30
$w_5^*$				640042.26
$w_6^*$				-1061800.52
$w_7^*$				1042400.18
$w_8^*$				-557682.99
$w_9^*$				125201.43



- Training & test error diverge for higher model orders
- Model 'overfits' to the noise in the training data
- Large amplitude weights with alternating polarity.
- $(\Phi^T \Phi)$  may be ill conditioned

## Amount of Training Data ( $M = 9$ )




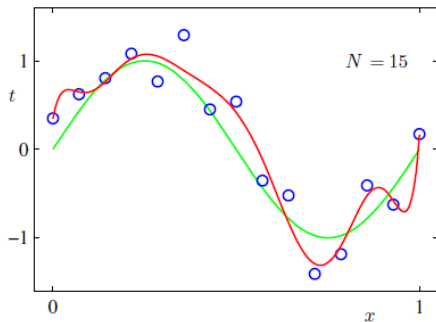


## Amount of Training Data ( $M = 9$ )





- Overfitting is less severe with increased amount of data.

## Amount of Training Data ( $M = 9$ )





- Overfitting is less severe with increased amount of data.
- Model order cannot be limited by the amount of data available!

## Amount of Training Data ( $M = 9$ )



- Overfitting is less severe with increased amount of data.
- Model order cannot be limited by the amount of data available!
- Model order should be based on complexity of task/pattern!

## Amount of Training Data ( $M = 9$ )



- Overfitting is less severe with increased amount of data.
- Model order cannot be limited by the amount of data available!
- Model order should be based on complexity of task/pattern!
- A way forward: arrest the growth of the model weights

# Regularized Least Squares

# Regularized Least Squares

- Add a penalty term to the error term to discourage weight growth

$$J(\mathbf{w}) = \underbrace{E_D(\mathbf{w})}_{\text{Data Term}} + \underbrace{\lambda E_W(\mathbf{w})}_{\text{Regularization Term}}$$

# Regularized Least Squares

- Add a penalty term to the error term to discourage weight growth

$$J(\mathbf{w}) = \underbrace{E_D(\mathbf{w})}_{\text{Data Term}} + \underbrace{\lambda E_W(\mathbf{w})}_{\text{Regularization Term}}$$

- $\lambda$  controls relative importance of the terms (bias vs variance)

# Regularized Least Squares

- Add a penalty term to the error term to discourage weight growth

$$J(\mathbf{w}) = \underbrace{E_D(\mathbf{w})}_{\text{Data Term}} + \underbrace{\lambda E_W(\mathbf{w})}_{\text{Regularization Term}}$$

- $\lambda$  controls relative importance of the terms (bias vs variance)
- Sum of squares error function with a quadratic regularizer

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left( t_n - \mathbf{w}^T \phi(\mathbf{x}_n) \right)^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$



# Regularized Least Squares

- Add a penalty term to the error term to discourage weight growth

$$J(\mathbf{w}) = \underbrace{E_D(\mathbf{w})}_{\text{Data Term}} + \underbrace{\lambda E_W(\mathbf{w})}_{\text{Regularization Term}}$$

- $\lambda$  controls relative importance of the terms (bias vs variance)
- Sum of squares error function with a quadratic regularizer

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left( t_n - \mathbf{w}^T \phi(\mathbf{x}_n) \right)^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$

- Equating  $\nabla_{\mathbf{w}} J(\mathbf{w}) = \mathbf{0} \implies -\Phi^T (\mathbf{t} - \Phi \mathbf{w}) + \lambda \mathbf{w} = \mathbf{0}$

$$\mathbf{w}_* = (\Phi^T \Phi + \lambda \mathbf{I})^{-1} \Phi^T \mathbf{t}$$

# Regularized Least Squares

- Add a penalty term to the error term to discourage weight growth

$$J(\mathbf{w}) = \underbrace{E_D(\mathbf{w})}_{\text{Data Term}} + \underbrace{\lambda E_W(\mathbf{w})}_{\text{Regularization Term}}$$

- $\lambda$  controls relative importance of the terms (bias vs variance)
- Sum of squares error function with a quadratic regularizer

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left( t_n - \mathbf{w}^T \phi(\mathbf{x}_n) \right)^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$

- Equating  $\nabla_{\mathbf{w}} J(\mathbf{w}) = \mathbf{0} \implies -\Phi^T (\mathbf{t} - \Phi \mathbf{w}) + \lambda \mathbf{w} = \mathbf{0}$

$$\mathbf{w}_* = (\Phi^T \Phi + \lambda \mathbf{I})^{-1} \Phi^T \mathbf{t}$$


- Regularization term conditions the autocorrelation matrix!

# Modified Error Surface

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left( t_n - \mathbf{w}^T \phi(\mathbf{x}_n) \right)^2 + \lambda \|\mathbf{w}\|_p$$


# Modified Error Surface

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left( t_n - \mathbf{w}^T \phi(\mathbf{x}_n) \right)^2 + \lambda \|\mathbf{w}\|_p$$




# Modified Error Surface

$$J(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left( t_n - \mathbf{w}^T \phi(\mathbf{x}_n) \right)^2 + \lambda \|\mathbf{w}\|_p$$





$L_2$  Regularizer





$L_1$  Regularizer

# $L_1$ vs $L_2$





# $L_1$ vs $L_2$



- $L_1$  regularization promotes sparser solutions



# $L_1$ vs $L_2$



- $L_1$  regularization promotes sparser solutions
- $L_1$  regularization  $\implies$  Laplacian priors





# $L_1$ vs $L_2$





- $L_1$  regularization promotes sparser solutions
- $L_1$  regularization  $\implies$  Laplacian priors
- $L_2$  regularization  $\implies$  Gaussian priors

## Effect of Regularization ( $N = 10, M = 9$ )





# Effect of Regularization ( $N = 10, M = 9$ )




	$\ln \lambda = -\infty$	$\ln \lambda = -18$	$\ln \lambda = 0$
$w_0^*$	0.35	0.35	0.13
$w_1^*$	232.37	4.74	-0.05
$w_2^*$	-5321.83	-0.77	-0.06
$w_3^*$	48568.31	-31.97	-0.05
$w_4^*$	-231639.30	-3.89	-0.03
$w_5^*$	640042.26	55.28	-0.02
$w_6^*$	-1061800.52	41.32	-0.01
$w_7^*$	1042400.18	-45.95	-0.00
$w_8^*$	-557682.99	-91.53	0.00
$w_9^*$	125201.43	72.68	0.01

# Effect of Regularization ( $N = 10, M = 9$ )



	$\ln \lambda = -\infty$	$\ln \lambda = -18$	$\ln \lambda = 0$
$w_0^*$	0.35	0.35	0.13
$w_1^*$	232.37	4.74	-0.05
$w_2^*$	-5321.83	-0.77	-0.06
$w_3^*$	48568.31	-31.97	-0.05
$w_4^*$	-231639.30	-3.89	-0.03
$w_5^*$	640042.26	55.28	-0.02
$w_6^*$	-1061800.52	41.32	-0.01
$w_7^*$	1042400.18	-45.95	-0.00
$w_8^*$	-557682.99	-91.53	0.00
$w_9^*$	125201.43	72.68	0.01



# Sequential Learning

# Sequential Learning

- LS approach involves considering entire training set in one go.

# Sequential Learning

- LS approach involves considering entire training set in one go.
- For HD data the matrix  $(\Phi^T \Phi)$  may be poorly conditioned

# Sequential Learning

- LS approach involves considering entire training set in one go.
- For HD data the matrix  $(\Phi^T \Phi)$  may be poorly conditioned
- Iteratively update  $\mathbf{w}^{(\tau+1)}$  by adding a correction factor

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} + \Delta \mathbf{w}^{(\tau)}$$



# Sequential Learning

- LS approach involves considering entire training set in one go.
- For HD data the matrix  $(\Phi^T \Phi)$  may be poorly conditioned
- Iteratively update  $\mathbf{w}^{(\tau+1)}$  by adding a correction factor

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} + \Delta \mathbf{w}^{(\tau)}$$

- Apply correction factor in the negative direction of gradient of  $J(\mathbf{w}^{(\tau)})$

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla J(\mathbf{w}^{(\tau)})$$

# Sequential Learning

- LS approach involves considering entire training set in one go.
- For HD data the matrix  $(\Phi^T \Phi)$  may be poorly conditioned
- Iteratively update  $\mathbf{w}^{(\tau+1)}$  by adding a correction factor

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} + \Delta \mathbf{w}^{(\tau)}$$

- Apply correction factor in the negative direction of gradient of  $J(\mathbf{w}^{(\tau)})$

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla J(\mathbf{w}^{(\tau)})$$

- Choose a random batch of points  $\mathcal{B}$  to update  $\mathbf{w}$ .  $J(\mathbf{w}^{(\tau)}) = \frac{1}{2} \sum_{n \in \mathcal{B}} e_n^2$

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} + \eta \sum_{n \in \mathcal{B}} \left( t_n - \mathbf{w}^{(\tau)T} \phi(\mathbf{x}_n) \right) \phi(\mathbf{x}_n)$$

# Sequential Learning

- LS approach involves considering entire training set in one go.
- For HD data the matrix  $(\Phi^T \Phi)$  may be poorly conditioned
- Iteratively update  $\mathbf{w}^{(\tau+1)}$  by adding a correction factor

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} + \Delta \mathbf{w}^{(\tau)}$$

- Apply correction factor in the negative direction of gradient of  $J(\mathbf{w}^{(\tau)})$





$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla J(\mathbf{w}^{(\tau)})$$

- Choose a random batch of points  $\mathcal{B}$  to update  $\mathbf{w}$ .  $J(\mathbf{w}^{(\tau)}) = \frac{1}{2} \sum_{n \in \mathcal{B}} e_n^2$



$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} + \eta \sum_{n \in \mathcal{B}} \left( t_n - \mathbf{w}^{(\tau)T} \phi(\mathbf{x}_n) \right) \phi(\mathbf{x}_n)$$

- $|\mathcal{B}| = N$ : Steepest descent       $|\mathcal{B}| = 1$ : LMS      Otherwise: SGD.



## SGD Error Dynamics: $w_1 = 1.6, w_2 = -0.5$



# Convergence of SGD



# Convergence of SGD




- SGD algorithm converges in mean:





$$\lim_{k \rightarrow \infty} \mathbb{E}[\mathbf{w}_k] \rightarrow (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t} \quad \eta \text{ is small enough}$$

- Expectation over multiple runs ( $k$ ) converges to true solution for convex error surfaces, provided  $\eta$  is sufficiently small

# Geometry of Error Surface vs Convergence Rate







# Geometry of Error Surface vs Convergence Rate


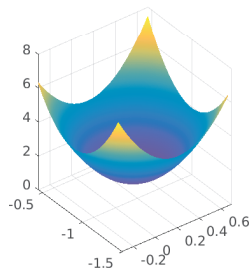






# Geometry of Error Surface vs Convergence Rate





# Geometry of Error Surface vs Convergence Rate





- Gradient magnitude depends on direction!





# Geometry of Error Surface vs Convergence Rate





- Gradient magnitude depends on direction!
- $\eta$  has to be fixed based on steepest direction.



# Geometry of Error Surface vs Convergence Rate



- Gradient magnitude depends on direction!
- $\eta$  has to be fixed based on steepest direction.
- Convergence along flatter dimension is too slow!



# Newton's Method

# Newton's Method

- The weights of the model are updated as

$$\mathbf{w}_{n+1} = \mathbf{w}_n + \Delta \mathbf{w}$$

# Newton's Method

- The weights of the model are updated as

$$\mathbf{w}_{n+1} = \mathbf{w}_n + \Delta \mathbf{w}$$

- Expanding the objective function using Taylor series

$$J(\mathbf{w}_{n+1}) = J(\mathbf{w}_n + \Delta \mathbf{w}) =$$

# Newton's Method

- The weights of the model are updated as

$$\mathbf{w}_{n+1} = \mathbf{w}_n + \Delta \mathbf{w}$$

- Expanding the objective function using Taylor series

$$J(\mathbf{w}_{n+1}) = J(\mathbf{w}_n + \Delta \mathbf{w}) = J(\mathbf{w}_n) + \Delta \mathbf{w}^T \nabla J(\mathbf{w}_n) + \frac{1}{2} \Delta \mathbf{w}^T \nabla^2 J(\mathbf{w}_n) \Delta \mathbf{w}$$



# Newton's Method

- The weights of the model are updated as

$$\mathbf{w}_{n+1} = \mathbf{w}_n + \Delta \mathbf{w}$$

- Expanding the objective function using Taylor series

$$J(\mathbf{w}_{n+1}) = J(\mathbf{w}_n + \Delta \mathbf{w}) = J(\mathbf{w}_n) + \Delta \mathbf{w}^T \nabla J(\mathbf{w}_n) + \frac{1}{2} \Delta \mathbf{w}^T \nabla^2 J(\mathbf{w}_n) \Delta \mathbf{w}$$

- Estimate  $\Delta \mathbf{w}$  s.t  $J(\mathbf{w}_n + \Delta \mathbf{w})$  is minimized

$$\frac{\partial}{\partial \Delta \mathbf{w}} \left( J(\mathbf{w}_n) + \Delta \mathbf{w}^T \nabla J(\mathbf{w}_n) + \frac{1}{2} \Delta \mathbf{w}^T \nabla^2 J(\mathbf{w}_n) \Delta \mathbf{w} \right) = 0$$

# Newton's Method

- The weights of the model are updated as

$$\mathbf{w}_{n+1} = \mathbf{w}_n + \Delta \mathbf{w}$$

- Expanding the objective function using Taylor series

$$J(\mathbf{w}_{n+1}) = J(\mathbf{w}_n + \Delta \mathbf{w}) = J(\mathbf{w}_n) + \Delta \mathbf{w}^T \nabla J(\mathbf{w}_n) + \frac{1}{2} \Delta \mathbf{w}^T \nabla^2 J(\mathbf{w}_n) \Delta \mathbf{w}$$

- Estimate  $\Delta \mathbf{w}$  s.t  $J(\mathbf{w}_n + \Delta \mathbf{w})$  is minimized

$$\frac{\partial}{\partial \Delta \mathbf{w}} \left( J(\mathbf{w}_n) + \Delta \mathbf{w}^T \nabla J(\mathbf{w}_n) + \frac{1}{2} \Delta \mathbf{w}^T \nabla^2 J(\mathbf{w}_n) \Delta \mathbf{w} \right) = 0$$

- Optimal update is given by  $\Delta \mathbf{w} = -\frac{\nabla J(\mathbf{w}_n)}{\nabla^2 J(\mathbf{w}_n)}$

$$\mathbf{w}_{n+1} = \mathbf{w}_n - \mathbf{H}^{-1}(\mathbf{w}_n) \nabla J(\mathbf{w}_n) \quad \mathbf{H}(\mathbf{w}_n) = \nabla^2 J(\mathbf{w}_n)$$

# Homework - 1

- Apply Newtons method to steepest-descent algorithm to the optimal step size  $\eta$ , and check how many iterations are required for convergence.

$$\mathbf{w}^{new} = \mathbf{w}^{old} + \eta \mathbf{X}^T(\mathbf{t} - \mathbf{X}\mathbf{w}) \Big|_{\mathbf{w}=\mathbf{w}^{old}}$$

## Homework - 2

- Suppose you are experimenting with  $L_1$  and  $L_2$  regularization. Further, imagine that you are running gradient descent and at some iteration your weight vector is  $w = [1, \epsilon] \in \mathbb{R}^2$  where  $\epsilon > 0$  is very small. With the help of this example explain why  $L_2$  norm does not encourage sparsity i.e., it will not try to drive  $\epsilon$  to 0 to produce a sparse weight vector. Give mathematical explanation.

## Homework - 3

- Till now we have been considering a scalar target  $t$  from a vector of input observations  $\mathbf{x}$ . How do you extend this approach for regressing a vector of targets  $\mathbf{t} = (t_1, t_2, \dots, t_P)$ . Derive the closed form solutions and write sequential update equations using SGD.