

# Linear Models for Classification

K Sri Rama Murty

IIT Hyderabad

`ksrm@ee.iith.ac.in`

March 22, 2022

# Classification

# Classification

- Given an observation vector  $\mathbf{x}$ , assign it to one of the  $K$  classes

# Classification

- Given an observation vector  $\mathbf{x}$ , assign it to one of the  $K$  classes
- The target can take one of the  $K$  discrete labels  $\mathcal{C}_k$ ,  $k = 1, 2, \dots, K$

# Classification

- Given an observation vector  $\mathbf{x}$ , assign it to one of the  $K$  classes
- The target can take one of the  $K$  discrete labels  $\mathcal{C}_k$ ,  $k = 1, 2, \dots, K$ 
  - Classify  $\mathbf{x} = (45\text{Kg}, 4.8\text{ft})$  as adult or kid

# Classification

- Given an observation vector  $\mathbf{x}$ , assign it to one of the  $K$  classes
- The target can take one of the  $K$  discrete labels  $\mathcal{C}_k$ ,  $k = 1, 2, \dots, K$ 
  - Classify  $\mathbf{x} = (45\text{Kg}, 4.8\text{ft})$  as adult or kid
  - Classify a given image as face or nonface

# Classification

- Given an observation vector  $\mathbf{x}$ , assign it to one of the  $K$  classes
- The target can take one of the  $K$  discrete labels  $\mathcal{C}_k$ ,  $k = 1, 2, \dots, K$ 
  - Classify  $\mathbf{x} = (45\text{Kg}, 4.8\text{ft})$  as adult or kid
  - Classify a given image as face or nonface
  - Classify a speech waveform as one of the 40 phonemes

# Classification

- Given an observation vector  $\mathbf{x}$ , assign it to one of the  $K$  classes
- The target can take one of the  $K$  discrete labels  $\mathcal{C}_k$ ,  $k = 1, 2, \dots, K$ 
  - Classify  $\mathbf{x} = (45\text{Kg}, 4.8\text{ft})$  as adult or kid
  - Classify a given image as face or nonface
  - Classify a speech waveform as one of the 40 phonemes
- Discriminant function to create separating hyperplane between classes  
Eg. Least squares, Fisher discriminant, perceptron, SVM



# Classification

- Given an observation vector  $\mathbf{x}$ , assign it to one of the  $K$  classes
- The target can take one of the  $K$  discrete labels  $\mathcal{C}_k$ ,  $k = 1, 2, \dots, K$ 
  - Classify  $\mathbf{x} = (45\text{Kg}, 4.8\text{ft})$  as adult or kid
  - Classify a given image as face or nonface
  - Classify a speech waveform as one of the 40 phonemes
- Discriminant function to create separating hyperplane between classes  
Eg. Least squares, Fisher discriminant, perceptron, SVM
- Probabilistic approaches to estimate posterior probabilities  $P[\mathcal{C}_k/\mathbf{x}]$

# Classification

- Given an observation vector  $\mathbf{x}$ , assign it to one of the  $K$  classes
- The target can take one of the  $K$  discrete labels  $\mathcal{C}_k$ ,  $k = 1, 2, \dots, K$ 
  - Classify  $\mathbf{x} = (45\text{Kg}, 4.8\text{ft})$  as adult or kid
  - Classify a given image as face or nonface
  - Classify a speech waveform as one of the 40 phonemes
- Discriminant function to create separating hyperplane between classes  
Eg. Least squares, Fisher discriminant, perceptron, SVM
- Probabilistic approaches to estimate posterior probabilities  $P[\mathcal{C}_k/\mathbf{x}]$ 
  - Discriminative models: Directly estimate  $P[\mathcal{C}_k/\mathbf{x}] = f(\mathbf{x}, \mathbf{w})$   
Eg. Logistic regression, DNN classifiers

# Classification

- Given an observation vector  $\mathbf{x}$ , assign it to one of the  $K$  classes
- The target can take one of the  $K$  discrete labels  $\mathcal{C}_k$ ,  $k = 1, 2, \dots, K$ 
  - Classify  $\mathbf{x} = (45\text{Kg}, 4.8\text{ft})$  as adult or kid
  - Classify a given image as face or nonface
  - Classify a speech waveform as one of the 40 phonemes
- Discriminant function to create separating hyperplane between classes  
Eg. Least squares, Fisher discriminant, perceptron, SVM
- Probabilistic approaches to estimate posterior probabilities  $P[\mathcal{C}_k/\mathbf{x}]$ 
  - Discriminative models: Directly estimate  $P[\mathcal{C}_k/\mathbf{x}] = f(\mathbf{x}, \mathbf{w})$   
Eg. Logistic regression, DNN classifiers
  - Generative models: Arrive at the posterior from joint density  $p(\mathbf{x}, \mathcal{C}_k)$   
Eg. PDF Estimation: GMM (RV), HMM (RP)

# Discriminant Functions (Binary Classification)

# Discriminant Functions (Binary Classification)

- For an input vector  $\mathbf{x} \in \mathbb{R}^D$ , Linear Discriminant Function is given by

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

# Discriminant Functions (Binary Classification)

- For an input vector  $\mathbf{x} \in \mathbb{R}^D$ , Linear Discriminant Function is given by

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

- Decision Rule: Assign  $\mathbf{x}$  to  $\mathcal{C}_1$  if  $y(\mathbf{x}) \geq 0$ , otherwise assign  $\mathbf{x}$  to  $\mathcal{C}_2$

# Discriminant Functions (Binary Classification)

- For an input vector  $\mathbf{x} \in \mathbb{R}^D$ , Linear Discriminant Function is given by

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

- Decision Rule: Assign  $\mathbf{x}$  to  $\mathcal{C}_1$  if  $y(\mathbf{x}) \geq 0$ , otherwise assign  $\mathbf{x}$  to  $\mathcal{C}_2$
- Decision boundary,  $y(\mathbf{x}) = 0$ , corresponds to a  $D - 1$  dim. hyperplane

# Discriminant Functions (Binary Classification)

- For an input vector  $\mathbf{x} \in \mathbb{R}^D$ , Linear Discriminant Function is given by

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

- Decision Rule: Assign  $\mathbf{x}$  to  $\mathcal{C}_1$  if  $y(\mathbf{x}) \geq 0$ , otherwise assign  $\mathbf{x}$  to  $\mathcal{C}_2$
- Decision boundary,  $y(\mathbf{x}) = 0$ , corresponds to a  $D - 1$  dim. hyperplane
- Weight vector  $\mathbf{w}$  is orthogonal to every vector on the decision surface



# Discriminant Functions (Binary Classification)

- For an input vector  $\mathbf{x} \in \mathbb{R}^D$ , Linear Discriminant Function is given by

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

- Decision Rule: Assign  $\mathbf{x}$  to  $\mathcal{C}_1$  if  $y(\mathbf{x}) \geq 0$ , otherwise assign  $\mathbf{x}$  to  $\mathcal{C}_2$
- Decision boundary,  $y(\mathbf{x}) = 0$ , corresponds to a  $D - 1$  dim. hyperplane
- Weight vector  $\mathbf{w}$  is orthogonal to every vector on the decision surface
- Weight vector  $\mathbf{w}$  determines the orientation of the decision surface

# Discriminant Functions (Binary Classification)

- For an input vector  $\mathbf{x} \in \mathbb{R}^D$ , Linear Discriminant Function is given by

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

- Decision Rule: Assign  $\mathbf{x}$  to  $\mathcal{C}_1$  if  $y(\mathbf{x}) \geq 0$ , otherwise assign  $\mathbf{x}$  to  $\mathcal{C}_2$
- Decision boundary,  $y(\mathbf{x}) = 0$ , corresponds to a  $D - 1$  dim. hyperplane
- Weight vector  $\mathbf{w}$  is orthogonal to every vector on the decision surface
- Weight vector  $\mathbf{w}$  determines the orientation of the decision surface
- Bias parameter  $w_0$  determines the location of the decision surface

# Discriminant Functions (Binary Classification)

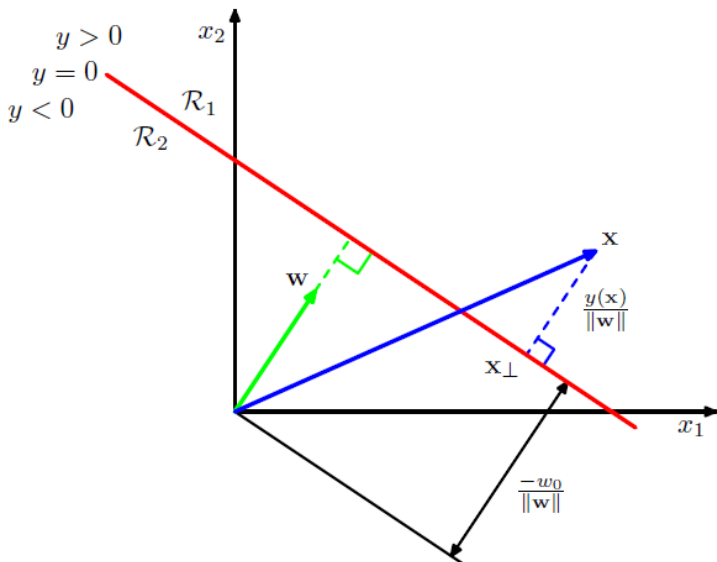
- For an input vector  $\mathbf{x} \in \mathbb{R}^D$ , Linear Discriminant Function is given by

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

- Decision Rule: Assign  $\mathbf{x}$  to  $\mathcal{C}_1$  if  $y(\mathbf{x}) \geq 0$ , otherwise assign  $\mathbf{x}$  to  $\mathcal{C}_2$
- Decision boundary,  $y(\mathbf{x}) = 0$ , corresponds to a  $D - 1$  dim. hyperplane
- Weight vector  $\mathbf{w}$  is orthogonal to every vector on the decision surface
- Weight vector  $\mathbf{w}$  determines the orientation of the decision surface
- Bias parameter  $w_0$  determines the location of the decision surface
- Normal distance from origin to the decision surface is given by

$$\frac{\mathbf{w}^T \mathbf{x}}{\|\mathbf{w}\|} = -\frac{w_0}{\|\mathbf{w}\|}$$

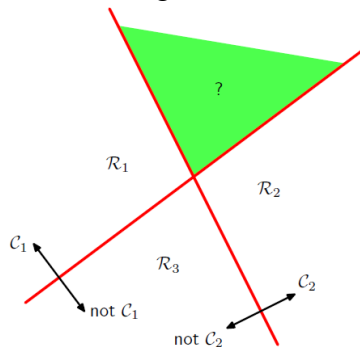
# Geometry of Decision Boundary in 2D



# Extending to Multiple Classes

# Extending to Multiple Classes

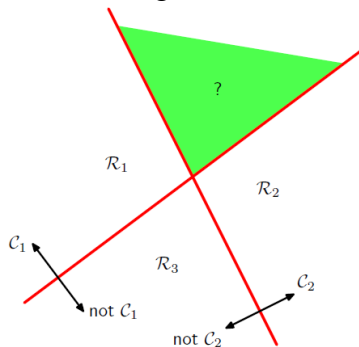
One Against Rest



K-1 Decision Surfaces

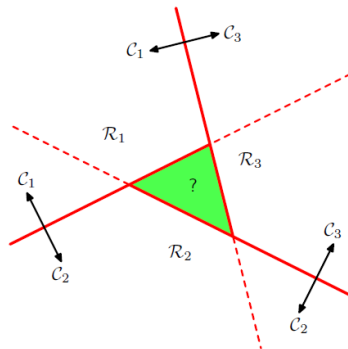
# Extending to Multiple Classes

One Against Rest



$K-1$  Decision Surfaces

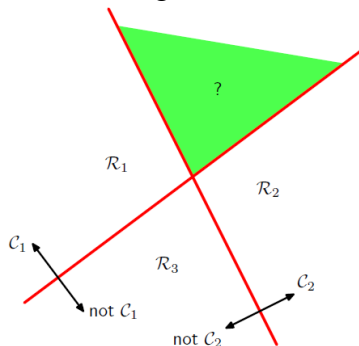
One Against One



$K(K-1)/2$  Decision Surfaces

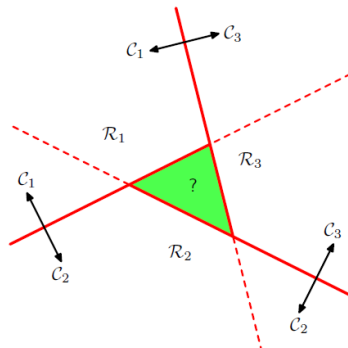
# Extending to Multiple Classes

One Against Rest



$K-1$  Decision Surfaces

One Against One

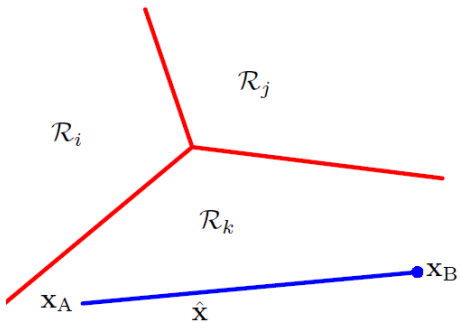


$K(K-1)/2$  Decision Surfaces

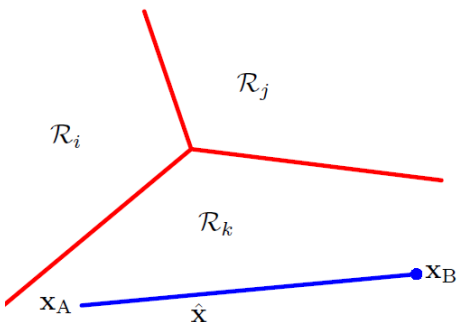
- Both the approaches lead to ambiguous regions!
- Decision regions should be singly connected and convex.



# Multiclass Linear Discriminant



# Multiclass Linear Discriminant



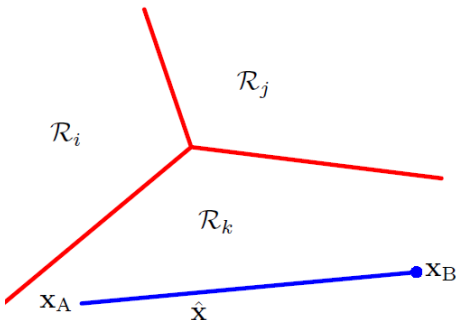
- K-Linear discriminant functions

$$y_k(\mathbf{x}) = \mathbf{w}_k^T \mathbf{x} + w_{k0} \quad k = 1, 2, \dots, K$$

- Decision rule: Assign  $\mathbf{x}$  to  $\mathcal{C}_k$  if

$$y_k(\mathbf{x}) > y_j(\mathbf{x}), \quad \forall j \neq k$$

# Multiclass Linear Discriminant



- Decision boundary  $\mathcal{C}_k$  &  $\mathcal{C}_j$ :

$$y_k(\mathbf{x}) = y_j(\mathbf{x})$$

$$(\mathbf{w}_k - \mathbf{w}_j)^T \mathbf{x} + (w_{k0} - w_{j0}) = 0$$

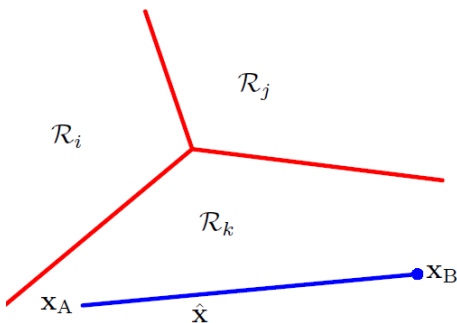
- K-Linear discriminant functions

$$y_k(\mathbf{x}) = \mathbf{w}_k^T \mathbf{x} + w_{k0} \quad k = 1, 2, \dots, K$$

- Decision rule: Assign  $\mathbf{x}$  to  $\mathcal{C}_k$  if

$$y_k(\mathbf{x}) > y_j(\mathbf{x}), \quad \forall j \neq k$$

# Multiclass Linear Discriminant



- Decision boundary  $\mathcal{C}_k$  &  $\mathcal{C}_j$ :

$$y_k(\mathbf{x}) = y_j(\mathbf{x})$$

$$(\mathbf{w}_k - \mathbf{w}_j)^T \mathbf{x} + (w_{k0} - w_{j0}) = 0$$

- Decision regions are convex

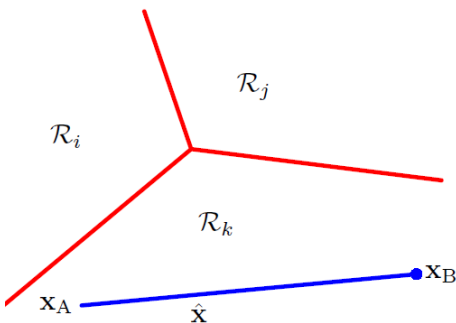
- K-Linear discriminant functions

$$y_k(\mathbf{x}) = \mathbf{w}_k^T \mathbf{x} + w_{k0} \quad k = 1, 2, \dots, K$$

- Decision rule: Assign  $\mathbf{x}$  to  $\mathcal{C}_k$  if

$$y_k(\mathbf{x}) > y_j(\mathbf{x}), \quad \forall j \neq k$$

# Multiclass Linear Discriminant



- K-Linear discriminant functions

$$y_k(\mathbf{x}) = \mathbf{w}_k^T \mathbf{x} + w_{k0} \quad k = 1, 2, \dots, K$$

- Decision rule: Assign  $\mathbf{x}$  to  $\mathcal{C}_k$  if

$$y_k(\mathbf{x}) > y_j(\mathbf{x}), \quad \forall j \neq k$$

- Decision boundary  $\mathcal{C}_k$  &  $\mathcal{C}_j$ :

$$y_k(\mathbf{x}) = y_j(\mathbf{x})$$

$$(\mathbf{w}_k - \mathbf{w}_j)^T \mathbf{x} + (w_{k0} - w_{j0}) = 0$$

- Decision regions are convex
- For two classes, we can either employ a single discriminant or two discriminants

# Least Squares for Classification

# Least Squares for Classification

- Let us apply Least Squares Regression approach to classification

# Least Squares for Classification

- Let us apply Least Squares Regression approach to classification
- Training data:  $\{(\mathbf{x}_1, \mathbf{t}_1), (\mathbf{x}_2, \mathbf{t}_2), \dots, (\mathbf{x}_N, \mathbf{t}_N)\}$



# Least Squares for Classification

- Let us apply Least Squares Regression approach to classification
- Training data:  $\{(\mathbf{x}_1, \mathbf{t}_1), (\mathbf{x}_2, \mathbf{t}_2), \dots, (\mathbf{x}_N, \mathbf{t}_N)\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  denotes input observation vectors

# Least Squares for Classification

- Let us apply Least Squares Regression approach to classification
- Training data:  $\{(\mathbf{x}_1, \mathbf{t}_1), (\mathbf{x}_2, \mathbf{t}_2), \dots, (\mathbf{x}_N, \mathbf{t}_N)\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  denotes input observation vectors
  - $\mathbf{t}_n$ , with 1-of-K binary coding, denotes the class label ( $K$  - Classes)

# Least Squares for Classification

- Let us apply Least Squares Regression approach to classification
- Training data:  $\{(\mathbf{x}_1, \mathbf{t}_1), (\mathbf{x}_2, \mathbf{t}_2), \dots, (\mathbf{x}_N, \mathbf{t}_N)\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  denotes input observation vectors
  - $\mathbf{t}_n$ , with 1-of-K binary coding, denotes the class label ( $K$  - Classes)
- Each class  $\mathcal{C}_k$  is described by its own linear discriminant

$$y_k(\mathbf{x}) = \mathbf{w}_k^T \mathbf{x} + w_{k0} \quad k = 1, 2, \dots, K$$

# Least Squares for Classification

- Let us apply Least Squares Regression approach to classification
- Training data:  $\{(\mathbf{x}_1, \mathbf{t}_1), (\mathbf{x}_2, \mathbf{t}_2), \dots, (\mathbf{x}_N, \mathbf{t}_N)\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  denotes input observation vectors
  - $\mathbf{t}_n$ , with 1-of-K binary coding, denotes the class label ( $K$  - Classes)
- Each class  $\mathcal{C}_k$  is described by its own linear discriminant

$$y_k(\mathbf{x}) = \mathbf{w}_k^T \mathbf{x} + w_{k0} \quad k = 1, 2, \dots, K$$

- Regressed target for  $\mathbf{x}_n$ :  $\hat{\mathbf{t}}_n = \mathbf{y}(\mathbf{x}_n) = \mathbf{W}^T \mathbf{x}_n$  (Augmented)

# Least Squares for Classification

- Let us apply Least Squares Regression approach to classification
- Training data:  $\{(\mathbf{x}_1, \mathbf{t}_1), (\mathbf{x}_2, \mathbf{t}_2), \dots, (\mathbf{x}_N, \mathbf{t}_N)\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  denotes input observation vectors
  - $\mathbf{t}_n$ , with 1-of-K binary coding, denotes the class label ( $K$  - Classes)
- Each class  $\mathcal{C}_k$  is described by its own linear discriminant

$$y_k(\mathbf{x}) = \mathbf{w}_k^T \mathbf{x} + w_{k0} \quad k = 1, 2, \dots, K$$

- Regressed target for  $\mathbf{x}_n$ :  $\hat{\mathbf{t}}_n = \mathbf{y}(\mathbf{x}_n) = \mathbf{W}^T \mathbf{x}_n$  (Augmented)
- Estimate  $\mathbf{W}$  to minimize sum of squares error

$$J(\mathbf{W}) = \frac{1}{2} \text{Tr} \left\{ (\mathbf{XW} - \mathbf{T})^T (\mathbf{XW} - \mathbf{T}) \right\}$$

# Least Squares for Classification

- Let us apply Least Squares Regression approach to classification
- Training data:  $\{(\mathbf{x}_1, \mathbf{t}_1), (\mathbf{x}_2, \mathbf{t}_2), \dots, (\mathbf{x}_N, \mathbf{t}_N)\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  denotes input observation vectors
  - $\mathbf{t}_n$ , with 1-of-K binary coding, denotes the class label ( $K$  - Classes)
- Each class  $\mathcal{C}_k$  is described by its own linear discriminant

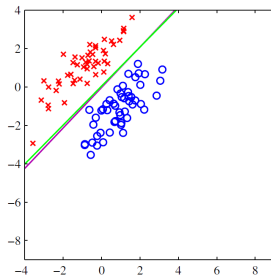
$$y_k(\mathbf{x}) = \mathbf{w}_k^\top \mathbf{x} + w_{k0} \quad k = 1, 2, \dots, K$$

- Regressed target for  $\mathbf{x}_n$ :  $\hat{\mathbf{t}}_n = \mathbf{y}(\mathbf{x}_n) = \mathbf{W}^\top \mathbf{x}_n$  (Augmented)
- Estimate  $\mathbf{W}$  to minimize sum of squares error

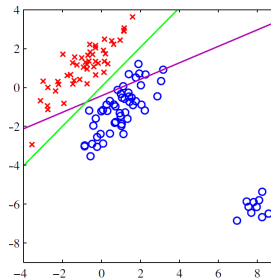
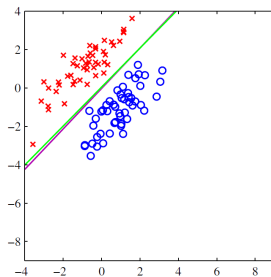
$$J(\mathbf{W}) = \frac{1}{2} \text{Tr} \left\{ (\mathbf{XW} - \mathbf{T})^\top (\mathbf{XW} - \mathbf{T}) \right\}$$

- Setting  $\nabla_{\mathbf{W}} J(\mathbf{W}) = \mathbf{0}$ , we get  $\mathbf{W}_* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{T}$

# Issues with LS

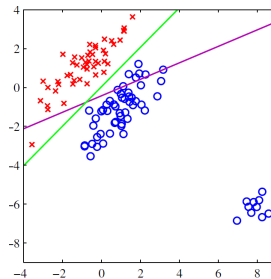
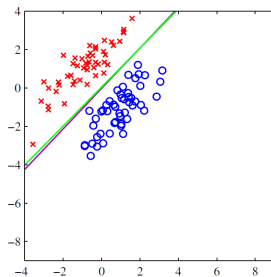


# Issues with LS



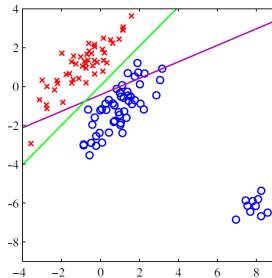
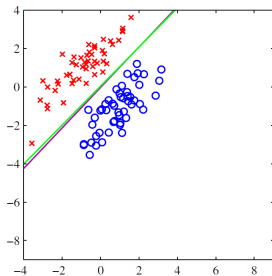


# Issues with LS



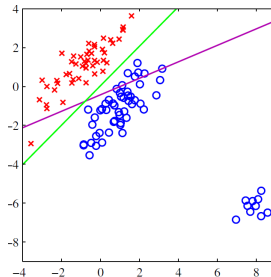
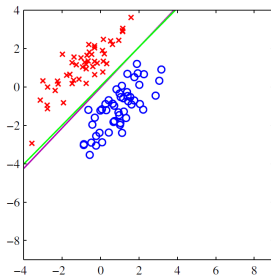
- Least squares solutions lack robustness to outliers

# Issues with LS



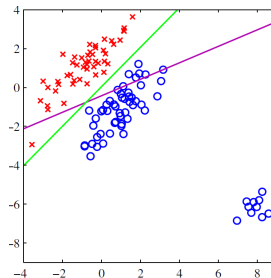
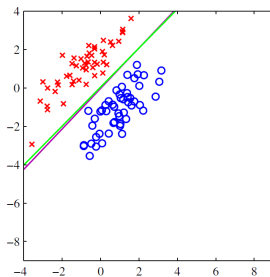
- Least squares solutions lack robustness to outliers
- Additional data-points resulted in significant change in boundary

# Issues with LS



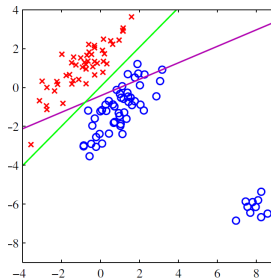
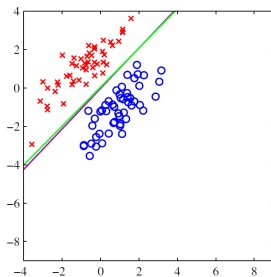
- Least squares solutions lack robustness to outliers
- Additional data-points resulted in significant change in boundary
- Sum of squares error penalizes predictions that are "too correct"

# Issues with LS



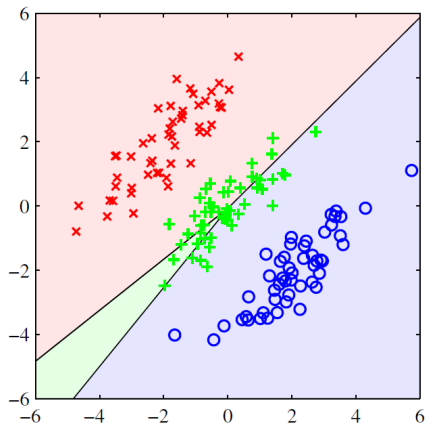
- Least squares solutions lack robustness to outliers
- Additional data-points resulted in significant change in boundary
- Sum of squares error penalizes predictions that are "too correct"
- Attempts to achieve "many-to-one" mapping through linearity!

# Issues with LS

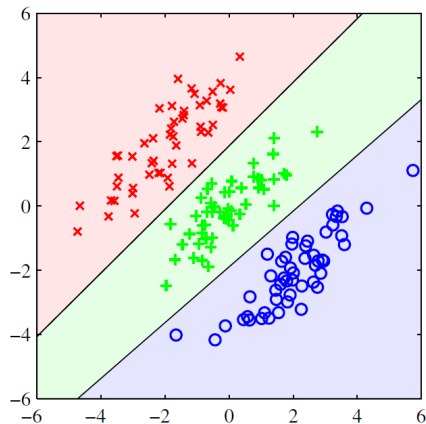


- Least squares solutions lack robustness to outliers
- Additional data-points resulted in significant change in boundary
- Sum of squares error penalizes predictions that are "too correct"
- Attempts to achieve "many-to-one" mapping through linearity!
- LS approach failed even for linearly separable classes

# Issues with LS



Least Squares Classifier



Logistic Regression

# Homework

# Homework

- **Property of LS:** If every target in the training set satisfies some linear constraint

$$\mathbf{a}^T \mathbf{t}_n + b = 0, \quad \forall n$$

for some arbitrary constants  $\mathbf{a}$  and  $b$ , then the model prediction for any value of  $\mathbf{x}$  satisfies the same constraint.

$$\mathbf{a}^T \mathbf{y}(\mathbf{x}) + b = 0$$



# Homework

- **Property of LS:** If every target in the training set satisfies some linear constraint

$$\mathbf{a}^T \mathbf{t}_n + b = 0, \quad \forall n$$

for some arbitrary constants  $\mathbf{a}$  and  $b$ , then the model prediction for any value of  $\mathbf{x}$  satisfies the same constraint.

$$\mathbf{a}^T \mathbf{y}(\mathbf{x}) + b = 0$$

- If we use 1-of-K coding for targets, then the predictions sum to 1.

$$\sum_{k=1}^K y_k(\mathbf{x}) = 1$$

# Homework

- **Property of LS:** If every target in the training set satisfies some linear constraint

$$\mathbf{a}^T \mathbf{t}_n + b = 0, \quad \forall n$$

for some arbitrary constants  $\mathbf{a}$  and  $b$ , then the model prediction for any value of  $\mathbf{x}$  satisfies the same constraint.

$$\mathbf{a}^T \mathbf{y}(\mathbf{x}) + b = 0$$

- If we use 1-of-K coding for targets, then the predictions sum to 1.

$$\sum_{k=1}^K y_k(\mathbf{x}) = 1$$

- However,  $y_k(\mathbf{x})$  cannot be interpreted as posterior probability. They can be negative!

# Linear Discriminant Analysis

# Linear Discriminant Analysis

- Dimensionality reduction interpretation to LS classifier (2 classes)

# Linear Discriminant Analysis

- Dimensionality reduction interpretation to LS classifier (2 classes)
  - Project the  $\mathbf{x} \in \mathbb{R}^D$  on to real line (1-D):  $y(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$

# Linear Discriminant Analysis

- Dimensionality reduction interpretation to LS classifier (2 classes)
  - Project the  $\mathbf{x} \in \mathbb{R}^D$  on to real line (1-D):  $y(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$
  - Assign  $\mathbf{x}$  to  $\mathcal{C}_1$  if  $y(\mathbf{x}) > -w_0$  and to  $\mathcal{C}_2$  otherwise

# Linear Discriminant Analysis

- Dimensionality reduction interpretation to LS classifier (2 classes)
  - Project the  $\mathbf{x} \in \mathbb{R}^D$  on to real line (1-D):  $y(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$
  - Assign  $\mathbf{x}$  to  $\mathcal{C}_1$  if  $y(\mathbf{x}) > -w_0$  and to  $\mathcal{C}_2$  otherwise
  - Projecting to 1-D, in general, leads to loss of information

# Linear Discriminant Analysis

- Dimensionality reduction interpretation to LS classifier (2 classes)
  - Project the  $\mathbf{x} \in \mathbb{R}^D$  on to real line (1-D):  $y(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$
  - Assign  $\mathbf{x}$  to  $\mathcal{C}_1$  if  $y(\mathbf{x}) > -w_0$  and to  $\mathcal{C}_2$  otherwise
  - Projecting to 1-D, in general, leads to loss of information
  - Adjust  $\mathbf{w}$  to select a projection that maximizes the class separation.



# Linear Discriminant Analysis

- Dimensionality reduction interpretation to LS classifier (2 classes)
  - Project the  $\mathbf{x} \in \mathbb{R}^D$  on to real line (1-D):  $y(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$
  - Assign  $\mathbf{x}$  to  $\mathcal{C}_1$  if  $y(\mathbf{x}) > -w_0$  and to  $\mathcal{C}_2$  otherwise
  - Projecting to 1-D, in general, leads to loss of information
  - Adjust  $\mathbf{w}$  to select a projection that maximizes the class separation.
- Consider  $N_1$  points from  $\mathcal{C}_1$  and  $N_2$  points from  $\mathcal{C}_2$  in  $D$ -dim space

# Linear Discriminant Analysis

- Dimensionality reduction interpretation to LS classifier (2 classes)
  - Project the  $\mathbf{x} \in \mathbb{R}^D$  on to real line (1-D):  $y(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$
  - Assign  $\mathbf{x}$  to  $\mathcal{C}_1$  if  $y(\mathbf{x}) > -w_0$  and to  $\mathcal{C}_2$  otherwise
  - Projecting to 1-D, in general, leads to loss of information
  - Adjust  $\mathbf{w}$  to select a projection that maximizes the class separation.
- Consider  $N_1$  points from  $\mathcal{C}_1$  and  $N_2$  points from  $\mathcal{C}_2$  in  $D$ -dim space
  - Mean vectors in original space:  $\mathbf{m}_1 = \frac{1}{N_1} \sum_{n \in \mathcal{C}_1} \mathbf{x}_n$     $\mathbf{m}_2 = \frac{1}{N_2} \sum_{n \in \mathcal{C}_2} \mathbf{x}_n$

# Linear Discriminant Analysis

- Dimensionality reduction interpretation to LS classifier (2 classes)
  - Project the  $\mathbf{x} \in \mathbb{R}^D$  on to real line (1-D):  $y(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$
  - Assign  $\mathbf{x}$  to  $\mathcal{C}_1$  if  $y(\mathbf{x}) > -w_0$  and to  $\mathcal{C}_2$  otherwise
  - Projecting to 1-D, in general, leads to loss of information
  - Adjust  $\mathbf{w}$  to select a projection that maximizes the class separation.
- Consider  $N_1$  points from  $\mathcal{C}_1$  and  $N_2$  points from  $\mathcal{C}_2$  in  $D$ -dim space
  - Mean vectors in original space:  $\mathbf{m}_1 = \frac{1}{N_1} \sum_{n \in \mathcal{C}_1} \mathbf{x}_n$     $\mathbf{m}_2 = \frac{1}{N_2} \sum_{n \in \mathcal{C}_2} \mathbf{x}_n$
  - Let the projected means be  $\mu_1 = \mathbf{w}^T \mathbf{m}_1$  and  $\mu_2 = \mathbf{w}^T \mathbf{m}_2$

# Linear Discriminant Analysis

- Dimensionality reduction interpretation to LS classifier (2 classes)
  - Project the  $\mathbf{x} \in \mathbb{R}^D$  on to real line (1-D):  $y(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$
  - Assign  $\mathbf{x}$  to  $\mathcal{C}_1$  if  $y(\mathbf{x}) > -w_0$  and to  $\mathcal{C}_2$  otherwise
  - Projecting to 1-D, in general, leads to loss of information
  - Adjust  $\mathbf{w}$  to select a projection that maximizes the class separation.
- Consider  $N_1$  points from  $\mathcal{C}_1$  and  $N_2$  points from  $\mathcal{C}_2$  in  $D$ -dim space
  - Mean vectors in original space:  $\mathbf{m}_1 = \frac{1}{N_1} \sum_{n \in \mathcal{C}_1} \mathbf{x}_n$     $\mathbf{m}_2 = \frac{1}{N_2} \sum_{n \in \mathcal{C}_2} \mathbf{x}_n$
  - Let the projected means be  $\mu_1 = \mathbf{w}^T \mathbf{m}_1$  and  $\mu_2 = \mathbf{w}^T \mathbf{m}_2$
  - Choose  $\mathbf{w}$  to maximize the separation between projected means

$$\mu_2 - \mu_1 = \mathbf{w}^T (\mathbf{m}_2 - \mathbf{m}_1) \quad \text{st.} \quad \mathbf{w}^T \mathbf{w} = 1$$

# Linear Discriminant Analysis

- Dimensionality reduction interpretation to LS classifier (2 classes)
  - Project the  $\mathbf{x} \in \mathbb{R}^D$  on to real line (1-D):  $y(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$
  - Assign  $\mathbf{x}$  to  $\mathcal{C}_1$  if  $y(\mathbf{x}) > -w_0$  and to  $\mathcal{C}_2$  otherwise
  - Projecting to 1-D, in general, leads to loss of information
  - Adjust  $\mathbf{w}$  to select a projection that maximizes the class separation.
- Consider  $N_1$  points from  $\mathcal{C}_1$  and  $N_2$  points from  $\mathcal{C}_2$  in  $D$ -dim space
  - Mean vectors in original space:  $\mathbf{m}_1 = \frac{1}{N_1} \sum_{n \in \mathcal{C}_1} \mathbf{x}_n$   $\mathbf{m}_2 = \frac{1}{N_2} \sum_{n \in \mathcal{C}_2} \mathbf{x}_n$
  - Let the projected means be  $\mu_1 = \mathbf{w}^T \mathbf{m}_1$  and  $\mu_2 = \mathbf{w}^T \mathbf{m}_2$
  - Choose  $\mathbf{w}$  to maximize the separation between projected means

$$\mu_2 - \mu_1 = \mathbf{w}^T (\mathbf{m}_2 - \mathbf{m}_1) \quad \text{st.} \quad \mathbf{w}^T \mathbf{w} = 1$$

- The direction for  $\mathbf{w}$  can be shown to be  $\mathbf{w} \propto (\mathbf{m}_2 - \mathbf{m}_1)$

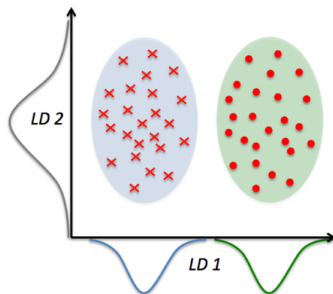
# Linear Discriminant Analysis

- Dimensionality reduction interpretation to LS classifier (2 classes)
  - Project the  $\mathbf{x} \in \mathbb{R}^D$  on to real line (1-D):  $y(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$
  - Assign  $\mathbf{x}$  to  $\mathcal{C}_1$  if  $y(\mathbf{x}) > -w_0$  and to  $\mathcal{C}_2$  otherwise
  - Projecting to 1-D, in general, leads to loss of information
  - Adjust  $\mathbf{w}$  to select a projection that maximizes the class separation.
- Consider  $N_1$  points from  $\mathcal{C}_1$  and  $N_2$  points from  $\mathcal{C}_2$  in  $D$ -dim space
  - Mean vectors in original space:  $\mathbf{m}_1 = \frac{1}{N_1} \sum_{n \in \mathcal{C}_1} \mathbf{x}_n$   $\mathbf{m}_2 = \frac{1}{N_2} \sum_{n \in \mathcal{C}_2} \mathbf{x}_n$
  - Let the projected means be  $\mu_1 = \mathbf{w}^T \mathbf{m}_1$  and  $\mu_2 = \mathbf{w}^T \mathbf{m}_2$
  - Choose  $\mathbf{w}$  to maximize the separation between projected means

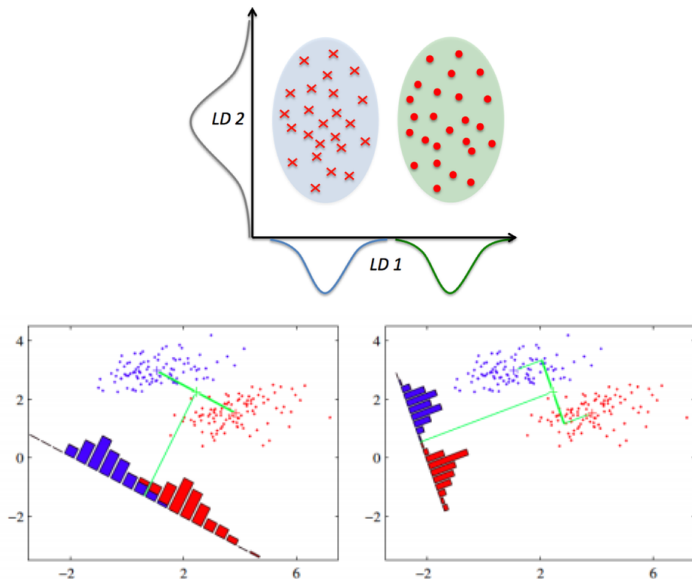
$$\mu_2 - \mu_1 = \mathbf{w}^T (\mathbf{m}_2 - \mathbf{m}_1) \quad \text{st.} \quad \mathbf{w}^T \mathbf{w} = 1$$

- The direction for  $\mathbf{w}$  can be shown to be  $\mathbf{w} \propto (\mathbf{m}_2 - \mathbf{m}_1)$
- This approach is not optimal for nondiagonal covariances

# LDA - Illustration



# LDA - Illustration





# Fisher Discriminant Analysis

- For nondiagonal covariances, spread of data should also be considered
- Project the data in a direction that
  - maximizes separation between the means of the projected classes
  - minimizes the variance within each projected class
- The objective function is given by

$$J(\mathbf{w}) = \frac{(\mu_2 - \mu_1)^2}{\sigma_1^2 + \sigma_2^2}$$

# Fisher Discriminant Analysis

- For nondiagonal covariances, spread of data should also be considered
- Project the data in a direction that
  - maximizes separation between the means of the projected classes
  - minimizes the variance within each projected class
- The objective function is given by

$$J(\mathbf{w}) = \frac{(\mu_2 - \mu_1)^2}{\sigma_1^2 + \sigma_2^2} = \frac{\mathbf{w}^T \mathbf{S}_B \mathbf{w}}{\mathbf{w}^T \mathbf{S}_W \mathbf{w}}$$

- $\mathbf{S}_B$  is *between-class* covariance:  $\mathbf{S}_B = (\mathbf{m}_2 - \mathbf{m}_1)(\mathbf{m}_2 - \mathbf{m}_1)^T$
- $\mathbf{S}_W$  is *within-class* covariance:  $\mathbf{S}_W = \sum_{k=1}^2 \sum_{n \in \mathcal{C}_k} (\mathbf{x}_n - \mathbf{m}_k)(\mathbf{x}_n - \mathbf{m}_k)^T$
- Optimal direction of  $\mathbf{w}$  can be obtained by maximizing  $J(\mathbf{w})$

# Solution to Fisher Criterion

# Solution to Fisher Criterion

- Equating  $\nabla J(\mathbf{w}) = \mathbf{0}$ , we obtain

$$(\mathbf{w}^T \mathbf{S}_B \mathbf{w}) \mathbf{S}_W \mathbf{w} = (\mathbf{w}^T \mathbf{S}_W \mathbf{w}) \mathbf{S}_B \mathbf{w}$$

# Solution to Fisher Criterion

- Equating  $\nabla J(\mathbf{w}) = \mathbf{0}$ , we obtain

$$(\mathbf{w}^T \mathbf{S}_B \mathbf{w}) \mathbf{S}_W \mathbf{w} = (\mathbf{w}^T \mathbf{S}_W \mathbf{w}) \mathbf{S}_B \mathbf{w}$$

- $\mathbf{w}^T \mathbf{S}_B \mathbf{w}$  and  $\mathbf{w}^T \mathbf{S}_W \mathbf{w}$  contribute to only scalar multiples

# Solution to Fisher Criterion

- Equating  $\nabla J(\mathbf{w}) = \mathbf{0}$ , we obtain

$$(\mathbf{w}^T \mathbf{S}_B \mathbf{w}) \mathbf{S}_W \mathbf{w} = (\mathbf{w}^T \mathbf{S}_W \mathbf{w}) \mathbf{S}_B \mathbf{w}$$

- $\mathbf{w}^T \mathbf{S}_B \mathbf{w}$  and  $\mathbf{w}^T \mathbf{S}_W \mathbf{w}$  contribute to only scalar multiples
- The direction of  $\mathbf{w}$  is given by  $\mathbf{S}_W \mathbf{w} \propto \mathbf{S}_B \mathbf{w}$

# Solution to Fisher Criterion

- Equating  $\nabla J(\mathbf{w}) = \mathbf{0}$ , we obtain

$$(\mathbf{w}^T \mathbf{S}_B \mathbf{w}) \mathbf{S}_W \mathbf{w} = (\mathbf{w}^T \mathbf{S}_W \mathbf{w}) \mathbf{S}_B \mathbf{w}$$

- $\mathbf{w}^T \mathbf{S}_B \mathbf{w}$  and  $\mathbf{w}^T \mathbf{S}_W \mathbf{w}$  contribute to only scalar multiples
- The direction of  $\mathbf{w}$  is given by  $\mathbf{S}_W \mathbf{w} \propto \mathbf{S}_B \mathbf{w}$
- $\mathbf{S}_B \mathbf{w}$  is always in the direction of  $\mathbf{m}_2 - \mathbf{m}_1$

# Solution to Fisher Criterion

- Equating  $\nabla J(\mathbf{w}) = \mathbf{0}$ , we obtain

$$(\mathbf{w}^T \mathbf{S}_B \mathbf{w}) \mathbf{S}_W \mathbf{w} = (\mathbf{w}^T \mathbf{S}_W \mathbf{w}) \mathbf{S}_B \mathbf{w}$$

- $\mathbf{w}^T \mathbf{S}_B \mathbf{w}$  and  $\mathbf{w}^T \mathbf{S}_W \mathbf{w}$  contribute to only scalar multiples
- The direction of  $\mathbf{w}$  is given by  $\mathbf{S}_W \mathbf{w} \propto \mathbf{S}_B \mathbf{w}$
- $\mathbf{S}_B \mathbf{w}$  is always in the direction of  $\mathbf{m}_2 - \mathbf{m}_1$
- The direction of the fisher discriminant:

$$\mathbf{w} \propto \mathbf{S}_W^{-1}(\mathbf{m}_2 - \mathbf{m}_1)$$



# Solution to Fisher Criterion

- Equating  $\nabla J(\mathbf{w}) = \mathbf{0}$ , we obtain

$$(\mathbf{w}^T \mathbf{S}_B \mathbf{w}) \mathbf{S}_W \mathbf{w} = (\mathbf{w}^T \mathbf{S}_W \mathbf{w}) \mathbf{S}_B \mathbf{w}$$

- $\mathbf{w}^T \mathbf{S}_B \mathbf{w}$  and  $\mathbf{w}^T \mathbf{S}_W \mathbf{w}$  contribute to only scalar multiples
- The direction of  $\mathbf{w}$  is given by  $\mathbf{S}_W \mathbf{w} \propto \mathbf{S}_B \mathbf{w}$
- $\mathbf{S}_B \mathbf{w}$  is always in the direction of  $\mathbf{m}_2 - \mathbf{m}_1$
- The direction of the fisher discriminant:  $\mathbf{w} \propto \mathbf{S}_W^{-1}(\mathbf{m}_2 - \mathbf{m}_1)$
- Multiplication by  $\mathbf{S}_W^{-1}$  can be interpreted as data whitening.

# Solution to Fisher Criterion

- Equating  $\nabla J(\mathbf{w}) = \mathbf{0}$ , we obtain

$$(\mathbf{w}^T \mathbf{S}_B \mathbf{w}) \mathbf{S}_W \mathbf{w} = (\mathbf{w}^T \mathbf{S}_W \mathbf{w}) \mathbf{S}_B \mathbf{w}$$

- $\mathbf{w}^T \mathbf{S}_B \mathbf{w}$  and  $\mathbf{w}^T \mathbf{S}_W \mathbf{w}$  contribute to only scalar multiples
- The direction of  $\mathbf{w}$  is given by  $\mathbf{S}_W \mathbf{w} \propto \mathbf{S}_B \mathbf{w}$
- $\mathbf{S}_B \mathbf{w}$  is always in the direction of  $\mathbf{m}_2 - \mathbf{m}_1$
- The direction of the fisher discriminant:  $\mathbf{w} \propto \mathbf{S}_W^{-1}(\mathbf{m}_2 - \mathbf{m}_1)$
- Multiplication by  $\mathbf{S}_W^{-1}$  can be interpreted as data whitening.
- FDA also is sensitive to outliers even if they are "too correct"

# Solution to Fisher Criterion

- Equating  $\nabla J(\mathbf{w}) = \mathbf{0}$ , we obtain

$$(\mathbf{w}^T \mathbf{S}_B \mathbf{w}) \mathbf{S}_W \mathbf{w} = (\mathbf{w}^T \mathbf{S}_W \mathbf{w}) \mathbf{S}_B \mathbf{w}$$

- $\mathbf{w}^T \mathbf{S}_B \mathbf{w}$  and  $\mathbf{w}^T \mathbf{S}_W \mathbf{w}$  contribute to only scalar multiples
- The direction of  $\mathbf{w}$  is given by  $\mathbf{S}_W \mathbf{w} \propto \mathbf{S}_B \mathbf{w}$
- $\mathbf{S}_B \mathbf{w}$  is always in the direction of  $\mathbf{m}_2 - \mathbf{m}_1$
- The direction of the fisher discriminant:  $\mathbf{w} \propto \mathbf{S}_W^{-1}(\mathbf{m}_2 - \mathbf{m}_1)$
- Multiplication by  $\mathbf{S}_W^{-1}$  can be interpreted as data whitening.
- FDA also is sensitive to outliers even if they are "too correct"
- FD is strictly not a "discriminant" in true sense.

# Solution to Fisher Criterion

- Equating  $\nabla J(\mathbf{w}) = \mathbf{0}$ , we obtain

$$(\mathbf{w}^T \mathbf{S}_B \mathbf{w}) \mathbf{S}_W \mathbf{w} = (\mathbf{w}^T \mathbf{S}_W \mathbf{w}) \mathbf{S}_B \mathbf{w}$$

- $\mathbf{w}^T \mathbf{S}_B \mathbf{w}$  and  $\mathbf{w}^T \mathbf{S}_W \mathbf{w}$  contribute to only scalar multiples
- The direction of  $\mathbf{w}$  is given by  $\mathbf{S}_W \mathbf{w} \propto \mathbf{S}_B \mathbf{w}$
- $\mathbf{S}_B \mathbf{w}$  is always in the direction of  $\mathbf{m}_2 - \mathbf{m}_1$
- The direction of the fisher discriminant:  $\mathbf{w} \propto \mathbf{S}_W^{-1}(\mathbf{m}_2 - \mathbf{m}_1)$
- Multiplication by  $\mathbf{S}_W^{-1}$  can be interpreted as data whitening.
- FDA also is sensitive to outliers even if they are "too correct"
- FD is strictly not a "discriminant" in true sense.
  - It just projects binary class data to 1-D

# Solution to Fisher Criterion

- Equating  $\nabla J(\mathbf{w}) = \mathbf{0}$ , we obtain

$$(\mathbf{w}^T \mathbf{S}_B \mathbf{w}) \mathbf{S}_W \mathbf{w} = (\mathbf{w}^T \mathbf{S}_W \mathbf{w}) \mathbf{S}_B \mathbf{w}$$

- $\mathbf{w}^T \mathbf{S}_B \mathbf{w}$  and  $\mathbf{w}^T \mathbf{S}_W \mathbf{w}$  contribute to only scalar multiples
- The direction of  $\mathbf{w}$  is given by  $\mathbf{S}_W \mathbf{w} \propto \mathbf{S}_B \mathbf{w}$
- $\mathbf{S}_B \mathbf{w}$  is always in the direction of  $\mathbf{m}_2 - \mathbf{m}_1$
- The direction of the fisher discriminant:  $\mathbf{w} \propto \mathbf{S}_W^{-1}(\mathbf{m}_2 - \mathbf{m}_1)$
- Multiplication by  $\mathbf{S}_W^{-1}$  can be interpreted as data whitening.
- FDA also is sensitive to outliers even if they are "too correct"
- FD is strictly not a "discriminant" in true sense.
  - It just projects binary class data to 1-D
  - We need to arrive at a threshold  $w_0$  to perform classification

# Homework: Relation to Least Squares

- In LS approach, linear discriminant is determined to make model predictions as close as possible to target values
- In FDA, the discriminant is derived to achieve maximum class separation in the projected space
- If we take targets for  $\mathcal{C}_1$  and  $\mathcal{C}_2$  as  $\frac{N}{N_1}$  and  $-\frac{N}{N_2}$ , respectively, where  $N = N_1 + N_2$ , show that LS approach yields the same solution as FD.

# Extending FDA to Multiple Classes

# Extending FDA to Multiple Classes

- Assumption: Data dimensionality  $D > K$  Number of classes



## Extending FDA to Multiple Classes

- Assumption: Data dimensionality  $D > K$  Number of classes
- For multi-class case, it is not enough to project data to 1-D

# Extending FDA to Multiple Classes

- Assumption: Data dimensionality  $D > K$  Number of classes
- For multi-class case, it is not enough to project data to 1-D
- Project the data to a lower dimensional space  $D' < D$

# Extending FDA to Multiple Classes

- Assumption: Data dimensionality  $D > K$  Number of classes
- For multi-class case, it is not enough to project data to 1-D
- Project the data to a lower dimensional space  $D' < D$ 
  - Let  $\mathbf{W} : \mathbb{R}^{D \times D'}$  be the linear map to achieve dimensionality reduction

# Extending FDA to Multiple Classes

- Assumption: Data dimensionality  $D > K$  Number of classes
- For multi-class case, it is not enough to project data to 1-D
- Project the data to a lower dimensional space  $D' < D$ 
  - Let  $\mathbf{W} : \mathbb{R}^{D \times D'}$  be the linear map to achieve dimensionality reduction
  - The lower dimensional representation of  $\mathbf{x}_n \in \mathbb{R}^D$  is  $\mathbf{y}_n = \mathbf{W}^T \mathbf{x}_n \in \mathbb{R}^{D'}$

# Extending FDA to Multiple Classes

- Assumption: Data dimensionality  $D > K$  Number of classes
- For multi-class case, it is not enough to project data to 1-D
- Project the data to a lower dimensional space  $D' < D$ 
  - Let  $\mathbf{W} : \mathbb{R}^{D \times D'}$  be the linear map to achieve dimensionality reduction
  - The lower dimensional representation of  $\mathbf{x}_n \in \mathbb{R}^D$  is  $\mathbf{y}_n = \mathbf{W}^T \mathbf{x}_n \in \mathbb{R}^{D'}$
  - Maximize between-class spread and minimize within-class spread in the projected space

## Extending FDA to Multiple Classes

- Assumption: Data dimensionality  $D > K$  Number of classes
- For multi-class case, it is not enough to project data to 1-D
- Project the data to a lower dimensional space  $D' < D$ 
  - Let  $\mathbf{W} : \mathbb{R}^{D \times D'}$  be the linear map to achieve dimensionality reduction
  - The lower dimensional representation of  $\mathbf{x}_n \in \mathbb{R}^D$  is  $\mathbf{y}_n = \mathbf{W}^T \mathbf{x}_n \in \mathbb{R}^{D'}$
  - Maximize between-class spread and minimize within-class spread in the projected space
  - Let  $\mathbf{S}_W$  and  $\mathbf{S}_B$  denote within-class and between-class covariance in original D-dim space

## Extending FDA to Multiple Classes

- Assumption: Data dimensionality  $D > K$  Number of classes
- For multi-class case, it is not enough to project data to 1-D
- Project the data to a lower dimensional space  $D' < D$ 
  - Let  $\mathbf{W} : \mathbb{R}^{D \times D'}$  be the linear map to achieve dimensionality reduction
  - The lower dimensional representation of  $\mathbf{x}_n \in \mathbb{R}^D$  is  $\mathbf{y}_n = \mathbf{W}^T \mathbf{x}_n \in \mathbb{R}^{D'}$
  - Maximize between-class spread and minimize within-class spread in the projected space
  - Let  $\mathbf{S}_W$  and  $\mathbf{S}_B$  denote within-class and between-class covariance in original D-dim space
  - Let  $\mathbf{\Sigma}_W$  and  $\mathbf{\Sigma}_B$  be their counterparts in projected space

# Extending FDA to Multiple Classes

- Assumption: Data dimensionality  $D > K$  Number of classes
- For multi-class case, it is not enough to project data to 1-D
- Project the data to a lower dimensional space  $D' < D$ 
  - Let  $\mathbf{W} : \mathbb{R}^{D \times D'}$  be the linear map to achieve dimensionality reduction
  - The lower dimensional representation of  $\mathbf{x}_n \in \mathbb{R}^D$  is  $\mathbf{y}_n = \mathbf{W}^T \mathbf{x}_n \in \mathbb{R}^{D'}$
  - Maximize between-class spread and minimize within-class spread in the projected space
  - Let  $\mathbf{S}_W$  and  $\mathbf{S}_B$  denote within-class and between-class covariance in original D-dim space
  - Let  $\boldsymbol{\Sigma}_W$  and  $\boldsymbol{\Sigma}_B$  be their counterparts in projected space
  - Estimate  $\mathbf{W}$  to maximize  $J(\mathbf{W}) = \text{Tr} \{ \boldsymbol{\Sigma}_W^{-1} \boldsymbol{\Sigma}_B \}$



# Extending FDA to Multiple Classes

- Assumption: Data dimensionality  $D > K$  Number of classes
- For multi-class case, it is not enough to project data to 1-D
- Project the data to a lower dimensional space  $D' < D$ 
  - Let  $\mathbf{W} : \mathbb{R}^{D \times D'}$  be the linear map to achieve dimensionality reduction
  - The lower dimensional representation of  $\mathbf{x}_n \in \mathbb{R}^D$  is  $\mathbf{y}_n = \mathbf{W}^T \mathbf{x}_n \in \mathbb{R}^{D'}$
  - Maximize between-class spread and minimize within-class spread in the projected space
  - Let  $\mathbf{S}_W$  and  $\mathbf{S}_B$  denote within-class and between-class covariance in original D-dim space
  - Let  $\boldsymbol{\Sigma}_W$  and  $\boldsymbol{\Sigma}_B$  be their counterparts in projected space
  - Estimate  $\mathbf{W}$  to maximize  $J(\mathbf{W}) = \text{Tr} \{ \boldsymbol{\Sigma}_W^{-1} \boldsymbol{\Sigma}_B \}$
- The columns of  $\mathbf{W}$  are given by eigenvectors corresponding to the  $D'$  largest eigenvalues of  $\mathbf{S}_W^{-1} \mathbf{S}_B$

# Statistics in Original $D$ -dim Space

## Statistics in Original $D$ -dim Space

- Class specific statistics in the original  $D$ -dim space

$$\mathbf{m}_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{x}_n$$

## Statistics in Original $D$ -dim Space

- Class specific statistics in the original  $D$ -dim space

$$\mathbf{m}_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{x}_n \quad \mathbf{S}_k = \sum_{n \in \mathcal{C}_k} (\mathbf{x}_n - \mathbf{m}_k)(\mathbf{x}_n - \mathbf{m}_k)^T$$

## Statistics in Original $D$ -dim Space

- Class specific statistics in the original  $D$ -dim space

$$\mathbf{m}_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{x}_n \quad \mathbf{S}_k = \sum_{n \in \mathcal{C}_k} (\mathbf{x}_n - \mathbf{m}_k)(\mathbf{x}_n - \mathbf{m}_k)^T$$

- Define *within-class* covariance as  $\mathbf{S}_W = \sum_{k=1}^K \mathbf{S}_k$

## Statistics in Original $D$ -dim Space

- Class specific statistics in the original  $D$ -dim space

$$\mathbf{m}_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{x}_n \quad \mathbf{S}_k = \sum_{n \in \mathcal{C}_k} (\mathbf{x}_n - \mathbf{m}_k)(\mathbf{x}_n - \mathbf{m}_k)^T$$

- Define *within-class* covariance as  $\mathbf{S}_W = \sum_{k=1}^K \mathbf{S}_k$
- Overall data statistics (without considering class labels)

$$\mathbf{m}_T = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$$

## Statistics in Original $D$ -dim Space

- Class specific statistics in the original  $D$ -dim space

$$\mathbf{m}_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{x}_n \quad \mathbf{S}_k = \sum_{n \in \mathcal{C}_k} (\mathbf{x}_n - \mathbf{m}_k)(\mathbf{x}_n - \mathbf{m}_k)^T$$

- Define *within-class* covariance as  $\mathbf{S}_W = \sum_{k=1}^K \mathbf{S}_k$
- Overall data statistics (without considering class labels)

$$\mathbf{m}_T = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n = \frac{1}{N} \sum_{k=1}^K N_k \mathbf{m}_k$$

## Statistics in Original $D$ -dim Space

- Class specific statistics in the original  $D$ -dim space

$$\mathbf{m}_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{x}_n \quad \mathbf{S}_k = \sum_{n \in \mathcal{C}_k} (\mathbf{x}_n - \mathbf{m}_k)(\mathbf{x}_n - \mathbf{m}_k)^T$$

- Define *within-class* covariance as  $\mathbf{S}_W = \sum_{k=1}^K \mathbf{S}_k$
- Overall data statistics (without considering class labels)

$$\mathbf{m}_T = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n = \frac{1}{N} \sum_{k=1}^K N_k \mathbf{m}_k \quad \mathbf{S}_T = \sum_{n=1}^N (\mathbf{x}_n - \mathbf{m}_T)(\mathbf{x}_n - \mathbf{m}_T)^T$$



## Statistics in Original $D$ -dim Space

- Class specific statistics in the original  $D$ -dim space

$$\mathbf{m}_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{x}_n \quad \mathbf{S}_k = \sum_{n \in \mathcal{C}_k} (\mathbf{x}_n - \mathbf{m}_k)(\mathbf{x}_n - \mathbf{m}_k)^\top$$

- Define *within-class* covariance as  $\mathbf{S}_W = \sum_{k=1}^K \mathbf{S}_k$
- Overall data statistics (without considering class labels)

$$\mathbf{m}_T = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n = \frac{1}{N} \sum_{k=1}^K N_k \mathbf{m}_k \quad \mathbf{S}_T = \sum_{n=1}^N (\mathbf{x}_n - \mathbf{m}_T)(\mathbf{x}_n - \mathbf{m}_T)^\top$$

- Assumption: Define *between-class* covariance using  $\mathbf{S}_T = \mathbf{S}_W + \mathbf{S}_B$

## Statistics in Original $D$ -dim Space

- Class specific statistics in the original  $D$ -dim space

$$\mathbf{m}_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{x}_n \quad \mathbf{S}_k = \sum_{n \in \mathcal{C}_k} (\mathbf{x}_n - \mathbf{m}_k)(\mathbf{x}_n - \mathbf{m}_k)^T$$

- Define *within-class* covariance as  $\mathbf{S}_W = \sum_{k=1}^K \mathbf{S}_k$
- Overall data statistics (without considering class labels)

$$\mathbf{m}_T = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n = \frac{1}{N} \sum_{k=1}^K N_k \mathbf{m}_k \quad \mathbf{S}_T = \sum_{n=1}^N (\mathbf{x}_n - \mathbf{m}_T)(\mathbf{x}_n - \mathbf{m}_T)^T$$

- Assumption: Define *between-class* covariance using  $\mathbf{S}_T = \mathbf{S}_W + \mathbf{S}_B$

$$\mathbf{S}_B = \sum_{k=1}^K N_k (\mathbf{m}_k - \mathbf{m}_T)(\mathbf{m}_k - \mathbf{m}_T)^T$$

# Statistics in Projected $D'$ -dim Space

# Statistics in Projected $D'$ -dim Space

- The projected data-points in  $\mathbb{R}^{D'}$  are given by  $\mathbf{y}_n = \mathbf{W}^T \mathbf{x}_n$

# Statistics in Projected $D'$ -dim Space

- The projected data-points in  $\mathbb{R}^{D'}$  are given by  $\mathbf{y}_n = \mathbf{W}^T \mathbf{x}_n$
- The class-specific statistics in the projected space are

$$\mu_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{y}_n$$

# Statistics in Projected $D'$ -dim Space

- The projected data-points in  $\mathbb{R}^{D'}$  are given by  $\mathbf{y}_n = \mathbf{W}^T \mathbf{x}_n$
- The class-specific statistics in the projected space are

$$\boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{y}_n \quad \boldsymbol{\Sigma}_k = \sum_{n \in \mathcal{C}_k} (\mathbf{y}_n - \boldsymbol{\mu}_k)(\mathbf{y}_n - \boldsymbol{\mu}_k)^T$$

# Statistics in Projected $D'$ -dim Space

- The projected data-points in  $\mathbb{R}^{D'}$  are given by  $\mathbf{y}_n = \mathbf{W}^T \mathbf{x}_n$
- The class-specific statistics in the projected space are

$$\boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{y}_n \quad \boldsymbol{\Sigma}_k = \sum_{n \in \mathcal{C}_k} (\mathbf{y}_n - \boldsymbol{\mu}_k)(\mathbf{y}_n - \boldsymbol{\mu}_k)^T$$

- Relation between the statistics of original and projected space

# Statistics in Projected $D'$ -dim Space

- The projected data-points in  $\mathbb{R}^{D'}$  are given by  $\mathbf{y}_n = \mathbf{W}^T \mathbf{x}_n$
- The class-specific statistics in the projected space are

$$\boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{y}_n \quad \boldsymbol{\Sigma}_k = \sum_{n \in \mathcal{C}_k} (\mathbf{y}_n - \boldsymbol{\mu}_k)(\mathbf{y}_n - \boldsymbol{\mu}_k)^T$$

- Relation between the statistics of original and projected space

$$\boldsymbol{\mu}_k = \mathbf{W}^T \mathbf{m}_k$$



# Statistics in Projected $D'$ -dim Space

- The projected data-points in  $\mathbb{R}^{D'}$  are given by  $\mathbf{y}_n = \mathbf{W}^T \mathbf{x}_n$
- The class-specific statistics in the projected space are

$$\boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{y}_n \quad \boldsymbol{\Sigma}_k = \sum_{n \in \mathcal{C}_k} (\mathbf{y}_n - \boldsymbol{\mu}_k)(\mathbf{y}_n - \boldsymbol{\mu}_k)^T$$

- Relation between the statistics of original and projected space

$$\boldsymbol{\mu}_k = \mathbf{W}^T \mathbf{m}_k \quad \boldsymbol{\Sigma}_k = \mathbf{W}^T \mathbf{S}_k \mathbf{W}$$

# Statistics in Projected $D'$ -dim Space

- The projected data-points in  $\mathbb{R}^{D'}$  are given by  $\mathbf{y}_n = \mathbf{W}^T \mathbf{x}_n$
- The class-specific statistics in the projected space are

$$\boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{y}_n \quad \boldsymbol{\Sigma}_k = \sum_{n \in \mathcal{C}_k} (\mathbf{y}_n - \boldsymbol{\mu}_k)(\mathbf{y}_n - \boldsymbol{\mu}_k)^T$$

- Relation between the statistics of original and projected space

$$\boldsymbol{\mu}_k = \mathbf{W}^T \mathbf{m}_k \quad \boldsymbol{\Sigma}_k = \mathbf{W}^T \mathbf{S}_k \mathbf{W} \quad \boldsymbol{\Sigma}_W = \mathbf{W}^T \mathbf{S}_W \mathbf{W} \quad \boldsymbol{\Sigma}_B = \mathbf{W}^T \mathbf{S}_B \mathbf{W}$$

# Statistics in Projected $D'$ -dim Space

- The projected data-points in  $\mathbb{R}^{D'}$  are given by  $\mathbf{y}_n = \mathbf{W}^T \mathbf{x}_n$
- The class-specific statistics in the projected space are

$$\boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{y}_n \quad \boldsymbol{\Sigma}_k = \sum_{n \in \mathcal{C}_k} (\mathbf{y}_n - \boldsymbol{\mu}_k)(\mathbf{y}_n - \boldsymbol{\mu}_k)^T$$

- Relation between the statistics of original and projected space

$$\boldsymbol{\mu}_k = \mathbf{W}^T \mathbf{m}_k \quad \boldsymbol{\Sigma}_k = \mathbf{W}^T \mathbf{S}_k \mathbf{W} \quad \boldsymbol{\Sigma}_W = \mathbf{W}^T \mathbf{S}_W \mathbf{W} \quad \boldsymbol{\Sigma}_B = \mathbf{W}^T \mathbf{S}_B \mathbf{W}$$

- Estimate  $\mathbf{W}$  to maximize  $J(\mathbf{W}) = \text{Tr} \left\{ (\mathbf{W}^T \mathbf{S}_W \mathbf{W})^{-1} (\mathbf{W}^T \mathbf{S}_B \mathbf{W}) \right\}$

# Statistics in Projected $D'$ -dim Space

- The projected data-points in  $\mathbb{R}^{D'}$  are given by  $\mathbf{y}_n = \mathbf{W}^T \mathbf{x}_n$
- The class-specific statistics in the projected space are

$$\boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{y}_n \quad \boldsymbol{\Sigma}_k = \sum_{n \in \mathcal{C}_k} (\mathbf{y}_n - \boldsymbol{\mu}_k)(\mathbf{y}_n - \boldsymbol{\mu}_k)^T$$

- Relation between the statistics of original and projected space

$$\boldsymbol{\mu}_k = \mathbf{W}^T \mathbf{m}_k \quad \boldsymbol{\Sigma}_k = \mathbf{W}^T \mathbf{S}_k \mathbf{W} \quad \boldsymbol{\Sigma}_W = \mathbf{W}^T \mathbf{S}_W \mathbf{W} \quad \boldsymbol{\Sigma}_B = \mathbf{W}^T \mathbf{S}_B \mathbf{W}$$

- Estimate  $\mathbf{W}$  to maximize  $J(\mathbf{W}) = \text{Tr} \left\{ (\mathbf{W}^T \mathbf{S}_W \mathbf{W})^{-1} (\mathbf{W}^T \mathbf{S}_B \mathbf{W}) \right\}$ 
  - Solution leads to the famous eigenvalue problem:  $\mathbf{S}_W \mathbf{W} \propto \mathbf{S}_B \mathbf{W}$

# Statistics in Projected $D'$ -dim Space

- The projected data-points in  $\mathbb{R}^{D'}$  are given by  $\mathbf{y}_n = \mathbf{W}^T \mathbf{x}_n$
- The class-specific statistics in the projected space are

$$\boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{y}_n \quad \boldsymbol{\Sigma}_k = \sum_{n \in \mathcal{C}_k} (\mathbf{y}_n - \boldsymbol{\mu}_k)(\mathbf{y}_n - \boldsymbol{\mu}_k)^T$$

- Relation between the statistics of original and projected space

$$\boldsymbol{\mu}_k = \mathbf{W}^T \mathbf{m}_k \quad \boldsymbol{\Sigma}_k = \mathbf{W}^T \mathbf{S}_k \mathbf{W} \quad \boldsymbol{\Sigma}_W = \mathbf{W}^T \mathbf{S}_W \mathbf{W} \quad \boldsymbol{\Sigma}_B = \mathbf{W}^T \mathbf{S}_B \mathbf{W}$$

- Estimate  $\mathbf{W}$  to maximize  $J(\mathbf{W}) = \text{Tr} \left\{ (\mathbf{W}^T \mathbf{S}_W \mathbf{W})^{-1} (\mathbf{W}^T \mathbf{S}_B \mathbf{W}) \right\}$ 
  - Solution leads to the famous eigenvalue problem:  $\mathbf{S}_W \mathbf{W} \propto \mathbf{S}_B \mathbf{W}$
  - Eigenvectors corresponding to  $D'$  largest eigenvalues forms  $\mathbf{W}$

## Statistics in Projected $D'$ -dim Space

- The projected data-points in  $\mathbb{R}^{D'}$  are given by  $\mathbf{y}_n = \mathbf{W}^T \mathbf{x}_n$
- The class-specific statistics in the projected space are

$$\boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{y}_n \quad \boldsymbol{\Sigma}_k = \sum_{n \in \mathcal{C}_k} (\mathbf{y}_n - \boldsymbol{\mu}_k)(\mathbf{y}_n - \boldsymbol{\mu}_k)^T$$

- Relation between the statistics of original and projected space

$$\boldsymbol{\mu}_k = \mathbf{W}^T \mathbf{m}_k \quad \boldsymbol{\Sigma}_k = \mathbf{W}^T \mathbf{S}_k \mathbf{W} \quad \boldsymbol{\Sigma}_W = \mathbf{W}^T \mathbf{S}_W \mathbf{W} \quad \boldsymbol{\Sigma}_B = \mathbf{W}^T \mathbf{S}_B \mathbf{W}$$

- Estimate  $\mathbf{W}$  to maximize  $J(\mathbf{W}) = \text{Tr} \left\{ (\mathbf{W}^T \mathbf{S}_W \mathbf{W})^{-1} (\mathbf{W}^T \mathbf{S}_B \mathbf{W}) \right\}$ 
  - Solution leads to the famous eigenvalue problem:  $\mathbf{S}_W \mathbf{W} \propto \mathbf{S}_B \mathbf{W}$
  - Eigenvectors corresponding to  $D'$  largest eigenvalues forms  $\mathbf{W}$
- $D'$  is bounded by rank of  $\mathbf{S}_B$ , and can be at most  $K - 1$

# The Perceptron

# The Perceptron

- Issues with least squares and linear discriminants



# The Perceptron

- Issues with least squares and linear discriminants
  - Linear relation cannot achieve *many-to-one* mapping

# The Perceptron

- Issues with least squares and linear discriminants
  - Linear relation cannot achieve *many-to-one* mapping
  - No inbuilt mechanism to ignore *too-correct* outliers

# The Perceptron

- Issues with least squares and linear discriminants
  - Linear relation cannot achieve *many-to-one* mapping
  - No inbuilt mechanism to ignore *too-correct* outliers
- Perceptron employs a *nonlinearity* to estimate target:  $y(\mathbf{x}) = f(\mathbf{w}^T \mathbf{x})$

# The Perceptron

- Issues with least squares and linear discriminants
  - Linear relation cannot achieve *many-to-one* mapping
  - No inbuilt mechanism to ignore *too-correct* outliers
- Perceptron employs a *nonlinearity* to estimate target:  $y(\mathbf{x}) = f(\mathbf{w}^T \mathbf{x})$
- $f(\cdot)$  is a hard-limiting nonlinear function given by

$$f(a) = \begin{cases} +1, & a \geq 0 \\ -1, & a < 0 \end{cases}$$

# The Perceptron

- Issues with least squares and linear discriminants
  - Linear relation cannot achieve *many-to-one* mapping
  - No inbuilt mechanism to ignore *too-correct* outliers
- Perceptron employs a *nonlinearity* to estimate target:  $y(\mathbf{x}) = f(\mathbf{w}^T \mathbf{x})$
- $f(\cdot)$  is a hard-limiting nonlinear function given by

$$f(a) = \begin{cases} +1, & a \geq 0 \\ -1, & a < 0 \end{cases}$$

- Desired targets are encoded as +1 for  $C_1$  and -1 for  $C_2$ .

# The Perceptron

- Issues with least squares and linear discriminants
  - Linear relation cannot achieve *many-to-one* mapping
  - No inbuilt mechanism to ignore *too-correct* outliers
- Perceptron employs a *nonlinearity* to estimate target:  $y(\mathbf{x}) = f(\mathbf{w}^T \mathbf{x})$
- $f(\cdot)$  is a hard-limiting nonlinear function given by

$$f(a) = \begin{cases} +1, & a \geq 0 \\ -1, & a < 0 \end{cases}$$

- Desired targets are encoded as +1 for  $\mathcal{C}_1$  and -1 for  $\mathcal{C}_2$ .
- Estimate  $\mathbf{w}$  s.t.  $\mathbf{w}^T \mathbf{x}_n \geq 0$  for  $\mathbf{x}_n \in \mathcal{C}_1$ , and  $\mathbf{w}^T \mathbf{x}_n < 0$  for  $\mathbf{x}_n \in \mathcal{C}_2$

# The Perceptron

- Issues with least squares and linear discriminants
  - Linear relation cannot achieve *many-to-one* mapping
  - No inbuilt mechanism to ignore *too-correct* outliers
- Perceptron employs a *nonlinearity* to estimate target:  $y(\mathbf{x}) = f(\mathbf{w}^T \mathbf{x})$
- $f(\cdot)$  is a hard-limiting nonlinear function given by

$$f(a) = \begin{cases} +1, & a \geq 0 \\ -1, & a < 0 \end{cases}$$

- Desired targets are encoded as +1 for  $\mathcal{C}_1$  and -1 for  $\mathcal{C}_2$ .
- Estimate  $\mathbf{w}$  s.t.  $\mathbf{w}^T \mathbf{x}_n \geq 0$  for  $\mathbf{x}_n \in \mathcal{C}_1$ , and  $\mathbf{w}^T \mathbf{x}_n < 0$  for  $\mathbf{x}_n \in \mathcal{C}_2$
- Both the targets and estimates are discrete

# The Perceptron

- Issues with least squares and linear discriminants
  - Linear relation cannot achieve *many-to-one* mapping
  - No inbuilt mechanism to ignore *too-correct* outliers
- Perceptron employs a *nonlinearity* to estimate target:  $y(\mathbf{x}) = f(\mathbf{w}^T \mathbf{x})$
- $f(\cdot)$  is a hard-limiting nonlinear function given by

$$f(a) = \begin{cases} +1, & a \geq 0 \\ -1, & a < 0 \end{cases}$$

- Desired targets are encoded as +1 for  $\mathcal{C}_1$  and -1 for  $\mathcal{C}_2$ .
- Estimate  $\mathbf{w}$  s.t.  $\mathbf{w}^T \mathbf{x}_n \geq 0$  for  $\mathbf{x}_n \in \mathcal{C}_1$ , and  $\mathbf{w}^T \mathbf{x}_n < 0$  for  $\mathbf{x}_n \in \mathcal{C}_2$
- Both the targets and estimates are discrete
- Number of misclassified examples is piece-wise constant



# The Perceptron Learning

# The Perceptron Learning

- Perceptron criterion: defined over set of misclassified examples  $\mathcal{M}$

$$J_p(\mathbf{w}) = - \sum_{n \in \mathcal{M}} (\mathbf{w}^T \mathbf{x}_n) t_n$$

# The Perceptron Learning

- Perceptron criterion: defined over set of misclassified examples  $\mathcal{M}$

$$J_p(\mathbf{w}) = - \sum_{n \in \mathcal{M}} (\mathbf{w}^T \mathbf{x}_n) t_n$$

- SGD can be employed to update the weight vector

$$\mathbf{w}^{new} = \mathbf{w}^{old} - \eta \nabla J_p(\mathbf{w})$$

# The Perceptron Learning

- Perceptron criterion: defined over set of misclassified examples  $\mathcal{M}$

$$J_p(\mathbf{w}) = - \sum_{n \in \mathcal{M}} (\mathbf{w}^T \mathbf{x}_n) t_n$$

- SGD can be employed to update the weight vector

$$\mathbf{w}^{new} = \mathbf{w}^{old} - \eta \nabla J_p(\mathbf{w}) = \mathbf{w}^{old} + \eta \sum_{n \in \mathcal{M}} \mathbf{x}_n t_n$$

# The Perceptron Learning

- Perceptron criterion: defined over set of misclassified examples  $\mathcal{M}$

$$J_p(\mathbf{w}) = - \sum_{n \in \mathcal{M}} (\mathbf{w}^T \mathbf{x}_n) t_n$$

- SGD can be employed to update the weight vector

$$\mathbf{w}^{new} = \mathbf{w}^{old} - \eta \nabla J_p(\mathbf{w}) = \mathbf{w}^{old} + \eta \sum_{n \in \mathcal{M}} \mathbf{x}_n t_n$$

- If  $x_n$  is correctly classified, do not update the weight vector

# The Perceptron Learning

- Perceptron criterion: defined over set of misclassified examples  $\mathcal{M}$

$$J_p(\mathbf{w}) = - \sum_{n \in \mathcal{M}} (\mathbf{w}^T \mathbf{x}_n) t_n$$

- SGD can be employed to update the weight vector

$$\mathbf{w}^{new} = \mathbf{w}^{old} - \eta \nabla J_p(\mathbf{w}) = \mathbf{w}^{old} + \eta \sum_{n \in \mathcal{M}} \mathbf{x}_n t_n$$

- If  $x_n$  is correctly classified, do not update the weight vector
- If  $\mathbf{x}_n \in \mathcal{C}_1$  is misclassified:  $\mathbf{w}^{new} = \mathbf{w}^{old} + \mathbf{x}_n$

# The Perceptron Learning

- Perceptron criterion: defined over set of misclassified examples  $\mathcal{M}$

$$J_p(\mathbf{w}) = - \sum_{n \in \mathcal{M}} (\mathbf{w}^T \mathbf{x}_n) t_n$$

- SGD can be employed to update the weight vector

$$\mathbf{w}^{new} = \mathbf{w}^{old} - \eta \nabla J_p(\mathbf{w}) = \mathbf{w}^{old} + \eta \sum_{n \in \mathcal{M}} \mathbf{x}_n t_n$$

- If  $x_n$  is correctly classified, do not update the weight vector
- If  $\mathbf{x}_n \in \mathcal{C}_1$  is misclassified:  $\mathbf{w}^{new} = \mathbf{w}^{old} + \mathbf{x}_n$
- If  $\mathbf{x}_n \in \mathcal{C}_2$  is misclassified:  $\mathbf{w}^{new} = \mathbf{w}^{old} - \mathbf{x}_n$

# The Perceptron Learning

- Perceptron criterion: defined over set of misclassified examples  $\mathcal{M}$

$$J_p(\mathbf{w}) = - \sum_{n \in \mathcal{M}} (\mathbf{w}^T \mathbf{x}_n) t_n$$

- SGD can be employed to update the weight vector

$$\mathbf{w}^{new} = \mathbf{w}^{old} - \eta \nabla J_p(\mathbf{w}) = \mathbf{w}^{old} + \eta \sum_{n \in \mathcal{M}} \mathbf{x}_n t_n$$

- If  $x_n$  is correctly classified, do not update the weight vector
  - If  $\mathbf{x}_n \in \mathcal{C}_1$  is misclassified:  $\mathbf{w}^{new} = \mathbf{w}^{old} + \mathbf{x}_n$
  - If  $\mathbf{x}_n \in \mathcal{C}_2$  is misclassified:  $\mathbf{w}^{new} = \mathbf{w}^{old} - \mathbf{x}_n$
- Perceptron algorithm is guaranteed to converge in finite steps, for linearly separable classes. [\[Demo\]](#)



# The Perceptron Learning

- Perceptron criterion: defined over set of misclassified examples  $\mathcal{M}$

$$J_p(\mathbf{w}) = - \sum_{n \in \mathcal{M}} (\mathbf{w}^T \mathbf{x}_n) t_n$$

- SGD can be employed to update the weight vector

$$\mathbf{w}^{new} = \mathbf{w}^{old} - \eta \nabla J_p(\mathbf{w}) = \mathbf{w}^{old} + \eta \sum_{n \in \mathcal{M}} \mathbf{x}_n t_n$$

- If  $x_n$  is correctly classified, do not update the weight vector
  - If  $\mathbf{x}_n \in \mathcal{C}_1$  is misclassified:  $\mathbf{w}^{new} = \mathbf{w}^{old} + \mathbf{x}_n$
  - If  $\mathbf{x}_n \in \mathcal{C}_2$  is misclassified:  $\mathbf{w}^{new} = \mathbf{w}^{old} - \mathbf{x}_n$
- Perceptron algorithm is guaranteed to converge in finite steps, for linearly separable classes. [\[Demo\]](#)
- Perceptron algorithm is not vulnerable to *too-correct* outliers

# Perceptron Convergence Theorem

# Perceptron Convergence Theorem

- Starting from a random initial guess  $\mathbf{w}_0$ , perform  $K$  updates:

$$\mathbf{w}_K = \mathbf{w}_{K-1} + \mathbf{x}_K t_K = \mathbf{w}_0 + \sum_{k=1}^K \mathbf{x}_k t_k$$

where  $\mathbf{x}_k$  is the randomly selected misclassified point at  $k^{th}$  iteration

# Perceptron Convergence Theorem

- Starting from a random initial guess  $\mathbf{w}_0$ , perform  $K$  updates:

$$\mathbf{w}_K = \mathbf{w}_{K-1} + \mathbf{x}_K t_K = \mathbf{w}_0 + \sum_{k=1}^K \mathbf{x}_k t_k$$

where  $\mathbf{x}_k$  is the randomly selected misclassified point at  $k^{th}$  iteration

- Letting  $\mathbf{w}_0 = \mathbf{0}$ , the norm of  $\mathbf{w}_K$  is bounded by

$$\|\mathbf{w}_K\| = \left\| \sum_{k=1}^K \mathbf{x}_k t_k \right\| \leq \sum_{k=1}^K \|\mathbf{x}_k\|$$

# Perceptron Convergence Theorem

- Starting from a random initial guess  $\mathbf{w}_0$ , perform  $K$  updates:

$$\mathbf{w}_K = \mathbf{w}_{K-1} + \mathbf{x}_K t_K = \mathbf{w}_0 + \sum_{k=1}^K \mathbf{x}_k t_k$$

where  $\mathbf{x}_k$  is the randomly selected misclassified point at  $k^{th}$  iteration

- Letting  $\mathbf{w}_0 = \mathbf{0}$ , the norm of  $\mathbf{w}_K$  is bounded by

$$\|\mathbf{w}_K\| = \left\| \sum_{k=1}^K \mathbf{x}_k t_k \right\| \leq \sum_{k=1}^K \|\mathbf{x}_k\|$$

- Let  $\alpha$  be the maximum norm in the training data:  $\alpha = \max_n \|\mathbf{x}_n\|$

# Perceptron Convergence Theorem

- Starting from a random initial guess  $\mathbf{w}_0$ , perform  $K$  updates:

$$\mathbf{w}_K = \mathbf{w}_{K-1} + \mathbf{x}_K t_K = \mathbf{w}_0 + \sum_{k=1}^K \mathbf{x}_k t_k$$

where  $\mathbf{x}_k$  is the randomly selected misclassified point at  $k^{th}$  iteration

- Letting  $\mathbf{w}_0 = \mathbf{0}$ , the norm of  $\mathbf{w}_K$  is bounded by

$$\|\mathbf{w}_K\| = \left\| \sum_{k=1}^K \mathbf{x}_k t_k \right\| \leq \sum_{k=1}^K \|\mathbf{x}_k\|$$

- Let  $\alpha$  be the maximum norm in the training data:  $\alpha = \max_n \|\mathbf{x}_n\|$
- The norm of the weight vector at  $K^{th}$  iteration is upper bounded by

$$\|\mathbf{w}_K\| \leq K\alpha$$

# Perceptron Convergence Theorem

# Perceptron Convergence Theorem

- Let  $\mathbf{w}_*$  be one of the solutions that separates classes exactly

$$\mathbf{w}_*^T \mathbf{w}_K = \sum_{k=1}^K \mathbf{w}_*^T \mathbf{x}_k t_k$$



# Perceptron Convergence Theorem

- Let  $\mathbf{w}_*$  be one of the solutions that separates classes exactly

$$\mathbf{w}_*^T \mathbf{w}_K = \sum_{k=1}^K \mathbf{w}_*^T \mathbf{x}_k t_k$$

- Since  $\mathbf{w}_*$  is a solution,  $\mathbf{w}_*^T \mathbf{x}_n t_n \geq 0$  for all  $n$  in the dataset

# Perceptron Convergence Theorem

- Let  $\mathbf{w}_*$  be one of the solutions that separates classes exactly

$$\mathbf{w}_*^T \mathbf{w}_K = \sum_{k=1}^K \mathbf{w}_*^T \mathbf{x}_k t_k$$

- Since  $\mathbf{w}_*$  is a solution,  $\mathbf{w}_*^T \mathbf{x}_n t_n \geq 0$  for all  $n$  in the dataset
- Let  $\beta$  be the minimum projection onto  $\mathbf{w}_*$ :  $\beta = \min_n \mathbf{w}_*^T \mathbf{x}_n t_n$

# Perceptron Convergence Theorem

- Let  $\mathbf{w}_*$  be one of the solutions that separates classes exactly

$$\mathbf{w}_*^T \mathbf{w}_K = \sum_{k=1}^K \mathbf{w}_*^T \mathbf{x}_k t_k$$

- Since  $\mathbf{w}_*$  is a solution,  $\mathbf{w}_*^T \mathbf{x}_n t_n \geq 0$  for all  $n$  in the dataset
- Let  $\beta$  be the minimum projection onto  $\mathbf{w}_*$ :  $\beta = \min_n \mathbf{w}_*^T \mathbf{x}_n t_n$
- Norm of the weight vector  $\mathbf{w}_K$  is lower bounded by

$$\mathbf{w}_*^T \mathbf{w}_K \geq K\beta \quad \|\mathbf{w}_K\| \geq \frac{1}{\|\mathbf{w}_*\|} K^2 \beta^2$$

# Perceptron Convergence Theorem

- Let  $\mathbf{w}_*$  be one of the solutions that separates classes exactly

$$\mathbf{w}_*^T \mathbf{w}_K = \sum_{k=1}^K \mathbf{w}_*^T \mathbf{x}_k t_k$$

- Since  $\mathbf{w}_*$  is a solution,  $\mathbf{w}_*^T \mathbf{x}_n t_n \geq 0$  for all  $n$  in the dataset
- Let  $\beta$  be the minimum projection onto  $\mathbf{w}_*$ :  $\beta = \min_n \mathbf{w}_*^T \mathbf{x}_n t_n$
- Norm of the weight vector  $\mathbf{w}_K$  is lower bounded by

$$\mathbf{w}_*^T \mathbf{w}_K \geq K\beta \quad \|\mathbf{w}_K\| \geq \frac{1}{\|\mathbf{w}_*\|} K^2 \beta^2$$

- The range of the norm of the weight vector is given by

$$K^2 \beta' \leq \|\mathbf{w}_K\| \leq K\alpha$$

# Perceptron Convergence Theorem

- Let  $\mathbf{w}_*$  be one of the solutions that separates classes exactly

$$\mathbf{w}_*^T \mathbf{w}_K = \sum_{k=1}^K \mathbf{w}_*^T \mathbf{x}_k t_k$$

- Since  $\mathbf{w}_*$  is a solution,  $\mathbf{w}_*^T \mathbf{x}_n t_n \geq 0$  for all  $n$  in the dataset
- Let  $\beta$  be the minimum projection onto  $\mathbf{w}_*$ :  $\beta = \min_n \mathbf{w}_*^T \mathbf{x}_n t_n$
- Norm of the weight vector  $\mathbf{w}_K$  is lower bounded by

$$\mathbf{w}_*^T \mathbf{w}_K \geq K\beta \quad \|\mathbf{w}_K\| \geq \frac{1}{\|\mathbf{w}_*\|} K^2 \beta^2$$

- The range of the norm of the weight vector is given by

$$K^2 \beta' \leq \|\mathbf{w}_K\| \leq K\alpha$$

- Lower bound is quadratic in  $K$ , upper bound is linear in  $K$ .

# Limitations of Perceptron

# Limitations of Perceptron

- Perceptron algorithm converges only for linearly separable data

# Limitations of Perceptron

- Perceptron algorithm converges only for linearly separable data
  - Number of iterations  $K$  could be very large



# Limitations of Perceptron

- Perceptron algorithm converges only for linearly separable data
  - Number of iterations  $K$  could be very large
  - Error does not decrease monotonically  $J_P(\mathbf{w}_{k+1}) \not\leq J_P(\mathbf{w}_k)$

# Limitations of Perceptron

- Perceptron algorithm converges only for linearly separable data
  - Number of iterations  $K$  could be very large
  - Error does not decrease monotonically  $J_P(\mathbf{w}_{k+1}) \not\leq J_P(\mathbf{w}_k)$
  - The final solution need not be optimal- one of the decision boundaries.

# Limitations of Perceptron

- Perceptron algorithm converges only for linearly separable data
  - Number of iterations  $K$  could be very large
  - Error does not decrease monotonically  $J_P(\mathbf{w}_{k+1}) \not\leq J_P(\mathbf{w}_k)$
  - The final solution need not be optimal- one of the decision boundaries.
- Does not converge if the classes are not linearly separable - XOR

# Limitations of Perceptron

- Perceptron algorithm converges only for linearly separable data
  - Number of iterations  $K$  could be very large
  - Error does not decrease monotonically  $J_P(\mathbf{w}_{k+1}) \not\leq J_P(\mathbf{w}_k)$
  - The final solution need not be optimal- one of the decision boundaries.
- Does not converge if the classes are not linearly separable - XOR
- Does not offer probabilistic interpretation or confidence intervals.

# Limitations of Perceptron

- Perceptron algorithm converges only for linearly separable data
  - Number of iterations  $K$  could be very large
  - Error does not decrease monotonically  $J_P(\mathbf{w}_{k+1}) \not\leq J_P(\mathbf{w}_k)$
  - The final solution need not be optimal- one of the decision boundaries.
- Does not converge if the classes are not linearly separable - XOR
- Does not offer probabilistic interpretation or confidence intervals.
- Cannot be readily extended to multiclass problems

# Decision Theory

# Decision Theory

- Probability theory offers a mathematical tool to deal with uncertainty.

# Decision Theory

- Probability theory offers a mathematical tool to deal with uncertainty.
- Supervised learning: Predict target  $\mathbf{t}$  from observed input  $\mathbf{x}$



# Decision Theory

- Probability theory offers a mathematical tool to deal with uncertainty.
- Supervised learning: Predict target  $\mathbf{t}$  from observed input  $\mathbf{x}$
- In such a case,  $p(\mathbf{x}, \mathbf{t})$  provides a complete summary of uncertainty

# Decision Theory

- Probability theory offers a mathematical tool to deal with uncertainty.
- Supervised learning: Predict target  $\mathbf{t}$  from observed input  $\mathbf{x}$
- In such a case,  $p(\mathbf{x}, \mathbf{t})$  provides a complete summary of uncertainty
- In classification task,  $\mathbf{t}$  can take discrete labels  $\mathcal{C}_k \quad k = 1, 2, \dots, K$

# Decision Theory

- Probability theory offers a mathematical tool to deal with uncertainty.
- Supervised learning: Predict target  $\mathbf{t}$  from observed input  $\mathbf{x}$
- In such a case,  $p(\mathbf{x}, \mathbf{t})$  provides a complete summary of uncertainty
- In classification task,  $\mathbf{t}$  can take discrete labels  $\mathcal{C}_k \quad k = 1, 2, \dots, K$ 
  - Estimating  $p(\mathbf{x}, \mathcal{C}_k)$  from training set is referred to as *Inference* step

# Decision Theory

- Probability theory offers a mathematical tool to deal with uncertainty.
- Supervised learning: Predict target  $\mathbf{t}$  from observed input  $\mathbf{x}$
- In such a case,  $p(\mathbf{x}, \mathbf{t})$  provides a complete summary of uncertainty
- In classification task,  $\mathbf{t}$  can take discrete labels  $\mathcal{C}_k \quad k = 1, 2, \dots, K$ 
  - Estimating  $p(\mathbf{x}, \mathcal{C}_k)$  from training set is referred to as *Inference* step
  - Assigning a test-point  $\mathbf{x}_t$  to one of the classes is *Decision* step

# Decision Theory

- Probability theory offers a mathematical tool to deal with uncertainty.
- Supervised learning: Predict target  $\mathbf{t}$  from observed input  $\mathbf{x}$
- In such a case,  $p(\mathbf{x}, \mathbf{t})$  provides a complete summary of uncertainty
- In classification task,  $\mathbf{t}$  can take discrete labels  $\mathcal{C}_k \quad k = 1, 2, \dots, K$ 
  - Estimating  $p(\mathbf{x}, \mathcal{C}_k)$  from training set is referred to as *Inference* step
  - Assigning a test-point  $\mathbf{x}_t$  to one of the classes is *Decision* step
- Decision rule divides the input space into  $K$  decision regions  $\mathcal{R}_k$

# Decision Theory

- Probability theory offers a mathematical tool to deal with uncertainty.
- Supervised learning: Predict target  $\mathbf{t}$  from observed input  $\mathbf{x}$
- In such a case,  $p(\mathbf{x}, \mathbf{t})$  provides a complete summary of uncertainty
- In classification task,  $\mathbf{t}$  can take discrete labels  $\mathcal{C}_k \quad k = 1, 2, \dots, K$ 
  - Estimating  $p(\mathbf{x}, \mathcal{C}_k)$  from training set is referred to as *Inference* step
  - Assigning a test-point  $\mathbf{x}_t$  to one of the classes is *Decision* step
- Decision rule divides the input space into  $K$  decision regions  $\mathcal{R}_k$
- Choose *decision boundaries* to minimize the probability of error

# Decision Theory

- Probability theory offers a mathematical tool to deal with uncertainty.
- Supervised learning: Predict target  $\mathbf{t}$  from observed input  $\mathbf{x}$
- In such a case,  $p(\mathbf{x}, \mathbf{t})$  provides a complete summary of uncertainty
- In classification task,  $\mathbf{t}$  can take discrete labels  $\mathcal{C}_k \quad k = 1, 2, \dots, K$ 
  - Estimating  $p(\mathbf{x}, \mathcal{C}_k)$  from training set is referred to as *Inference* step
  - Assigning a test-point  $\mathbf{x}_t$  to one of the classes is *Decision* step
- Decision rule divides the input space into  $K$  decision regions  $\mathcal{R}_k$
- Choose *decision boundaries* to minimize the probability of error
- The probability of error for binary classification is given by

$$p(\text{mistake}) = \int_{\mathcal{R}_1} p(\mathbf{x}, \mathcal{C}_2) d\mathbf{x} + \int_{\mathcal{R}_2} p(\mathbf{x}, \mathcal{C}_1) d\mathbf{x}$$

# Decision Theory

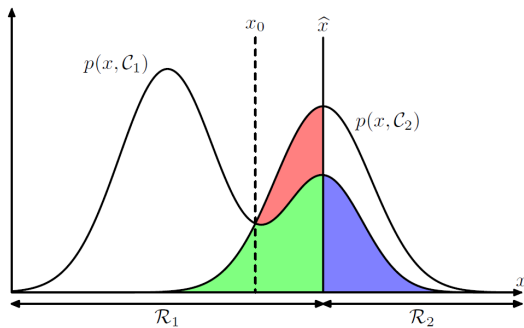
- Probability theory offers a mathematical tool to deal with uncertainty.
- Supervised learning: Predict target  $\mathbf{t}$  from observed input  $\mathbf{x}$
- In such a case,  $p(\mathbf{x}, \mathbf{t})$  provides a complete summary of uncertainty
- In classification task,  $\mathbf{t}$  can take discrete labels  $\mathcal{C}_k \quad k = 1, 2, \dots, K$ 
  - Estimating  $p(\mathbf{x}, \mathcal{C}_k)$  from training set is referred to as *Inference* step
  - Assigning a test-point  $\mathbf{x}_t$  to one of the classes is *Decision* step
- Decision rule divides the input space into  $K$  decision regions  $\mathcal{R}_k$
- Choose *decision boundaries* to minimize the probability of error
- The probability of error for binary classification is given by

$$p(\text{mistake}) = \int_{\mathcal{R}_1} p(\mathbf{x}, \mathcal{C}_2) d\mathbf{x} + \int_{\mathcal{R}_2} p(\mathbf{x}, \mathcal{C}_1) d\mathbf{x}$$

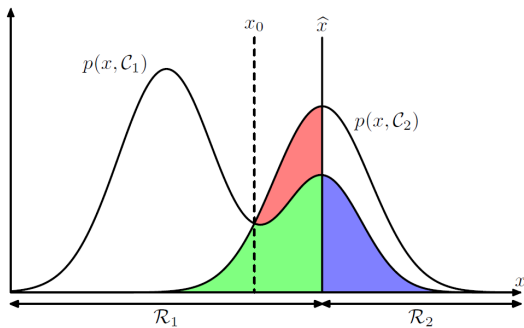
- Divide input space into  $\mathcal{R}_1$  and  $\mathcal{R}_2$  s.t.  $p(\text{mistake})$  is minimized



# Choice of Decision Regions (Boundaries)

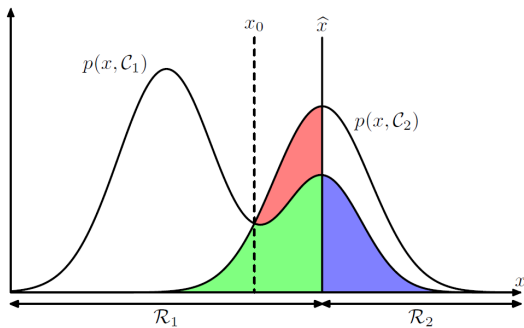


## Choice of Decision Regions (Boundaries)



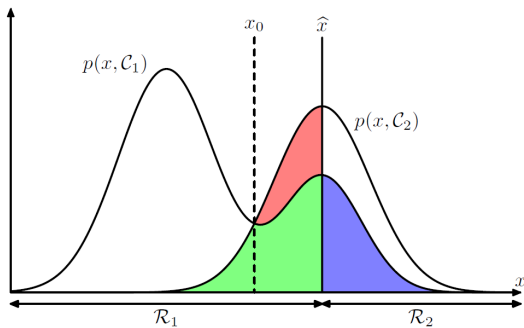
- Let  $x = \hat{x}$  be decision boundary.  $\mathcal{C}_1 : x \leq \hat{x}$  and  $\mathcal{C}_2 : x > \hat{x}$

## Choice of Decision Regions (Boundaries)



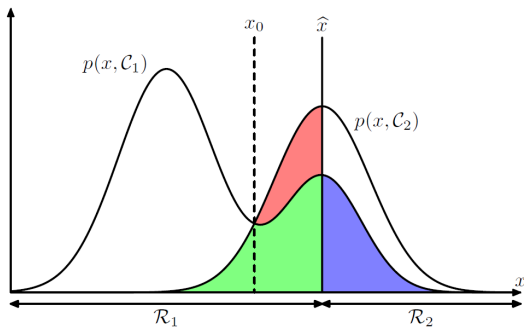
- Let  $x = \hat{x}$  be decision boundary.  $\mathcal{C}_1 : x \leq \hat{x}$  and  $\mathcal{C}_2 : x > \hat{x}$
- Errors arise from area under the blue, red and green regions

## Choice of Decision Regions (Boundaries)



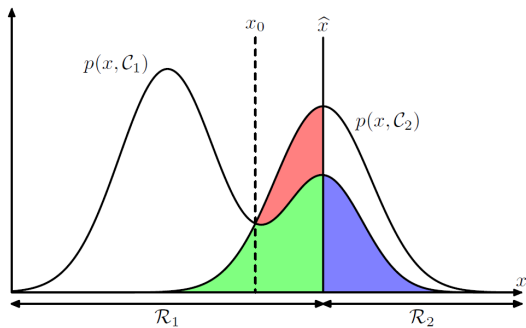
- Let  $x = \hat{x}$  be decision boundary.  $C_1 : x \leq \hat{x}$  and  $C_2 : x > \hat{x}$
- Errors arise from area under the blue, red and green regions
  - Blue: Points from  $C_1$  but misclassified as  $C_2$

# Choice of Decision Regions (Boundaries)



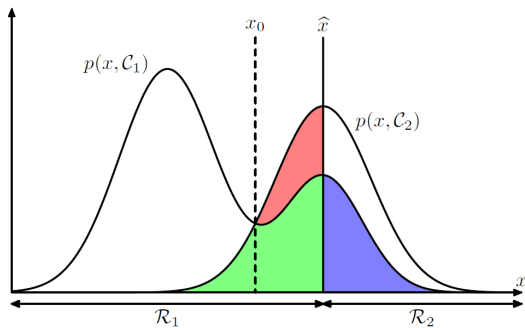
- Let  $x = \hat{x}$  be decision boundary.  $C_1 : x \leq \hat{x}$  and  $C_2 : x > \hat{x}$
- Errors arise from area under the blue, red and green regions
  - Blue: Points from  $C_1$  but misclassified as  $C_2$
  - Green + Red: Points from  $C_2$ , but misclassified as  $C_1$

# Choice of Decision Regions (Boundaries)



- Let  $x = \hat{x}$  be decision boundary.  $C_1 : x \leq \hat{x}$  and  $C_2 : x > \hat{x}$
- Errors arise from area under the blue, red and green regions
  - Blue: Points from  $C_1$  but misclassified as  $C_2$
  - Green + Red: Points from  $C_2$ , but misclassified as  $C_1$
  - Adjust the boundary  $\hat{x}$  to minimize the area under blue+green+red

# Choice of Decision Regions (Boundaries)



- Let  $x = \hat{x}$  be decision boundary.  $C_1 : x \leq \hat{x}$  and  $C_2 : x > \hat{x}$
- Errors arise from area under the blue, red and green regions
  - Blue: Points from  $C_1$  but misclassified as  $C_2$
  - Green + Red: Points from  $C_2$ , but misclassified as  $C_1$
  - Adjust the boundary  $\hat{x}$  to minimize the area under blue+green+red
  - Area under blue+green remains constant irrespective of choice of  $\hat{x}$

# Minimizing Classification Error

$$p(\text{mistake}) = \int_{\mathcal{R}_1} p(\mathbf{x}, \mathcal{C}_2) d\mathbf{x} + \int_{\mathcal{R}_2} p(\mathbf{x}, \mathcal{C}_1) d\mathbf{x}$$



# Minimizing Classification Error

$$\begin{aligned} p(\text{mistake}) &= \int_{\mathcal{R}_1} p(\mathbf{x}, C_2) d\mathbf{x} + \int_{\mathcal{R}_2} p(\mathbf{x}, C_1) d\mathbf{x} \\ &= 1 + \int_{\mathcal{R}_1} p(\mathbf{x}, C_2) d\mathbf{x} - \int_{\mathcal{R}_1} p(\mathbf{x}, C_1) d\mathbf{x} \end{aligned}$$

# Minimizing Classification Error

$$\begin{aligned} p(\text{mistake}) &= \int_{\mathcal{R}_1} p(\mathbf{x}, C_2) d\mathbf{x} + \int_{\mathcal{R}_2} p(\mathbf{x}, C_1) d\mathbf{x} \\ &= 1 + \int_{\mathcal{R}_1} p(\mathbf{x}, C_2) d\mathbf{x} - \int_{\mathcal{R}_1} p(\mathbf{x}, C_1) d\mathbf{x} \\ &= 1 - \int_{\mathcal{R}_1} \left( p(\mathbf{x}, C_1) - p(\mathbf{x}, C_2) \right) d\mathbf{x} \end{aligned}$$

# Minimizing Classification Error

$$\begin{aligned} p(\text{mistake}) &= \int_{\mathcal{R}_1} p(\mathbf{x}, C_2) d\mathbf{x} + \int_{\mathcal{R}_2} p(\mathbf{x}, C_1) d\mathbf{x} \\ &= 1 + \int_{\mathcal{R}_1} p(\mathbf{x}, C_2) d\mathbf{x} - \int_{\mathcal{R}_1} p(\mathbf{x}, C_1) d\mathbf{x} \\ &= 1 - \int_{\mathcal{R}_1} \left( p(\mathbf{x}, C_1) - p(\mathbf{x}, C_2) \right) d\mathbf{x} \end{aligned}$$

- Choose region  $\mathcal{R}_1$  such that  $p(\mathbf{x}, C_1) > p(\mathbf{x}, C_2)$

# Minimizing Classification Error

$$\begin{aligned} p(\text{mistake}) &= \int_{\mathcal{R}_1} p(\mathbf{x}, C_2) d\mathbf{x} + \int_{\mathcal{R}_2} p(\mathbf{x}, C_1) d\mathbf{x} \\ &= 1 + \int_{\mathcal{R}_1} p(\mathbf{x}, C_2) d\mathbf{x} - \int_{\mathcal{R}_1} p(\mathbf{x}, C_1) d\mathbf{x} \\ &= 1 - \int_{\mathcal{R}_1} \left( p(\mathbf{x}, C_1) - p(\mathbf{x}, C_2) \right) d\mathbf{x} \end{aligned}$$

- Choose region  $\mathcal{R}_1$  such that  $p(\mathbf{x}, C_1) > p(\mathbf{x}, C_2)$
- Decision rule: Assign  $\mathbf{x}$  to  $C_1$  if  $p[C_1/\mathbf{x}] > p[C_2/\mathbf{x}]$

# Minimizing Classification Error

$$\begin{aligned} p(\text{mistake}) &= \int_{\mathcal{R}_1} p(\mathbf{x}, C_2) d\mathbf{x} + \int_{\mathcal{R}_2} p(\mathbf{x}, C_1) d\mathbf{x} \\ &= 1 + \int_{\mathcal{R}_1} p(\mathbf{x}, C_2) d\mathbf{x} - \int_{\mathcal{R}_1} p(\mathbf{x}, C_1) d\mathbf{x} \\ &= 1 - \int_{\mathcal{R}_1} (p(\mathbf{x}, C_1) - p(\mathbf{x}, C_2)) d\mathbf{x} \end{aligned}$$

- Choose region  $\mathcal{R}_1$  such that  $p(\mathbf{x}, C_1) > p(\mathbf{x}, C_2)$
- Decision rule: Assign  $\mathbf{x}$  to  $C_1$  if  $p[C_1/\mathbf{x}] > p[C_2/\mathbf{x}]$
- For general case of  $K$  classes, it is easier to maximize  $p(\text{correct})$

$$p(\text{correct}) = \sum_{k=1}^K \int_{\mathcal{R}_k} p(\mathbf{x}, C_k) d\mathbf{x}$$

# Minimizing Classification Error

$$\begin{aligned} p(\text{mistake}) &= \int_{\mathcal{R}_1} p(\mathbf{x}, C_2) d\mathbf{x} + \int_{\mathcal{R}_2} p(\mathbf{x}, C_1) d\mathbf{x} \\ &= 1 + \int_{\mathcal{R}_1} p(\mathbf{x}, C_2) d\mathbf{x} - \int_{\mathcal{R}_1} p(\mathbf{x}, C_1) d\mathbf{x} \\ &= 1 - \int_{\mathcal{R}_1} (p(\mathbf{x}, C_1) - p(\mathbf{x}, C_2)) d\mathbf{x} \end{aligned}$$

- Choose region  $\mathcal{R}_1$  such that  $p(\mathbf{x}, C_1) > p(\mathbf{x}, C_2)$
- Decision rule: Assign  $\mathbf{x}$  to  $C_1$  if  $p[C_1/\mathbf{x}] > p[C_2/\mathbf{x}]$
- For general case of  $K$  classes, it is easier to maximize  $p(\text{correct})$

$$p(\text{correct}) = \sum_{k=1}^K \int_{\mathcal{R}_k} p(\mathbf{x}, C_k) d\mathbf{x}$$

- Decision Rule: Assign  $\mathbf{x}$  to the class with highest posterior  $p[C_k/\mathbf{x}]$

# Expected Loss

# Expected Loss

- In reality, the errors may have varying degrees of consequences



## Expected Loss

- In reality, the errors may have varying degrees of consequences

- Penalty for Healthy vs Cancer classification could be 
$$\begin{matrix} & H & C \\ \begin{matrix} H \\ C \end{matrix} & \begin{pmatrix} 0 & 1 \\ 100 & 0 \end{pmatrix} \end{matrix}$$

# Expected Loss

- In reality, the errors may have varying degrees of consequences

- Penalty for Healthy vs Cancer classification could be 
$$\begin{matrix} & H & C \\ \begin{matrix} H \\ C \end{matrix} & \begin{pmatrix} 0 & 1 \\ 100 & 0 \end{pmatrix} \end{matrix}$$

- Expected loss can be obtained by weighing error with penalty

$$\mathbb{E}[L] = \sum_k \sum_j \int_{\mathcal{R}_j} L_{kj} p(\mathbf{x}, \mathcal{C}_k) d\mathbf{x}$$

# Expected Loss

- In reality, the errors may have varying degrees of consequences

- Penalty for Healthy vs Cancer classification could be 
$$\begin{matrix} & H & C \\ \begin{matrix} H \\ C \end{matrix} & \begin{pmatrix} 0 & 1 \\ 100 & 0 \end{pmatrix} \end{matrix}$$

- Expected loss can be obtained by weighing error with penalty

$$\mathbb{E}[L] = \sum_k \sum_j \int_{\mathcal{R}_j} L_{kj} p(\mathbf{x}, C_k) d\mathbf{x}$$

- Choose region  $\mathcal{R}_j$  to minimize  $\sum_k L_{kj} p(\mathbf{x}, C_k)$

# Expected Loss

- In reality, the errors may have varying degrees of consequences

- Penalty for Healthy vs Cancer classification could be 
$$\begin{matrix} & H & C \\ \begin{matrix} H \\ C \end{matrix} & \begin{pmatrix} 0 & 1 \\ 100 & 0 \end{pmatrix} \end{matrix}$$

- Expected loss can be obtained by weighing error with penalty

$$\mathbb{E}[L] = \sum_k \sum_j \int_{\mathcal{R}_j} L_{kj} p(\mathbf{x}, \mathcal{C}_k) d\mathbf{x}$$

- Choose region  $\mathcal{R}_j$  to minimize  $\sum_k L_{kj} p(\mathbf{x}, \mathcal{C}_k)$
- Decision Rule: Assign  $\mathbf{x}$  to  $\mathcal{C}_j$  that minimizes  $\sum_k L_{kj} p[\mathcal{C}_k/\mathbf{x}]$

# Expected Loss

- In reality, the errors may have varying degrees of consequences

- Penalty for Healthy vs Cancer classification could be 
$$\begin{matrix} & H & C \\ \begin{matrix} H \\ C \end{matrix} & \begin{pmatrix} 0 & 1 \\ 100 & 0 \end{pmatrix} \end{matrix}$$

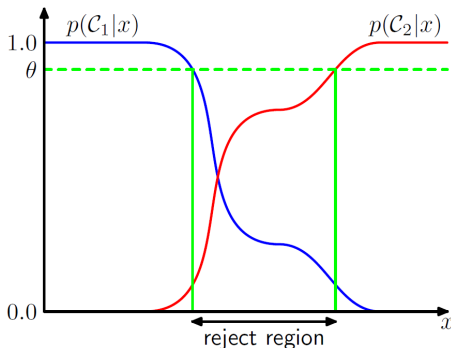
- Expected loss can be obtained by weighing error with penalty

$$\mathbb{E}[L] = \sum_k \sum_j \int_{\mathcal{R}_j} L_{kj} p(\mathbf{x}, \mathcal{C}_k) d\mathbf{x}$$

- Choose region  $\mathcal{R}_j$  to minimize  $\sum_k L_{kj} p(\mathbf{x}, \mathcal{C}_k)$
- Decision Rule: Assign  $\mathbf{x}$  to  $\mathcal{C}_j$  that minimizes  $\sum_k L_{kj} p[\mathcal{C}_k/\mathbf{x}]$
- This is a trivial assignment once posterior probabilities are estimated

# Reject Option

- Errors arises from regions where  $\max_k p[\mathcal{C}_k/\mathbf{x}] \ll 1$ 
  - That means, all the posteriors are in similar range
  - In those regions the classifier is relatively uncertain
- In such cases, it is better to avoid decision making
  - Reject the test samples  $\mathbf{x}$  for which  $\max_k p[\mathcal{C}_k/\mathbf{x}] < \theta$



## Homework: Expected Loss with Reject Option

- Consider a classification problem in which the loss incurred when an input vector from class  $\mathcal{C}_k$  is classified as belonging to class  $\mathcal{C}_j$  is given by the loss matrix  $L_{kj}$ , and for which the loss incurred in selecting the reject option is  $\lambda$ . Find the decision criterion that will give the minimum expected loss. Verify that this reduces to the reject criterion discussed earlier when the loss matrix is given by  $L_{kj} = 1 - I_{kj}$ . What is the relationship between  $\lambda$  and the rejection threshold  $\theta$ ?

# Inference & Decision Stages

- Classification problem can be broken into inference and decision stages



# Inference & Decision Stages

- Classification problem can be broken into inference and decision stages
- Generative models
  - Inference: Estimate posterior  $p[C_k/\mathbf{x}]$  from the joint density  $p(\mathbf{x}, C_k)$

$$p[C_k/\mathbf{x}] = \frac{p[C_k]p(\mathbf{x}/C_k)}{p(\mathbf{x})}$$

- Such a model can generate synthetic examples of a class from  $p(\mathbf{x}, C_k)$

# Inference & Decision Stages

- Classification problem can be broken into inference and decision stages
- Generative models
  - Inference: Estimate posterior  $p[\mathcal{C}_k/\mathbf{x}]$  from the joint density  $p(\mathbf{x}, \mathcal{C}_k)$

$$p[\mathcal{C}_k/\mathbf{x}] = \frac{p[\mathcal{C}_k]p(\mathbf{x}/\mathcal{C}_k)}{p(\mathbf{x})}$$

- Such a model can generate synthetic examples of a class from  $p(\mathbf{x}, \mathcal{C}_k)$
- Discriminative models
  - Inference: Estimate posterior  $p[\mathcal{C}_k/\mathbf{x}]$  as a parametric function  $y(\mathbf{x}, \mathbf{w})$
  - Decision:  $p[\mathcal{C}_k/\mathbf{x}]$  can be used for decision with loss and reject option

# Inference & Decision Stages

- Classification problem can be broken into inference and decision stages
- Generative models
  - Inference: Estimate posterior  $p[C_k/\mathbf{x}]$  from the joint density  $p(\mathbf{x}, C_k)$

$$p[C_k/\mathbf{x}] = \frac{p[C_k]p(\mathbf{x}/C_k)}{p(\mathbf{x})}$$

- Such a model can generate synthetic examples of a class from  $p(\mathbf{x}, C_k)$
- Discriminative models
  - Inference: Estimate posterior  $p[C_k/\mathbf{x}]$  as a parametric function  $y(\mathbf{x}, \mathbf{w})$
  - Decision:  $p[C_k/\mathbf{x}]$  can be used for decision with loss and reject option
- Discriminant functions
  - Find a function  $y(\mathbf{x}, \mathbf{w})$  that maps input  $\mathbf{x}$  to a class label
  - Inference and decision stages cannot be separated

# Pros & Cons

Feature  
Computation

Generate  
High

Discriminative  
Moderate

Discriminant  
Low

# Pros & Cons

| Feature     | Generate  | Discriminative | Discriminant |
|-------------|-----------|----------------|--------------|
| Computation | High      | Moderate       | Low          |
| Data Req.   | Very high | Moderate       | Low          |

## Pros & Cons

| Feature           | Generate            | Discriminative | Discriminant |
|-------------------|---------------------|----------------|--------------|
| Computation       | High                | Moderate       | Low          |
| Data Req.         | Very high           | Moderate       | Low          |
| Outlier detection | Yes $p(\mathbf{x})$ | No             | No           |

## Pros & Cons

| Feature           | Generate            | Discriminative | Discriminant |
|-------------------|---------------------|----------------|--------------|
| Computation       | High                | Moderate       | Low          |
| Data Req.         | Very high           | Moderate       | Low          |
| Outlier detection | Yes $p(\mathbf{x})$ | No             | No           |
| Accuracy          | Reasonable          | Higher         | Low          |

## Pros & Cons

| Feature           | Generate            | Discriminative | Discriminant |
|-------------------|---------------------|----------------|--------------|
| Computation       | High                | Moderate       | Low          |
| Data Req.         | Very high           | Moderate       | Low          |
| Outlier detection | Yes $p(\mathbf{x})$ | No             | No           |
| Accuracy          | Reasonable          | Higher         | Low          |
| Minimizing Risk   | Easy                | Easy           | Not SF       |



## Pros & Cons

| Feature           | Generate            | Discriminative | Discriminant |
|-------------------|---------------------|----------------|--------------|
| Computation       | High                | Moderate       | Low          |
| Data Req.         | Very high           | Moderate       | Low          |
| Outlier detection | Yes $p(\mathbf{x})$ | No             | No           |
| Accuracy          | Reasonable          | Higher         | Low          |
| Minimizing Risk   | Easy                | Easy           | Not SF       |
| Reject option     | Easy                | Easy           | Not SF       |

## Pros & Cons

| Feature           | Generate            | Discriminative | Discriminant |
|-------------------|---------------------|----------------|--------------|
| Computation       | High                | Moderate       | Low          |
| Data Req.         | Very high           | Moderate       | Low          |
| Outlier detection | Yes $p(\mathbf{x})$ | No             | No           |
| Accuracy          | Reasonable          | Higher         | Low          |
| Minimizing Risk   | Easy                | Easy           | Not SF       |
| Reject option     | Easy                | Easy           | Not SF       |
| Modifying Priors  | Easy                | Not SF         | No           |

# Pros & Cons

| Feature           | Generate            | Discriminative | Discriminant |
|-------------------|---------------------|----------------|--------------|
| Computation       | High                | Moderate       | Low          |
| Data Req.         | Very high           | Moderate       | Low          |
| Outlier detection | Yes $p(\mathbf{x})$ | No             | No           |
| Accuracy          | Reasonable          | Higher         | Low          |
| Minimizing Risk   | Easy                | Easy           | Not SF       |
| Reject option     | Easy                | Easy           | Not SF       |
| Modifying Priors  | Easy                | Not SF         | No           |
| Model Fusion      | Easy                | Easy           | Not SF       |

# Pros & Cons

| Feature           | Generate            | Discriminative | Discriminant |
|-------------------|---------------------|----------------|--------------|
| Computation       | High                | Moderate       | Low          |
| Data Req.         | Very high           | Moderate       | Low          |
| Outlier detection | Yes $p(\mathbf{x})$ | No             | No           |
| Accuracy          | Reasonable          | Higher         | Low          |
| Minimizing Risk   | Easy                | Easy           | Not SF       |
| Reject option     | Easy                | Easy           | Not SF       |
| Modifying Priors  | Easy                | Not SF         | No           |
| Model Fusion      | Easy                | Easy           | Not SF       |

$$P[C_k/\mathbf{x}_A, \mathbf{x}_B] \propto p[C_k]p(\mathbf{x}_A, \mathbf{x}_B/C_k)$$

## Pros & Cons

| Feature           | Generate            | Discriminative | Discriminant |
|-------------------|---------------------|----------------|--------------|
| Computation       | High                | Moderate       | Low          |
| Data Req.         | Very high           | Moderate       | Low          |
| Outlier detection | Yes $p(\mathbf{x})$ | No             | No           |
| Accuracy          | Reasonable          | Higher         | Low          |
| Minimizing Risk   | Easy                | Easy           | Not SF       |
| Reject option     | Easy                | Easy           | Not SF       |
| Modifying Priors  | Easy                | Not SF         | No           |
| Model Fusion      | Easy                | Easy           | Not SF       |

$$\begin{aligned}P[C_k/\mathbf{x}_A, \mathbf{x}_B] &\propto p[C_k]p(\mathbf{x}_A, \mathbf{x}_B/C_k) \\ &\propto p[C_k]p(\mathbf{x}_A/C_k)p(\mathbf{x}_B/C_k) \quad \text{Cond. Ind.}\end{aligned}$$

## Pros & Cons

| Feature           | Generate            | Discriminative | Discriminant |
|-------------------|---------------------|----------------|--------------|
| Computation       | High                | Moderate       | Low          |
| Data Req.         | Very high           | Moderate       | Low          |
| Outlier detection | Yes $p(\mathbf{x})$ | No             | No           |
| Accuracy          | Reasonable          | Higher         | Low          |
| Minimizing Risk   | Easy                | Easy           | Not SF       |
| Reject option     | Easy                | Easy           | Not SF       |
| Modifying Priors  | Easy                | Not SF         | No           |
| Model Fusion      | Easy                | Easy           | Not SF       |

$$\begin{aligned}
 P[C_k/\mathbf{x}_A, \mathbf{x}_B] &\propto p[C_k]p(\mathbf{x}_A, \mathbf{x}_B/C_k) \\
 &\propto p[C_k]p(\mathbf{x}_A/C_k)p(\mathbf{x}_B/C_k) \quad \text{Cond. Ind.} \\
 &\propto \frac{p[C_k/\mathbf{x}_A]p[C_k/\mathbf{x}_B]}{p[C_k]}
 \end{aligned}$$

# Maximum Likelihood Density Estimation

# Maximum Likelihood Density Estimation

- Consider data  $X = \{x_1, x_2, \dots, x_N\}$  drawn from unknown distribution



# Maximum Likelihood Density Estimation

- Consider data  $X = \{x_1, x_2, \dots, x_N\}$  drawn from unknown distribution
- Underlying distribution  $p(x_n)$  be approximated by a parametric form

# Maximum Likelihood Density Estimation

- Consider data  $X = \{x_1, x_2, \dots, x_N\}$  drawn from unknown distribution
- Underlying distribution  $p(x_n)$  be approximated by a parametric form
  - Gaussian assumption:  $p(x_n) \sim \mathcal{N}(x_n/\mu, \sigma^2)$

# Maximum Likelihood Density Estimation

- Consider data  $X = \{x_1, x_2, \dots, x_N\}$  drawn from unknown distribution
- Underlying distribution  $p(x_n)$  be approximated by a parametric form
  - Gaussian assumption:  $p(x_n) \sim \mathcal{N}(x_n/\mu, \sigma^2)$
  - Laplacian assumption:  $p(x_n) \sim \mathcal{L}(x_n/\mu, b)$

# Maximum Likelihood Density Estimation

- Consider data  $X = \{x_1, x_2, \dots, x_N\}$  drawn from unknown distribution
- Underlying distribution  $p(x_n)$  be approximated by a parametric form
  - Gaussian assumption:  $p(x_n) \sim \mathcal{N}(x_n/\mu, \sigma^2)$
  - Laplacian assumption:  $p(x_n) \sim \mathcal{L}(x_n/\mu, b)$
- Assuming i.i.d, the likelihood function can be written as

$$\begin{aligned} p(X/\mu, \sigma^2) &= \prod_{n=1}^N p(x_n/\mu, \sigma^2) \\ &= \prod_{n=1}^N \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x_n - \mu)^2}{2\sigma^2}\right) \end{aligned}$$

# Maximum Likelihood Density Estimation

- Consider data  $X = \{x_1, x_2, \dots, x_N\}$  drawn from unknown distribution
- Underlying distribution  $p(x_n)$  be approximated by a parametric form
  - Gaussian assumption:  $p(x_n) \sim \mathcal{N}(x_n/\mu, \sigma^2)$
  - Laplacian assumption:  $p(x_n) \sim \mathcal{L}(x_n/\mu, b)$
- Assuming i.i.d, the likelihood function can be written as

$$\begin{aligned} p(X/\mu, \sigma^2) &= \prod_{n=1}^N p(x_n/\mu, \sigma^2) \\ &= \prod_{n=1}^N \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x_n - \mu)^2}{2\sigma^2}\right) \end{aligned}$$

- Estimate  $\mu$  and  $\sigma$  to maximize the likelihood function, or eqv.

$$J(\mu, \sigma^2) = \log p(X/\mu, \sigma^2)$$

# Parameter Estimation

# Parameter Estimation

- $\mu$  and  $\sigma$  can be determined by equating the derivatives to zero

$$\hat{\mu} = \frac{1}{N} \sum_{n=1}^N x_n \quad \hat{\sigma}^2 = \frac{1}{N} \sum_{n=1}^N (x_n - \hat{\mu})^2$$

# Parameter Estimation

- $\mu$  and  $\sigma$  can be determined by equating the derivatives to zero

$$\hat{\mu} = \frac{1}{N} \sum_{n=1}^N x_n \quad \hat{\sigma}^2 = \frac{1}{N} \sum_{n=1}^N (x_n - \hat{\mu})^2$$

- Checking for *bias* in estimation: Taking expectation over estimates

$$\mathbb{E}[\hat{\mu}] = \mathbb{E} \left[ \frac{1}{N} \sum_{n=1}^N x_n \right] = \frac{1}{N} \sum_{n=1}^N \mathbb{E}[x_n] = \mu \quad \text{unbiased}$$

$$\mathbb{E}[\hat{\sigma}^2] = \frac{1}{N} \sum_{n=1}^N \mathbb{E}[(x_n - \hat{\mu})^2] = \frac{N-1}{N} \sigma^2 \quad \text{biased}$$



# Parameter Estimation

- $\mu$  and  $\sigma$  can be determined by equating the derivatives to zero

$$\hat{\mu} = \frac{1}{N} \sum_{n=1}^N x_n \quad \hat{\sigma}^2 = \frac{1}{N} \sum_{n=1}^N (x_n - \hat{\mu})^2$$

- Checking for *bias* in estimation: Taking expectation over estimates

$$\mathbb{E}[\hat{\mu}] = \mathbb{E}\left[\frac{1}{N} \sum_{n=1}^N x_n\right] = \frac{1}{N} \sum_{n=1}^N \mathbb{E}[x_n] = \mu \quad \text{unbiased}$$

$$\mathbb{E}[\hat{\sigma}^2] = \frac{1}{N} \sum_{n=1}^N \mathbb{E}[(x_n - \hat{\mu})^2] = \frac{N-1}{N} \sigma^2 \quad \text{biased}$$

- Variance is underestimated in ML approach.  $\hat{\sigma}^2 < \sigma^2$

# Parameter Estimation

- $\mu$  and  $\sigma$  can be determined by equating the derivatives to zero

$$\hat{\mu} = \frac{1}{N} \sum_{n=1}^N x_n \quad \hat{\sigma}^2 = \frac{1}{N} \sum_{n=1}^N (x_n - \hat{\mu})^2$$

- Checking for *bias* in estimation: Taking expectation over estimates

$$\mathbb{E}[\hat{\mu}] = \mathbb{E}\left[\frac{1}{N} \sum_{n=1}^N x_n\right] = \frac{1}{N} \sum_{n=1}^N \mathbb{E}[x_n] = \mu \quad \text{unbiased}$$

$$\mathbb{E}[\hat{\sigma}^2] = \frac{1}{N} \sum_{n=1}^N \mathbb{E}[(x_n - \hat{\mu})^2] = \frac{N-1}{N} \sigma^2 \quad \text{biased}$$

- Variance is underestimated in ML approach.  $\hat{\sigma}^2 < \sigma^2$
- The variance estimation can be corrected as  $\tilde{\sigma}^2 = \frac{1}{N-1} \sum_{n=1}^N (x_n - \hat{\mu})^2$

# ML approach to Generative Models

# ML approach to Generative Models

- Training Data:  $\{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_n, t_n), \dots, (\mathbf{x}_N, t_N)\}$

# ML approach to Generative Models

- Training Data:  $\{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_n, t_n), \dots, (\mathbf{x}_N, t_N)\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  are input variables,  $t_n \in \{0, 1\}$  are targets

# ML approach to Generative Models

- Training Data:  $\{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_n, t_n), \dots, (\mathbf{x}_N, t_N)\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  are input variables,  $t_n \in \{0, 1\}$  are targets
  - Let  $N_1$  points are from  $\mathcal{C}_1$  and  $N_2$  points are from  $\mathcal{C}_2$ :  $N = N_1 + N_2$

# ML approach to Generative Models

- Training Data:  $\{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_n, t_n), \dots, (\mathbf{x}_N, t_N)\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  are input variables,  $t_n \in \{0, 1\}$  are targets
  - Let  $N_1$  points are from  $\mathcal{C}_1$  and  $N_2$  points are from  $\mathcal{C}_2$ :  $N = N_1 + N_2$
- Generative models require estimation of joint density  $p(\mathbf{x}_n, \mathcal{C}_k)$

$$p(\mathbf{x}_n, \mathcal{C}_k) = P[\mathcal{C}_k]p(\mathbf{x}_n/\mathcal{C}_k)$$

# ML approach to Generative Models

- Training Data:  $\{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_n, t_n), \dots, (\mathbf{x}_N, t_N)\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  are input variables,  $t_n \in \{0, 1\}$  are targets
  - Let  $N_1$  points are from  $\mathcal{C}_1$  and  $N_2$  points are from  $\mathcal{C}_2$ :  $N = N_1 + N_2$
- Generative models require estimation of joint density  $p(\mathbf{x}_n, \mathcal{C}_k)$

$$p(\mathbf{x}_n, \mathcal{C}_k) = P[\mathcal{C}_k]p(\mathbf{x}_n/\mathcal{C}_k)$$

- Let the priors for  $\mathcal{C}_1$  and  $\mathcal{C}_2$  be  $\pi$  and  $1 - \pi$ , respectively



# ML approach to Generative Models

- Training Data:  $\{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_n, t_n), \dots, (\mathbf{x}_N, t_N)\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  are input variables,  $t_n \in \{0, 1\}$  are targets
  - Let  $N_1$  points are from  $\mathcal{C}_1$  and  $N_2$  points are from  $\mathcal{C}_2$ :  $N = N_1 + N_2$
- Generative models require estimation of joint density  $p(\mathbf{x}_n, \mathcal{C}_k)$

$$p(\mathbf{x}_n, \mathcal{C}_k) = P[\mathcal{C}_k]p(\mathbf{x}_n/\mathcal{C}_k)$$

- Let the priors for  $\mathcal{C}_1$  and  $\mathcal{C}_2$  be  $\pi$  and  $1 - \pi$ , respectively
- Let class conditional densities  $p(\mathbf{x}_n/\mathcal{C}_k)$  follow Gaussian distributions

$$p(\mathbf{x}_n/\mathcal{C}_k) \sim \mathcal{N}(\mathbf{x}_n/\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

# ML approach to Generative Models

- Training Data:  $\{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_n, t_n), \dots, (\mathbf{x}_N, t_N)\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  are input variables,  $t_n \in \{0, 1\}$  are targets
  - Let  $N_1$  points are from  $\mathcal{C}_1$  and  $N_2$  points are from  $\mathcal{C}_2$ :  $N = N_1 + N_2$
- Generative models require estimation of joint density  $p(\mathbf{x}_n, \mathcal{C}_k)$

$$p(\mathbf{x}_n, \mathcal{C}_k) = P[\mathcal{C}_k]p(\mathbf{x}_n/\mathcal{C}_k)$$

- Let the priors for  $\mathcal{C}_1$  and  $\mathcal{C}_2$  be  $\pi$  and  $1 - \pi$ , respectively
- Let class conditional densities  $p(\mathbf{x}_n/\mathcal{C}_k)$  follow Gaussian distributions

$$p(\mathbf{x}_n/\mathcal{C}_k) \sim \mathcal{N}(\mathbf{x}_n/\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

- The joint densities are given by

$$\text{For } \mathbf{x}_n \in \mathcal{C}_1 \quad p(\mathbf{x}_n/\mathcal{C}_1) = \pi \mathcal{N}(\mathbf{x}_n/\boldsymbol{\mu}_1, \boldsymbol{\Sigma}_1)$$

$$\text{For } \mathbf{x}_n \in \mathcal{C}_2 \quad p(\mathbf{x}_n/\mathcal{C}_1) = (1 - \pi) \mathcal{N}(\mathbf{x}_n/\boldsymbol{\mu}_2, \boldsymbol{\Sigma}_2)$$

# ML Solution

# ML Solution

- Estimate  $\theta = (\pi, \mu_1, \mu_2, \Sigma_1, \Sigma_2)$  to maximize the likelihood function

$$p(\mathbf{t}/\theta) = \prod_{n=1}^N [\pi \mathcal{N}(\mathbf{x}_n/\mu_1, \Sigma_1)]^{t_n} [(1 - \pi) \mathcal{N}(\mathbf{x}_n/\mu_2, \Sigma_2)]^{1-t_n}$$

# ML Solution

- Estimate  $\theta = (\pi, \mu_1, \mu_2, \Sigma_1, \Sigma_2)$  to maximize the likelihood function

$$p(\mathbf{t}/\theta) = \prod_{n=1}^N [\pi \mathcal{N}(\mathbf{x}_n/\mu_1, \Sigma_1)]^{t_n} [(1 - \pi) \mathcal{N}(\mathbf{x}_n/\mu_2, \Sigma_2)]^{1-t_n}$$

- It is convenient to maximize log-likelihood  $\log p(\mathbf{t}/\theta)$

# ML Solution

- Estimate  $\theta = (\pi, \mu_1, \mu_2, \Sigma_1, \Sigma_2)$  to maximize the likelihood function

$$p(\mathbf{t}/\theta) = \prod_{n=1}^N [\pi \mathcal{N}(\mathbf{x}_n/\mu_1, \Sigma_1)]^{t_n} [(1 - \pi) \mathcal{N}(\mathbf{x}_n/\mu_2, \Sigma_2)]^{1-t_n}$$

- It is convenient to maximize log-likelihood  $\log p(\mathbf{t}/\theta)$
- The terms involving  $\pi$  in log-likelihood  $\mathcal{L}(\theta)$  are

$$\sum_{n=1}^N t_n \log \pi + (1 - t_n) \log(1 - \pi)$$

# ML Solution

- Estimate  $\theta = (\pi, \mu_1, \mu_2, \Sigma_1, \Sigma_2)$  to maximize the likelihood function

$$p(\mathbf{t}/\theta) = \prod_{n=1}^N [\pi \mathcal{N}(\mathbf{x}_n/\mu_1, \Sigma_1)]^{t_n} [(1 - \pi) \mathcal{N}(\mathbf{x}_n/\mu_2, \Sigma_2)]^{1-t_n}$$

- It is convenient to maximize log-likelihood  $\log p(\mathbf{t}/\theta)$
- The terms involving  $\pi$  in log-likelihood  $\mathcal{L}(\theta)$  are

$$\sum_{n=1}^N t_n \log \pi + (1 - t_n) \log(1 - \pi)$$

- Setting the derivative of  $\mathcal{L}(\theta)$  w.r.t to  $\pi$  to zero,

$$\pi = \frac{N_1}{N} \implies p[C_1] = \frac{N_1}{N} \text{ and } p[C_2] = \frac{N_2}{N}$$

# ML Solution

- Estimate  $\theta = (\pi, \mu_1, \mu_2, \Sigma_1, \Sigma_2)$  to maximize the likelihood function

$$p(\mathbf{t}/\theta) = \prod_{n=1}^N [\pi \mathcal{N}(\mathbf{x}_n/\mu_1, \Sigma_1)]^{t_n} [(1 - \pi) \mathcal{N}(\mathbf{x}_n/\mu_2, \Sigma_2)]^{1-t_n}$$

- It is convenient to maximize log-likelihood  $\log p(\mathbf{t}/\theta)$
- The terms involving  $\pi$  in log-likelihood  $\mathcal{L}(\theta)$  are

$$\sum_{n=1}^N t_n \log \pi + (1 - t_n) \log(1 - \pi)$$

- Setting the derivative of  $\mathcal{L}(\theta)$  w.r.t to  $\pi$  to zero,

$$\pi = \frac{N_1}{N} \implies p[C_1] = \frac{N_1}{N} \text{ and } p[C_2] = \frac{N_2}{N}$$



# ML Solution

# ML Solution

- $\mu_1, \mu_2, \Sigma_1, \Sigma_2$  can be evaluated in a similar manner

$$\mu_1 = \frac{1}{N_1} \sum_{n \in \mathcal{C}_1} \mathbf{x}_n \quad \Sigma_1 = \frac{1}{N_1} \sum_{n \in \mathcal{C}_1} (\mathbf{x}_n - \mu_1)^T (\mathbf{x}_n - \mu_1)$$

$$\mu_2 = \frac{1}{N_2} \sum_{n \in \mathcal{C}_2} \mathbf{x}_n \quad \Sigma_2 = \frac{1}{N_2} \sum_{n \in \mathcal{C}_2} (\mathbf{x}_n - \mu_2)^T (\mathbf{x}_n - \mu_2)$$

# ML Solution

- $\mu_1, \mu_2, \Sigma_1, \Sigma_2$  can be evaluated in a similar manner

$$\mu_1 = \frac{1}{N_1} \sum_{n \in \mathcal{C}_1} \mathbf{x}_n \quad \Sigma_1 = \frac{1}{N_1} \sum_{n \in \mathcal{C}_1} (\mathbf{x}_n - \mu_1)^T (\mathbf{x}_n - \mu_1)$$

$$\mu_2 = \frac{1}{N_2} \sum_{n \in \mathcal{C}_2} \mathbf{x}_n \quad \Sigma_2 = \frac{1}{N_2} \sum_{n \in \mathcal{C}_2} (\mathbf{x}_n - \mu_2)^T (\mathbf{x}_n - \mu_2)$$

- Equivalent to estimating individual CCDs

# ML Solution

- $\mu_1, \mu_2, \Sigma_1, \Sigma_2$  can be evaluated in a similar manner

$$\mu_1 = \frac{1}{N_1} \sum_{n \in \mathcal{C}_1} \mathbf{x}_n \quad \Sigma_1 = \frac{1}{N_1} \sum_{n \in \mathcal{C}_1} (\mathbf{x}_n - \mu_1)^T (\mathbf{x}_n - \mu_1)$$

$$\mu_2 = \frac{1}{N_2} \sum_{n \in \mathcal{C}_2} \mathbf{x}_n \quad \Sigma_2 = \frac{1}{N_2} \sum_{n \in \mathcal{C}_2} (\mathbf{x}_n - \mu_2)^T (\mathbf{x}_n - \mu_2)$$

- Equivalent to estimating individual CCDs
- While estimating CCD of one class, data from other class is not considered

# ML Solution

- $\mu_1, \mu_2, \Sigma_1, \Sigma_2$  can be evaluated in a similar manner

$$\mu_1 = \frac{1}{N_1} \sum_{n \in \mathcal{C}_1} \mathbf{x}_n \quad \Sigma_1 = \frac{1}{N_1} \sum_{n \in \mathcal{C}_1} (\mathbf{x}_n - \mu_1)^T (\mathbf{x}_n - \mu_1)$$

$$\mu_2 = \frac{1}{N_2} \sum_{n \in \mathcal{C}_2} \mathbf{x}_n \quad \Sigma_2 = \frac{1}{N_2} \sum_{n \in \mathcal{C}_2} (\mathbf{x}_n - \mu_2)^T (\mathbf{x}_n - \mu_2)$$

- Equivalent to estimating individual CCDs
- While estimating CCD of one class, data from other class is not considered
- Offers a descriptive model, but need not be discriminative

# Home Work - Shared Covariance Matrix

# Home Work - Shared Covariance Matrix

- Let both the classes share the same covariance matrix

$$\Sigma_1 = \Sigma_2 = \Sigma$$

# Home Work - Shared Covariance Matrix

- Let both the classes share the same covariance matrix

$$\Sigma_1 = \Sigma_2 = \Sigma$$

- The shape of the distributions is same for both classes



# Home Work - Shared Covariance Matrix

- Let both the classes share the same covariance matrix

$$\Sigma_1 = \Sigma_2 = \Sigma$$

- The shape of the distributions is same for both classes
- The ML solution to the shared covariance matrix is given by

$$\Sigma = \frac{N_1}{N} \Sigma_1 + \frac{N_2}{N} \Sigma_2$$

# Home Work - Shared Covariance Matrix

- Let both the classes share the same covariance matrix

$$\Sigma_1 = \Sigma_2 = \Sigma$$

- The shape of the distributions is same for both classes
- The ML solution to the shared covariance matrix is given by

$$\Sigma = \frac{N_1}{N} \Sigma_1 + \frac{N_2}{N} \Sigma_2$$

- Shared covariance is given by a weighted combination of class-specific covariances.

# Decision Boundary(Shared Covariance)

## Decision Boundary(Shared Covariance)

- The decision boundary is locus of points satisfying  $P[\mathcal{C}_1/\mathbf{x}] = P[\mathcal{C}_2/\mathbf{x}]$

$$P[\mathcal{C}_1] \mathcal{N}(\mathbf{x}/\boldsymbol{\mu}_1, \boldsymbol{\Sigma}) = p[\mathcal{C}_2] \mathcal{N}(\mathbf{x}/\boldsymbol{\mu}_2, \boldsymbol{\Sigma})$$

## Decision Boundary(Shared Covariance)

- The decision boundary is locus of points satisfying  $P[\mathcal{C}_1/\mathbf{x}] = P[\mathcal{C}_2/\mathbf{x}]$

$$P[\mathcal{C}_1] \mathcal{N}(\mathbf{x}/\boldsymbol{\mu}_1, \boldsymbol{\Sigma}) = p[\mathcal{C}_2] \mathcal{N}(\mathbf{x}/\boldsymbol{\mu}_2, \boldsymbol{\Sigma})$$

- The quadratic terms cancel because of shared covariance

## Decision Boundary(Shared Covariance)

- The decision boundary is locus of points satisfying  $P[\mathcal{C}_1/\mathbf{x}] = P[\mathcal{C}_2/\mathbf{x}]$

$$P[\mathcal{C}_1] \mathcal{N}(\mathbf{x}/\boldsymbol{\mu}_1, \boldsymbol{\Sigma}) = p[\mathcal{C}_2] \mathcal{N}(\mathbf{x}/\boldsymbol{\mu}_2, \boldsymbol{\Sigma})$$

- The quadratic terms cancel because of shared covariance
- The decision boundary is a linear in  $\mathbf{x}$ :  $\mathbf{w}^T \mathbf{x} + w_0$

$$\mathbf{w} = \boldsymbol{\Sigma}^{-1}(\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2)$$

## Decision Boundary(Shared Covariance)

- The decision boundary is locus of points satisfying  $P[C_1/\mathbf{x}] = P[C_2/\mathbf{x}]$

$$P[C_1] \mathcal{N}(\mathbf{x}/\mu_1, \Sigma) = p[C_2] \mathcal{N}(\mathbf{x}/\mu_2, \Sigma)$$

- The quadratic terms cancel because of shared covariance
- The decision boundary is a linear in  $x$ :  $\mathbf{w}^T \mathbf{x} + w_0$

$$\mathbf{w} = \Sigma^{-1}(\mu_1 - \mu_2) \quad w_0 = -\frac{1}{2}\mu_1^T \Sigma^{-1} \mu_1 + \frac{1}{2}\mu_2^T \Sigma^{-1} \mu_2 + \log \frac{P[C_1]}{P[C_2]}$$

## Decision Boundary(Shared Covariance)

- The decision boundary is locus of points satisfying  $P[C_1/\mathbf{x}] = P[C_2/\mathbf{x}]$

$$P[C_1] \mathcal{N}(\mathbf{x}/\mu_1, \Sigma) = p[C_2] \mathcal{N}(\mathbf{x}/\mu_2, \Sigma)$$

- The quadratic terms cancel because of shared covariance
- The decision boundary is a linear in  $x$ :  $\mathbf{w}^T \mathbf{x} + w_0$

$$\mathbf{w} = \Sigma^{-1}(\mu_1 - \mu_2) \quad w_0 = -\frac{1}{2}\mu_1^T \Sigma^{-1} \mu_1 + \frac{1}{2}\mu_2^T \Sigma^{-1} \mu_2 + \log \frac{P[C_1]}{P[C_2]}$$

- Priors affect only the bias parameters, not the orientation



## Decision Boundary(Shared Covariance)

- The decision boundary is locus of points satisfying  $P[C_1/\mathbf{x}] = P[C_2/\mathbf{x}]$

$$P[C_1] \mathcal{N}(\mathbf{x}/\mu_1, \Sigma) = p[C_2] \mathcal{N}(\mathbf{x}/\mu_2, \Sigma)$$

- The quadratic terms cancel because of shared covariance
- The decision boundary is a linear in  $x$ :  $\mathbf{w}^T \mathbf{x} + w_0$

$$\mathbf{w} = \Sigma^{-1}(\mu_1 - \mu_2) \quad w_0 = -\frac{1}{2}\mu_1^T \Sigma^{-1} \mu_1 + \frac{1}{2}\mu_2^T \Sigma^{-1} \mu_2 + \log \frac{P[C_1]}{P[C_2]}$$

- Priors affect only the bias parameters, not the orientation
- The posterior density of  $C_1$  is given by  $P[C_1/\mathbf{x}] = \sigma(\mathbf{w}^T \mathbf{x} + w_0)$

## Decision Boundary(Shared Covariance)

- The decision boundary is locus of points satisfying  $P[C_1/\mathbf{x}] = P[C_2/\mathbf{x}]$

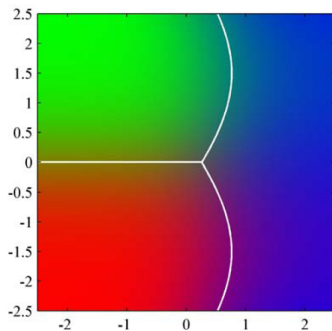
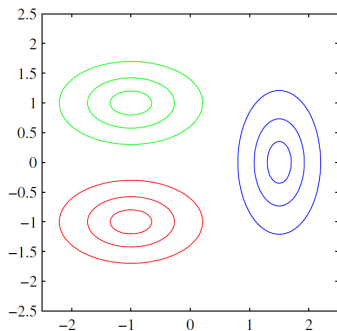
$$P[C_1] \mathcal{N}(\mathbf{x}/\boldsymbol{\mu}_1, \boldsymbol{\Sigma}) = p[C_2] \mathcal{N}(\mathbf{x}/\boldsymbol{\mu}_2, \boldsymbol{\Sigma})$$

- The quadratic terms cancel because of shared covariance
- The decision boundary is a linear in  $\mathbf{x}$ :  $\mathbf{w}^T \mathbf{x} + w_0$

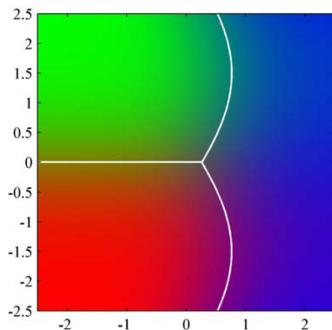
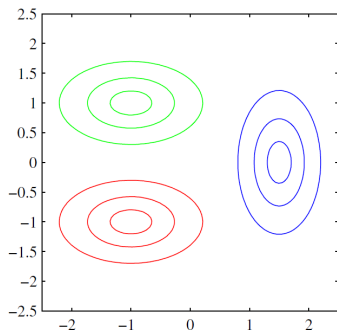
$$\mathbf{w} = \boldsymbol{\Sigma}^{-1}(\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2) \quad w_0 = -\frac{1}{2}\boldsymbol{\mu}_1^T \boldsymbol{\Sigma}^{-1} \boldsymbol{\mu}_1 + \frac{1}{2}\boldsymbol{\mu}_2^T \boldsymbol{\Sigma}^{-1} \boldsymbol{\mu}_2 + \log \frac{P[C_1]}{P[C_2]}$$

- Priors affect only the bias parameters, not the orientation
- The posterior density of  $C_1$  is given by  $P[C_1/\mathbf{x}] = \sigma(\mathbf{w}^T \mathbf{x} + w_0)$
- The decision boundary would be quadratic, if covariance is not shared

# Illustration of Decision Boundaries

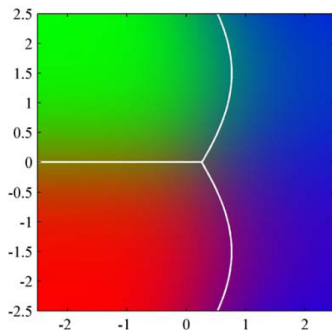
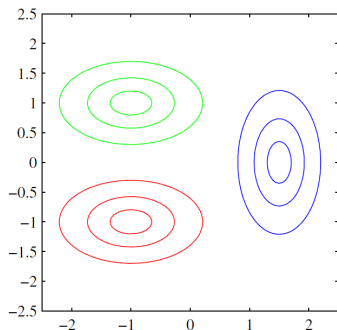


# Illustration of Decision Boundaries



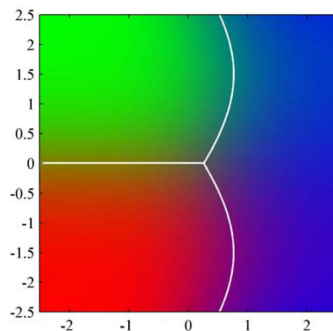
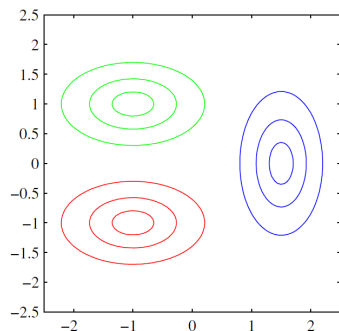
- Red and Green classes share the same covariance matrix

# Illustration of Decision Boundaries



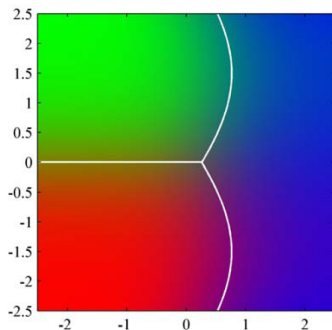
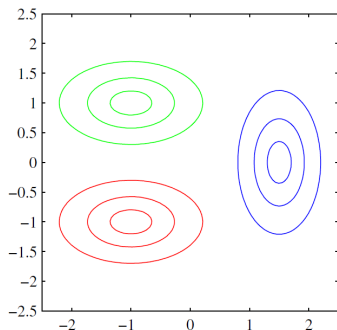
- Red and Green classes share the same covariance matrix
  - Decision boundary is linear

# Illustration of Decision Boundaries



- Red and Green classes share the same covariance matrix
  - Decision boundary is linear
- Blue has a different covariance - decision boundaries are quadratic

# Illustration of Decision Boundaries



- Red and Green classes share the same covariance matrix
  - Decision boundary is linear
- Blue has a different covariance - decision boundaries are quadratic
- Nonlinear decision boundaries can be modeled with pdfs having higher order moments!

# Probabilistic Discriminative Models



# Probabilistic Discriminative Models

- Discriminative models impose a parametric function on posterior

$$P[C_k/\mathbf{x}] = y(\mathbf{x}, \mathbf{W})$$

# Probabilistic Discriminative Models

- Discriminative models impose a parametric function on posterior

$$P[\mathcal{C}_k/\mathbf{x}] = y(\mathbf{x}, \mathbf{W})$$

- The function  $y(\cdot)$  should satisfy the axioms of probability

# Probabilistic Discriminative Models

- Discriminative models impose a parametric function on posterior

$$P[C_k/\mathbf{x}] = y(\mathbf{x}, \mathbf{W})$$

- The function  $y(\cdot)$  should satisfy the axioms of probability
- The posterior probability of the  $k^{\text{th}}$  class is given by (Bayes)

$$P[C_k/\mathbf{x}] = \frac{P[C_k]p(\mathbf{x}/C_k)}{\sum_{j=1}^K P[C_j]p(\mathbf{x}/C_j)}$$

# Probabilistic Discriminative Models

- Discriminative models impose a parametric function on posterior

$$P[C_k/\mathbf{x}] = y(\mathbf{x}, \mathbf{W})$$

- The function  $y(\cdot)$  should satisfy the axioms of probability
- The posterior probability of the  $k^{\text{th}}$  class is given by (Bayes)

$$P[C_k/\mathbf{x}] = \frac{P[C_k]p(\mathbf{x}/C_k)}{\sum_{j=1}^K P[C_j]p(\mathbf{x}/C_j)}$$

- Let  $a_k = \log P[C_k]p(\mathbf{x}/C_k)$  be parameterized as  $a_k = \mathbf{w}_k^T \mathbf{x}$

# Probabilistic Discriminative Models

- Discriminative models impose a parametric function on posterior

$$P[C_k/\mathbf{x}] = y(\mathbf{x}, \mathbf{W})$$

- The function  $y(\cdot)$  should satisfy the axioms of probability
- The posterior probability of the  $k^{\text{th}}$  class is given by (Bayes)

$$P[C_k/\mathbf{x}] = \frac{P[C_k]p(\mathbf{x}/C_k)}{\sum_{j=1}^K P[C_j]p(\mathbf{x}/C_j)}$$

- Let  $a_k = \log P[C_k]p(\mathbf{x}/C_k)$  be parameterized as  $a_k = \mathbf{w}_k^T \mathbf{x}$
- Posterior probability can be expressed as a softmax over activations  $a_k$

$$P[C_k/\mathbf{x}] = \frac{\exp(a_k)}{\sum_{j=1}^K \exp(a_j)} \quad \text{Softmax Fn.}$$

# Binary Classifier

# Binary Classifier

- For binary classifier, we need not evaluate two weight vectors.

# Binary Classifier

- For binary classifier, we need not evaluate two weight vectors.
- The posterior probability can be expressed as

$$p[C_1/\mathbf{x}] = \frac{P[C_1]p(\mathbf{x}/C_1)}{P[C_1]p(\mathbf{x}/C_1) + P[C_2]p(\mathbf{x}/C_2)}$$



# Binary Classifier

- For binary classifier, we need not evaluate two weight vectors.
- The posterior probability can be expressed as

$$p[C_1/\mathbf{x}] = \frac{P[C_1]p(\mathbf{x}/C_1)}{P[C_1]p(\mathbf{x}/C_1) + P[C_2]p(\mathbf{x}/C_2)}$$

- Let  $a = -\log \frac{P[C_1]p(\mathbf{x}/C_1)}{P[C_2]p(\mathbf{x}/C_2)}$  be parameterized as  $a = \mathbf{w}^T \mathbf{x}$

# Binary Classifier

- For binary classifier, we need not evaluate two weight vectors.
- The posterior probability can be expressed as

$$p[\mathcal{C}_1/\mathbf{x}] = \frac{P[\mathcal{C}_1]p(\mathbf{x}/\mathcal{C}_1)}{P[\mathcal{C}_1]p(\mathbf{x}/\mathcal{C}_1) + P[\mathcal{C}_2]p(\mathbf{x}/\mathcal{C}_2)}$$

- Let  $a = -\log \frac{P[\mathcal{C}_1]p(\mathbf{x}/\mathcal{C}_1)}{P[\mathcal{C}_2]p(\mathbf{x}/\mathcal{C}_2)}$  be parameterized as  $a = \mathbf{w}^T \mathbf{x}$
- Posterior probability of  $\mathcal{C}_1$  can be expressed as Sigmoid over  $a$

$$P[\mathcal{C}_1/\mathbf{x}] = \frac{1}{1 + \exp(-a)} = \sigma(a)$$

# Binary Classifier

- For binary classifier, we need not evaluate two weight vectors.
- The posterior probability can be expressed as

$$p[\mathcal{C}_1/\mathbf{x}] = \frac{P[\mathcal{C}_1]p(\mathbf{x}/\mathcal{C}_1)}{P[\mathcal{C}_1]p(\mathbf{x}/\mathcal{C}_1) + P[\mathcal{C}_2]p(\mathbf{x}/\mathcal{C}_2)}$$

- Let  $a = -\log \frac{P[\mathcal{C}_1]p(\mathbf{x}/\mathcal{C}_1)}{P[\mathcal{C}_2]p(\mathbf{x}/\mathcal{C}_2)}$  be parameterized as  $a = \mathbf{w}^T \mathbf{x}$
- Posterior probability of  $\mathcal{C}_1$  can be expressed as Sigmoid over  $a$

$$P[\mathcal{C}_1/\mathbf{x}] = \frac{1}{1 + \exp(-a)} = \sigma(a)$$

- Posterior probability of  $\mathcal{C}_2$  is given by  $P[\mathcal{C}_2/\mathbf{x}] = 1 - \sigma(a)$

# Binary Classifier

- For binary classifier, we need not evaluate two weight vectors.
- The posterior probability can be expressed as

$$p[\mathcal{C}_1/\mathbf{x}] = \frac{P[\mathcal{C}_1]p(\mathbf{x}/\mathcal{C}_1)}{P[\mathcal{C}_1]p(\mathbf{x}/\mathcal{C}_1) + P[\mathcal{C}_2]p(\mathbf{x}/\mathcal{C}_2)}$$

- Let  $a = -\log \frac{P[\mathcal{C}_1]p(\mathbf{x}/\mathcal{C}_1)}{P[\mathcal{C}_2]p(\mathbf{x}/\mathcal{C}_2)}$  be parameterized as  $a = \mathbf{w}^T \mathbf{x}$
- Posterior probability of  $\mathcal{C}_1$  can be expressed as Sigmoid over  $a$

$$P[\mathcal{C}_1/\mathbf{x}] = \frac{1}{1 + \exp(-a)} = \sigma(a)$$

- Posterior probability of  $\mathcal{C}_2$  is given by  $P[\mathcal{C}_2/\mathbf{x}] = 1 - \sigma(a)$
- Derivative of sigmoid function:

$$\frac{d\sigma}{da} = \sigma(a)(1 - \sigma(a))$$

# Logistic Regression

# Logistic Regression

- Training Data:  $\{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_n, t_n), \dots, (\mathbf{x}_N, t_N)\}$

# Logistic Regression

- Training Data:  $\{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_n, t_n), \dots, (\mathbf{x}_N, t_N)\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  are input variables,  $t_n \in \{0, 1\}$  are targets:  $P[C_1/\mathbf{x}_n]$

# Logistic Regression

- Training Data:  $\{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_n, t_n), \dots, (\mathbf{x}_N, t_N)\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  are input variables,  $t_n \in \{0, 1\}$  are targets:  $P[\mathcal{C}_1/\mathbf{x}_n]$
  - Let the target for  $\mathcal{C}_1$  be 1 and  $\mathcal{C}_2$  be 0



# Logistic Regression

- Training Data:  $\{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_n, t_n), \dots, (\mathbf{x}_N, t_N)\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  are input variables,  $t_n \in \{0, 1\}$  are targets:  $P[\mathcal{C}_1/\mathbf{x}_n]$
  - Let the target for  $\mathcal{C}_1$  be 1 and  $\mathcal{C}_2$  be 0
- Let the posterior probability be estimated as

$$\hat{P}[\mathcal{C}_1/\mathbf{x}_n] = y(\mathbf{x}_n, \mathbf{w}) = \sigma(\mathbf{w}^T \mathbf{x}_n) \quad \hat{P}[\mathcal{C}_2/\mathbf{x}_n] = 1 - \sigma(\mathbf{w}^T \mathbf{x}_n)$$

# Logistic Regression

- Training Data:  $\{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_n, t_n), \dots, (\mathbf{x}_N, t_N)\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  are input variables,  $t_n \in \{0, 1\}$  are targets:  $P[\mathcal{C}_1/\mathbf{x}_n]$
  - Let the target for  $\mathcal{C}_1$  be 1 and  $\mathcal{C}_2$  be 0
- Let the posterior probability be estimated as

$$\hat{P}[\mathcal{C}_1/\mathbf{x}_n] = y(\mathbf{x}_n, \mathbf{w}) = \sigma(\mathbf{w}^T \mathbf{x}_n) \quad \hat{P}[\mathcal{C}_2/\mathbf{x}_n] = 1 - \sigma(\mathbf{w}^T \mathbf{x}_n)$$

- Assuming that data points are i.i.d., the likelihood of data is given by

$$P(\mathbf{t}/\mathbf{X}, \mathbf{w}) = \prod_{n=1}^N y_n^{t_n} (1 - y_n)^{1-t_n}$$

# Logistic Regression

- Training Data:  $\{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_n, t_n), \dots, (\mathbf{x}_N, t_N)\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  are input variables,  $t_n \in \{0, 1\}$  are targets:  $P[\mathcal{C}_1/\mathbf{x}_n]$
  - Let the target for  $\mathcal{C}_1$  be 1 and  $\mathcal{C}_2$  be 0
- Let the posterior probability be estimated as

$$\hat{P}[\mathcal{C}_1/\mathbf{x}_n] = y(\mathbf{x}_n, \mathbf{w}) = \sigma(\mathbf{w}^T \mathbf{x}_n) \quad \hat{P}[\mathcal{C}_2/\mathbf{x}_n] = 1 - \sigma(\mathbf{w}^T \mathbf{x}_n)$$

- Assuming that data points are i.i.d., the likelihood of data is given by

$$P(\mathbf{t}/\mathbf{X}, \mathbf{w}) = \prod_{n=1}^N y_n^{t_n} (1 - y_n)^{1-t_n}$$

- $\mathbf{w}$  can be estimated by minimizing negative log of the likelihood, also referred to as *cross-entropy* loss

$$J(\mathbf{w}) = -\log P(\mathbf{t}/\mathbf{X}, \mathbf{w}) = -\sum_{n=1}^N \{t_n \log y_n + (1 - t_n) \log(1 - y_n)\}$$

# Parameter Estimation

# Parameter Estimation

- Model parameters  $\mathbf{w}$  can be updated using gradient descent

$$\mathbf{w}^{new} = \mathbf{w}^{old} - \eta \nabla J(\mathbf{w})$$

# Parameter Estimation

- Model parameters  $\mathbf{w}$  can be updated using gradient descent

$$\mathbf{w}^{new} = \mathbf{w}^{old} - \eta \nabla J(\mathbf{w})$$

- The 1<sup>st</sup> and 2<sup>nd</sup> order derivatives of the loss function are given by

# Parameter Estimation

- Model parameters  $\mathbf{w}$  can be updated using gradient descent

$$\mathbf{w}^{new} = \mathbf{w}^{old} - \eta \nabla J(\mathbf{w})$$

- The 1<sup>st</sup> and 2<sup>nd</sup> order derivatives of the loss function are given by

$$\nabla J(\mathbf{w}) = \sum_{n=1}^N \mathbf{x}_n (y_n - t_n) = \mathbf{X}^T (\mathbf{y} - \mathbf{t})$$

# Parameter Estimation

- Model parameters  $\mathbf{w}$  can be updated using gradient descent

$$\mathbf{w}^{new} = \mathbf{w}^{old} - \eta \nabla J(\mathbf{w})$$

- The 1<sup>st</sup> and 2<sup>nd</sup> order derivatives of the loss function are given by

$$\nabla J(\mathbf{w}) = \sum_{n=1}^N \mathbf{x}_n (y_n - t_n) = \mathbf{X}^T (\mathbf{y} - \mathbf{t})$$

$$\mathbf{H} = \nabla^2 J(\mathbf{w}) = \sum_{n=1}^N \mathbf{x}_n y_n (1 - y_n) \mathbf{x}_n^T = \mathbf{X}^T \mathbf{R} \mathbf{X}$$



# Parameter Estimation

- Model parameters  $\mathbf{w}$  can be updated using gradient descent

$$\mathbf{w}^{new} = \mathbf{w}^{old} - \eta \nabla J(\mathbf{w})$$

- The 1<sup>st</sup> and 2<sup>nd</sup> order derivatives of the loss function are given by

$$\nabla J(\mathbf{w}) = \sum_{n=1}^N \mathbf{x}_n (y_n - t_n) = \mathbf{X}^T (\mathbf{y} - \mathbf{t})$$

$$\mathbf{H} = \nabla^2 J(\mathbf{w}) = \sum_{n=1}^N \mathbf{x}_n y_n (1 - y_n) \mathbf{x}_n^T = \mathbf{X}^T \mathbf{R} \mathbf{X}$$

where  $\mathbf{R}$  is a diagonal matrix with entries  $R_{nn} = y_n(1 - y_n)$

# Parameter Estimation

- Model parameters  $\mathbf{w}$  can be updated using gradient descent

$$\mathbf{w}^{new} = \mathbf{w}^{old} - \eta \nabla J(\mathbf{w})$$

- The 1<sup>st</sup> and 2<sup>nd</sup> order derivatives of the loss function are given by

$$\nabla J(\mathbf{w}) = \sum_{n=1}^N \mathbf{x}_n (y_n - t_n) = \mathbf{X}^T (\mathbf{y} - \mathbf{t})$$

$$\mathbf{H} = \nabla^2 J(\mathbf{w}) = \sum_{n=1}^N \mathbf{x}_n y_n (1 - y_n) \mathbf{x}_n^T = \mathbf{X}^T \mathbf{R} \mathbf{X}$$

where  $\mathbf{R}$  is a diagonal matrix with entries  $R_{nn} = y_n(1 - y_n)$

- Hessian matrix varies depending on parameters  $\mathbf{w}$  through  $\mathbf{R}$

# Parameter Estimation

- Model parameters  $\mathbf{w}$  can be updated using gradient descent

$$\mathbf{w}^{new} = \mathbf{w}^{old} - \eta \nabla J(\mathbf{w})$$

- The 1<sup>st</sup> and 2<sup>nd</sup> order derivatives of the loss function are given by

$$\nabla J(\mathbf{w}) = \sum_{n=1}^N \mathbf{x}_n (y_n - t_n) = \mathbf{X}^T (\mathbf{y} - \mathbf{t})$$

$$\mathbf{H} = \nabla^2 J(\mathbf{w}) = \sum_{n=1}^N \mathbf{x}_n y_n (1 - y_n) \mathbf{x}_n^T = \mathbf{X}^T \mathbf{R} \mathbf{X}$$

where  $\mathbf{R}$  is a diagonal matrix with entries  $R_{nn} = y_n(1 - y_n)$

- Hessian matrix varies depending on parameters  $\mathbf{w}$  through  $\mathbf{R}$
- Since  $0 < y_n < 1$ , Hessian matrix  $\mathbf{H}$  is positive definite:  $\mathbf{u}^T \mathbf{H} \mathbf{u} > 0$

# Parameter Estimation

- Model parameters  $\mathbf{w}$  can be updated using gradient descent

$$\mathbf{w}^{new} = \mathbf{w}^{old} - \eta \nabla J(\mathbf{w})$$

- The 1<sup>st</sup> and 2<sup>nd</sup> order derivatives of the loss function are given by

$$\nabla J(\mathbf{w}) = \sum_{n=1}^N \mathbf{x}_n (y_n - t_n) = \mathbf{X}^T (\mathbf{y} - \mathbf{t})$$

$$\mathbf{H} = \nabla^2 J(\mathbf{w}) = \sum_{n=1}^N \mathbf{x}_n y_n (1 - y_n) \mathbf{x}_n^T = \mathbf{X}^T \mathbf{R} \mathbf{X}$$

where  $\mathbf{R}$  is a diagonal matrix with entries  $R_{nn} = y_n(1 - y_n)$

- Hessian matrix varies depending on parameters  $\mathbf{w}$  through  $\mathbf{R}$
- Since  $0 < y_n < 1$ , Hessian matrix  $\mathbf{H}$  is positive definite:  $\mathbf{u}^T \mathbf{H} \mathbf{u} > 0$
- Error function  $J(\mathbf{w})$  is convex in  $\mathbf{w} \implies$  admits unique minimum

# Iterative Reweighted Least Squares

$$\mathbf{w}^{new} = \mathbf{w}^{old} - \left( \mathbf{X}^T \mathbf{R} \mathbf{X} \right)^{-1} \mathbf{X}^T (\mathbf{y} - \mathbf{t})$$

# Iterative Reweighted Least Squares

$$\begin{aligned}\mathbf{w}^{new} &= \mathbf{w}^{old} - \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \mathbf{X}^T (\mathbf{y} - \mathbf{t}) \\ &= \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \left(\mathbf{X}^T \mathbf{R} \mathbf{X} \mathbf{w}^{old} - \mathbf{X}^T (\mathbf{y} - \mathbf{t})\right)\end{aligned}$$

# Iterative Reweighted Least Squares

$$\begin{aligned}\mathbf{w}^{new} &= \mathbf{w}^{old} - \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \mathbf{X}^T (\mathbf{y} - \mathbf{t}) \\ &= \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \left(\mathbf{X}^T \mathbf{R} \mathbf{X} \mathbf{w}^{old} - \mathbf{X}^T (\mathbf{y} - \mathbf{t})\right) \\ &= \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \mathbf{X}^T \mathbf{R} \left(\mathbf{X} \mathbf{w}^{old} - \mathbf{R}^{-1} (\mathbf{y} - \mathbf{t})\right)\end{aligned}$$

# Iterative Reweighted Least Squares

$$\begin{aligned}\mathbf{w}^{new} &= \mathbf{w}^{old} - \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \mathbf{X}^T (\mathbf{y} - \mathbf{t}) \\ &= \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \left(\mathbf{X}^T \mathbf{R} \mathbf{X} \mathbf{w}^{old} - \mathbf{X}^T (\mathbf{y} - \mathbf{t})\right) \\ &= \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \mathbf{X}^T \mathbf{R} \left(\mathbf{X} \mathbf{w}^{old} - \mathbf{R}^{-1} (\mathbf{y} - \mathbf{t})\right) \\ &= \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \mathbf{X}^T \mathbf{R} \mathbf{z}\end{aligned}$$



# Iterative Reweighted Least Squares

$$\begin{aligned}\mathbf{w}^{new} &= \mathbf{w}^{old} - \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \mathbf{X}^T (\mathbf{y} - \mathbf{t}) \\ &= \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \left(\mathbf{X}^T \mathbf{R} \mathbf{X} \mathbf{w}^{old} - \mathbf{X}^T (\mathbf{y} - \mathbf{t})\right) \\ &= \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \mathbf{X}^T \mathbf{R} \left(\mathbf{X} \mathbf{w}^{old} - \mathbf{R}^{-1} (\mathbf{y} - \mathbf{t})\right) \\ &= \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \mathbf{X}^T \mathbf{R} \mathbf{z}\end{aligned}$$

where  $\mathbf{z} = (\mathbf{X} \mathbf{w}^{old} - \mathbf{R}^{-1} (\mathbf{y} - \mathbf{t}))$ .

# Iterative Reweighted Least Squares

$$\begin{aligned}\mathbf{w}^{new} &= \mathbf{w}^{old} - \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \mathbf{X}^T (\mathbf{y} - \mathbf{t}) \\ &= \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \left(\mathbf{X}^T \mathbf{R} \mathbf{X} \mathbf{w}^{old} - \mathbf{X}^T (\mathbf{y} - \mathbf{t})\right) \\ &= \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \mathbf{X}^T \mathbf{R} \left(\mathbf{X} \mathbf{w}^{old} - \mathbf{R}^{-1} (\mathbf{y} - \mathbf{t})\right) \\ &= \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \mathbf{X}^T \mathbf{R} \mathbf{z}\end{aligned}$$

where  $\mathbf{z} = (\mathbf{X} \mathbf{w}^{old} - \mathbf{R}^{-1} (\mathbf{y} - \mathbf{t}))$ .

- The solution takes the form of normal equations of weighted LS.

# Iterative Reweighted Least Squares

$$\begin{aligned}\mathbf{w}^{new} &= \mathbf{w}^{old} - \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \mathbf{X}^T (\mathbf{y} - \mathbf{t}) \\ &= \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \left(\mathbf{X}^T \mathbf{R} \mathbf{X} \mathbf{w}^{old} - \mathbf{X}^T (\mathbf{y} - \mathbf{t})\right) \\ &= \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \mathbf{X}^T \mathbf{R} \left(\mathbf{X} \mathbf{w}^{old} - \mathbf{R}^{-1} (\mathbf{y} - \mathbf{t})\right) \\ &= \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \mathbf{X}^T \mathbf{R} \mathbf{z}\end{aligned}$$

where  $\mathbf{z} = (\mathbf{X} \mathbf{w}^{old} - \mathbf{R}^{-1} (\mathbf{y} - \mathbf{t}))$ .

- The solution takes the form of normal equations of weighted LS.
- However, the weighing matrix  $\mathbf{R}$  is not constant, but depends on  $\mathbf{w}$

# Iterative Reweighted Least Squares

$$\begin{aligned}\mathbf{w}^{new} &= \mathbf{w}^{old} - \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \mathbf{X}^T (\mathbf{y} - \mathbf{t}) \\ &= \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \left(\mathbf{X}^T \mathbf{R} \mathbf{X} \mathbf{w}^{old} - \mathbf{X}^T (\mathbf{y} - \mathbf{t})\right) \\ &= \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \mathbf{X}^T \mathbf{R} \left(\mathbf{X} \mathbf{w}^{old} - \mathbf{R}^{-1} (\mathbf{y} - \mathbf{t})\right) \\ &= \left(\mathbf{X}^T \mathbf{R} \mathbf{X}\right)^{-1} \mathbf{X}^T \mathbf{R} \mathbf{z}\end{aligned}$$

where  $\mathbf{z} = (\mathbf{X} \mathbf{w}^{old} - \mathbf{R}^{-1} (\mathbf{y} - \mathbf{t}))$ .

- The solution takes the form of normal equations of weighted LS.
- However, the weighing matrix  $\mathbf{R}$  is not constant, but depends on  $\mathbf{w}$
- Hence, the normal equations need to be applied iteratively.

# Multiclass Logistic Regression (Homework)

# Multiclass Logistic Regression (Homework)

- Consider multiclass data examples denoted by  $X = \{(\mathbf{x}_{1:N}, \mathbf{t}_{1:N})\}$

# Multiclass Logistic Regression (Homework)

- Consider multiclass data examples denoted by  $X = \{(\mathbf{x}_{1:N}, \mathbf{t}_{1:N})\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  represents input observations

# Multiclass Logistic Regression (Homework)

- Consider multiclass data examples denoted by  $X = \{(\mathbf{x}_{1:N}, \mathbf{t}_{1:N})\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  represents input observations
  - $\mathbf{t}_n$  is a  $K$ -dim one-hot vector denoting the class posteriors



# Multiclass Logistic Regression (Homework)

- Consider multiclass data examples denoted by  $X = \{(\mathbf{x}_{1:N}, \mathbf{t}_{1:N})\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  represents input observations
  - $\mathbf{t}_n$  is a  $K$ -dim one-hot vector denoting the class posteriors
- For this case, the posterior probabilities can be estimated as

$$y(\mathbf{x}_n, \mathbf{w}_k) = \hat{P}[C_k/\mathbf{x}_n] = \frac{\exp(a_{nk})}{\sum_{j=1}^K \exp(a_{nj})} \quad a_{nk} = \mathbf{w}_k^T \mathbf{x}_n$$

# Multiclass Logistic Regression (Homework)

- Consider multiclass data examples denoted by  $X = \{(\mathbf{x}_{1:N}, \mathbf{t}_{1:N})\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  represents input observations
  - $\mathbf{t}_n$  is a  $K$ -dim one-hot vector denoting the class posteriors
- For this case, the posterior probabilities can be estimated as

$$y(\mathbf{x}_n, \mathbf{w}_k) = \hat{P}[C_k/\mathbf{x}_n] = \frac{\exp(a_{nk})}{\sum_{j=1}^K \exp(a_{nj})} \quad a_{nk} = \mathbf{w}_k^T \mathbf{x}_n$$

- The likelihood function can be written as

$$P[\mathbf{T}/\mathbf{X}, \mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_K] = \prod_{n=1}^N \prod_{k=1}^K y_{nk}^{t_{nk}}$$

# Multiclass Logistic Regression (Homework)

- Consider multiclass data examples denoted by  $X = \{(\mathbf{x}_{1:N}, \mathbf{t}_{1:N})\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  represents input observations
  - $\mathbf{t}_n$  is a  $K$ -dim one-hot vector denoting the class posteriors
- For this case, the posterior probabilities can be estimated as

$$y(\mathbf{x}_n, \mathbf{w}_k) = \hat{P}[C_k/\mathbf{x}_n] = \frac{\exp(a_{nk})}{\sum_{j=1}^K \exp(a_{nj})} \quad a_{nk} = \mathbf{w}_k^T \mathbf{x}_n$$

- The likelihood function can be written as

$$P[\mathbf{T}/\mathbf{X}, \mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_K] = \prod_{n=1}^N \prod_{k=1}^K y_{nk}^{t_{nk}}$$
$$J(\mathbf{W}) = -P[\mathbf{T}/\mathbf{X}, \mathbf{W}] = -\sum_{n=1}^N \sum_{k=1}^K t_{nk} \log y_{nk} \quad (\text{CE})$$

# Multiclass Logistic Regression (Homework)

- Consider multiclass data examples denoted by  $X = \{(\mathbf{x}_{1:N}, \mathbf{t}_{1:N})\}$ 
  - $\mathbf{x}_n \in \mathbb{R}^D$  represents input observations
  - $\mathbf{t}_n$  is a  $K$ -dim one-hot vector denoting the class posteriors
- For this case, the posterior probabilities can be estimated as

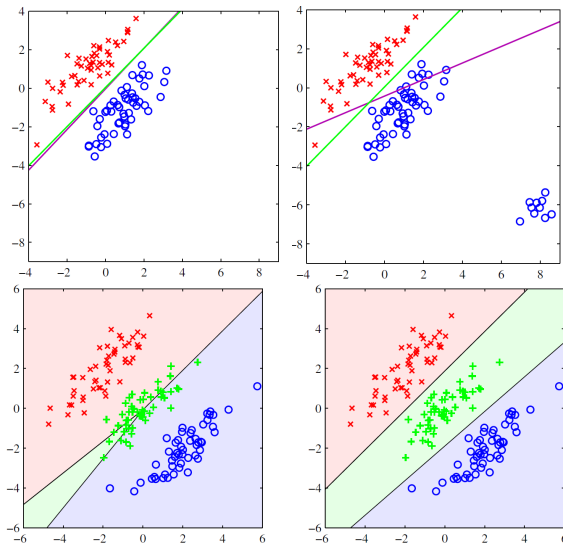
$$y(\mathbf{x}_n, \mathbf{w}_k) = \hat{P}[C_k/\mathbf{x}_n] = \frac{\exp(a_{nk})}{\sum_{j=1}^K \exp(a_{nj})} \quad a_{nk} = \mathbf{w}_k^T \mathbf{x}_n$$

- The likelihood function can be written as

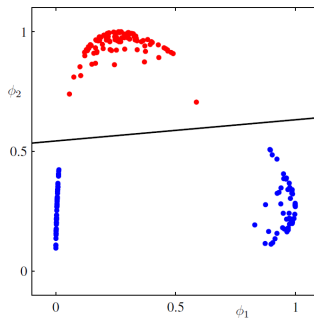
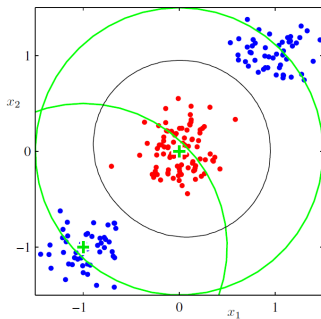
$$P[\mathbf{T}/\mathbf{X}, \mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_K] = \prod_{n=1}^N \prod_{k=1}^K y_{nk}^{t_{nk}}$$

$$J(\mathbf{W}) = -P[\mathbf{T}/\mathbf{X}, \mathbf{W}] = -\sum_{n=1}^N \sum_{k=1}^K t_{nk} \log y_{nk} \quad (\text{CE})$$

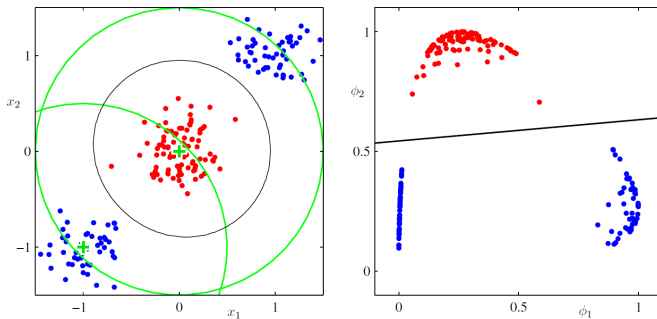
# Illustration of Logistic Regression



# Nonlinear Decision Boundary (Transformed Space $\phi(\cdot)$ )

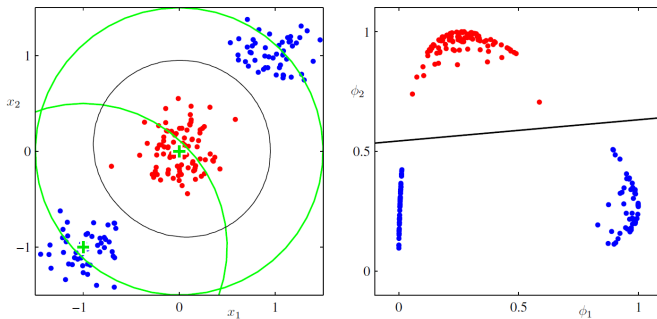


# Nonlinear Decision Boundary (Transformed Space $\phi(.)$ )



- Two Gaussian kernels are used to transform the data

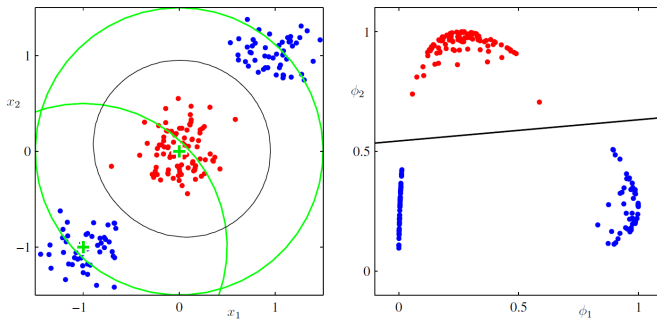
# Nonlinear Decision Boundary (Transformed Space $\phi(\cdot)$ )



- Two Gaussian kernels are used to transform the data
- In general, we need to design the kernel  $\phi$  from the data.

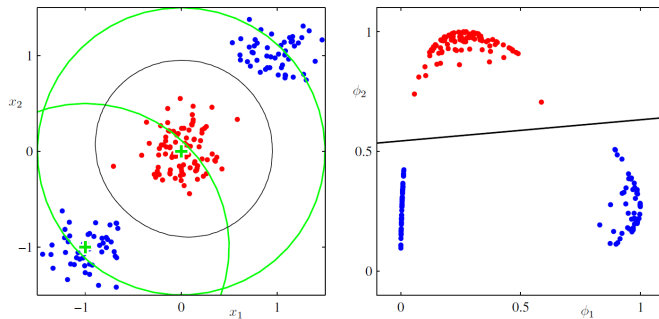


## Nonlinear Decision Boundary (Transformed Space $\phi(\cdot)$ )



- Two Gaussian kernels are used to transform the data
- In general, we need to design the kernel  $\phi$  from the data.
- DNN can be used to learn the optimal transformation from the data

# Nonlinear Decision Boundary (Transformed Space $\phi(\cdot)$ )



- Two Gaussian kernels are used to transform the data
- In general, we need to design the kernel  $\phi$  from the data.
- DNN can be used to learn the optimal transformation from the data
- Last layer of a DNN classifier performs logistic regression

# Summary of Linear Classifiers

- Assumption: Classes are separable by linear hyperplanes

# Summary of Linear Classifiers

- Assumption: Classes are separable by linear hyperplanes
- Linear discriminant functions model the separating hyperplanes
  - Least squares, Fisher discriminant, Perceptron, SVM (later)
  - May not work even if classes are separable because of outliers

# Summary of Linear Classifiers

- Assumption: Classes are separable by linear hyperplanes
- Linear discriminant functions model the separating hyperplanes
  - Least squares, Fisher discriminant, Perceptron, SVM (later)
  - May not work even if classes are separable because of outliers
- Generative models estimate posteriors from priors and CCDs
  - ML or MAP estimators are employed to model CCD
  - Don't consider other class examples while estimating CCD

# Summary of Linear Classifiers

- Assumption: Classes are separable by linear hyperplanes
- Linear discriminant functions model the separating hyperplanes
  - Least squares, Fisher discriminant, Perceptron, SVM (later)
  - May not work even if classes are separable because of outliers
- Generative models estimate posteriors from priors and CCDs
  - ML or MAP estimators are employed to model CCD
  - Don't consider other class examples while estimating CCD
- Discriminative models directly estimate posterior probabilities
  - Binary classes - logistic activation - binary cross entropy
  - Multiple classes - softmax activation - cross entropy
  - Rely on discriminative features - vulnerable to adversarial examples

# Summary of Linear Classifiers

- Assumption: Classes are separable by linear hyperplanes
- Linear discriminant functions model the separating hyperplanes
  - Least squares, Fisher discriminant, Perceptron, SVM (later)
  - May not work even if classes are separable because of outliers
- Generative models estimate posteriors from priors and CCDs
  - ML or MAP estimators are employed to model CCD
  - Don't consider other class examples while estimating CCD
- Discriminative models directly estimate posterior probabilities
  - Binary classes - logistic activation - binary cross entropy
  - Multiple classes - softmax activation - cross entropy
  - Rely on discriminative features - vulnerable to adversarial examples
- If decision boundary is not linear, apply these techniques on  $\phi(\mathbf{x})$ 
  - Neural networks offer a way of learning  $\phi(\mathbf{x})$  from data

# Thank You!